


# PostgreSQL Distributed

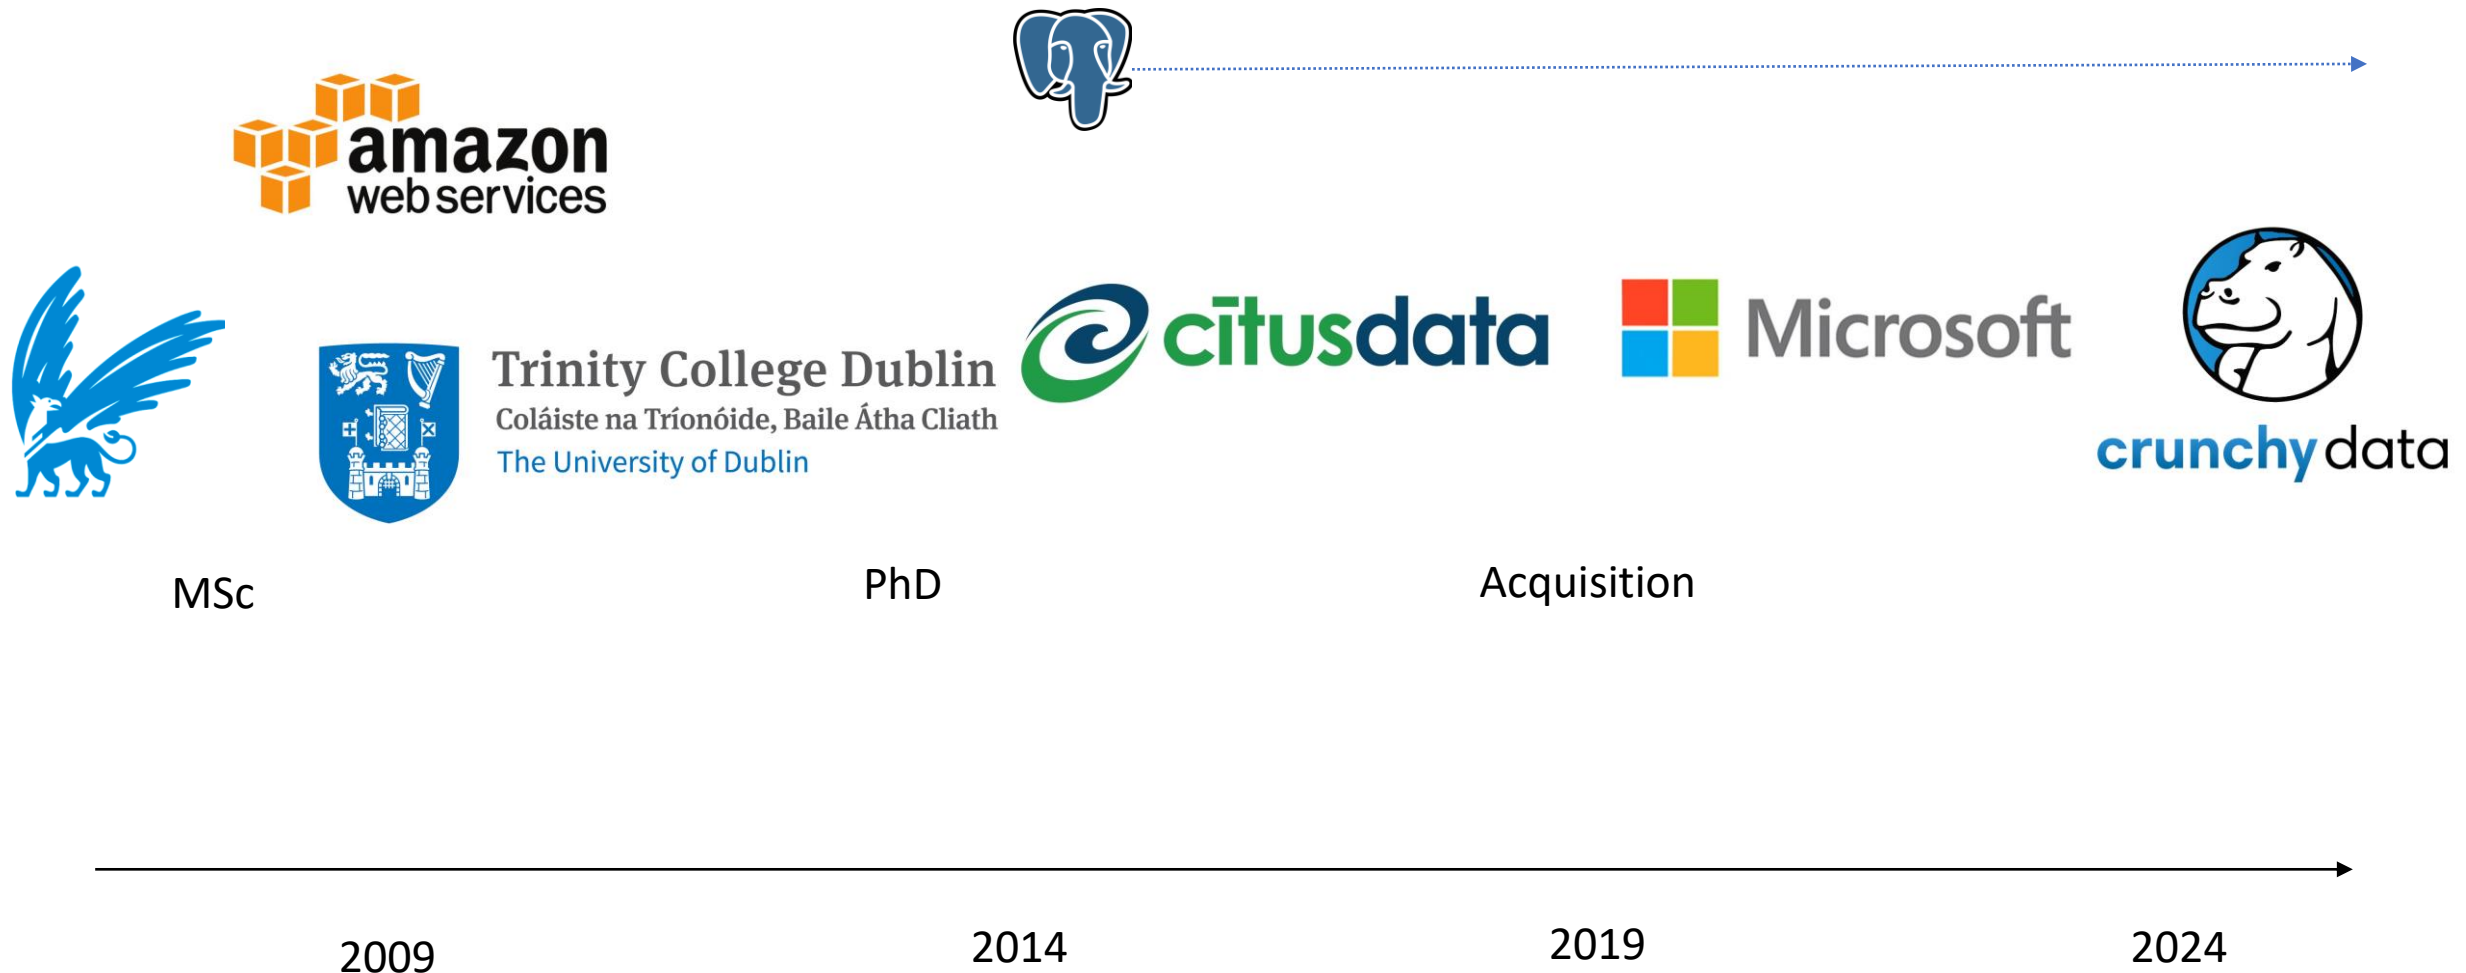
An overview of architectures and trade-offs

**Marco Slot** - [marco.slot@gmail.com](mailto:marco.slot@gmail.com)

Formerly: founding engineer at Citus Data, architect at Microsoft



# My career timeline



# My background in Distributed PostgreSQL

Developed Citus since 2014 <https://github.com/citusdata/citus>

Citus is a PostgreSQL *extension* that adds the ability to transparently distribute or replicate tables across a cluster of PostgreSQL servers.

See “Citus: Distributed PostgreSQL for Data-Intensive Applications”  
- SIGMOD '21

Also pg\_cron, azure\_storage, other PostgreSQL extensions.



# Today's talk on PostgreSQL Distributed

Many distributed database talks discuss algorithms for distributed query planning, transactions, etc.

In distributed systems, trade-offs are more important than algorithms.

Vendors and even many papers rarely talk about trade-offs.

Moreover, many different PostgreSQL distributed system architectures with different trade-offs exist.

**Experiment:** Discuss PostgreSQL distributed systems architecture trade-offs by example.

# Single machine PostgreSQL

PostgreSQL on a single machine can be incredibly fast


No network latency

Millions of IOPS

Microsecond disk latency

Low cost / fast hardware

Can co-locate application server



# Single machine PostgreSQL?

PostgreSQL on a single machine comes with operational hazards


Machine/DC failure (downtime)

Disk failure (data loss)

System overload (difficult to scale)

Accidental data deletion (need backups)

Disk full (downtime)






# PostgreSQL Distributed (in the cloud)

Fixing the operational hazards of single node PostgreSQL requires a distributed set up.

The cloud enables flexible distributed set ups, with resources shared between customers for high efficiency and resiliency.



# Goals of distributed database architecture

Goal: Offer same functionality and transactional semantics as single node RDBMS, with superior properties

Mechanisms:

<b>Replication</b>	- Place copies of data on different machines
<b>Distribution</b>	- Place partitions of data on different machines
<b>Decentralization</b>	- Place different DBMS activities on different machines

Reality: Concessions in terms of performance, transactional semantics, functionality, and/or operational complexity



# PostgreSQL Distributed Layers

Distributed architectures can hook in at different layers — many are orthogonal!

Client	Manual sharding, load balancing, write to multiple endpoints
Pooler	Load balancing and sharding (e.g. pgbouncer, pgcat)
Query engine	Transparent sharding (e.g. Citus, Aurora limitless)
Logical data layer	Active-active, federation (e.g. BDR, postgres_fdw)
Storage manager	DBMS-optimized cloud storage (e.g. Aurora, Neon)
Data files, WAL	Read replicas, hot standby
Disk	Cloud block storage (e.g. Amazon EBS, Azure Premium SSD)

# Practical view of Distributed PostgreSQL

Today we will cover:

- Network-attached cloud storage
- Read replicas & hot standby
- DBMS-optimized cloud storage
- Transparent Sharding
- Active-active deployments
- Distributed key-value stores with SQL

## Two questions:

1) What are the trade-offs?


Latency, Efficiency, Cost, Scalability, Availability, Consistency, Complexity, ...

2) For which workloads?

Lookups, analytical queries, small updates, large transforms, batch loads, ...

# The perils of latency: Synchronous protocol

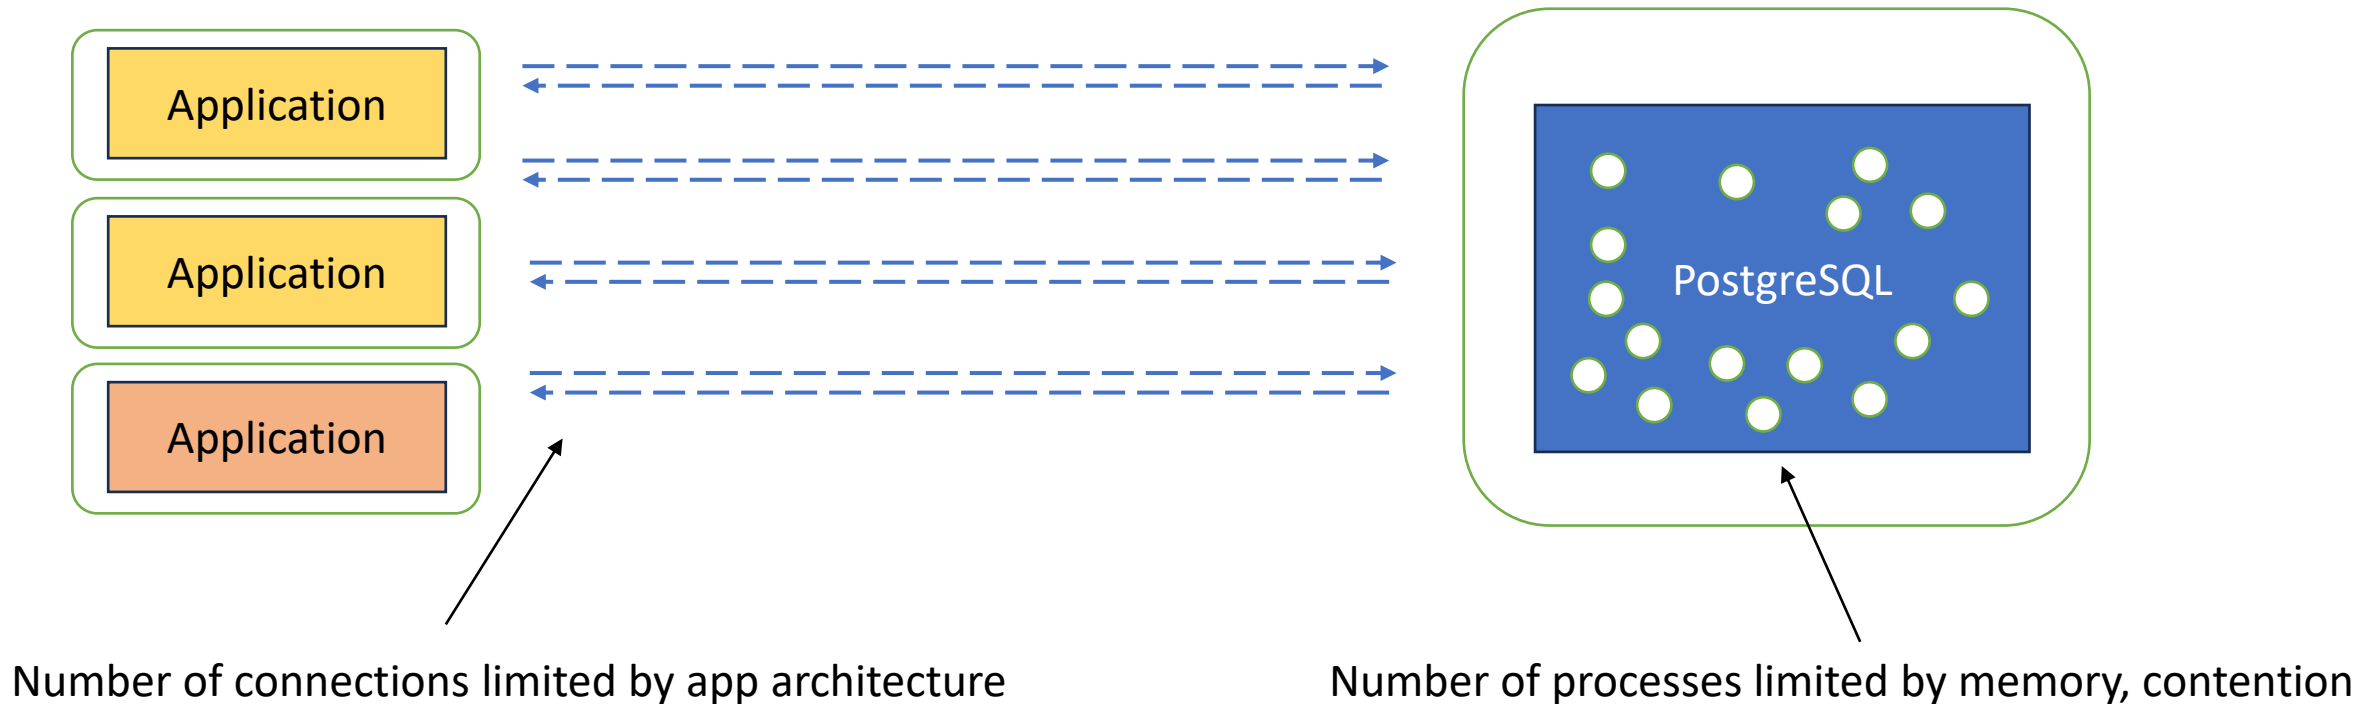
Transactions are performed step-by-step on each session.




Max throughput per session =  $1 / \text{avg. response time}$

# The perils of latency: Connection limits

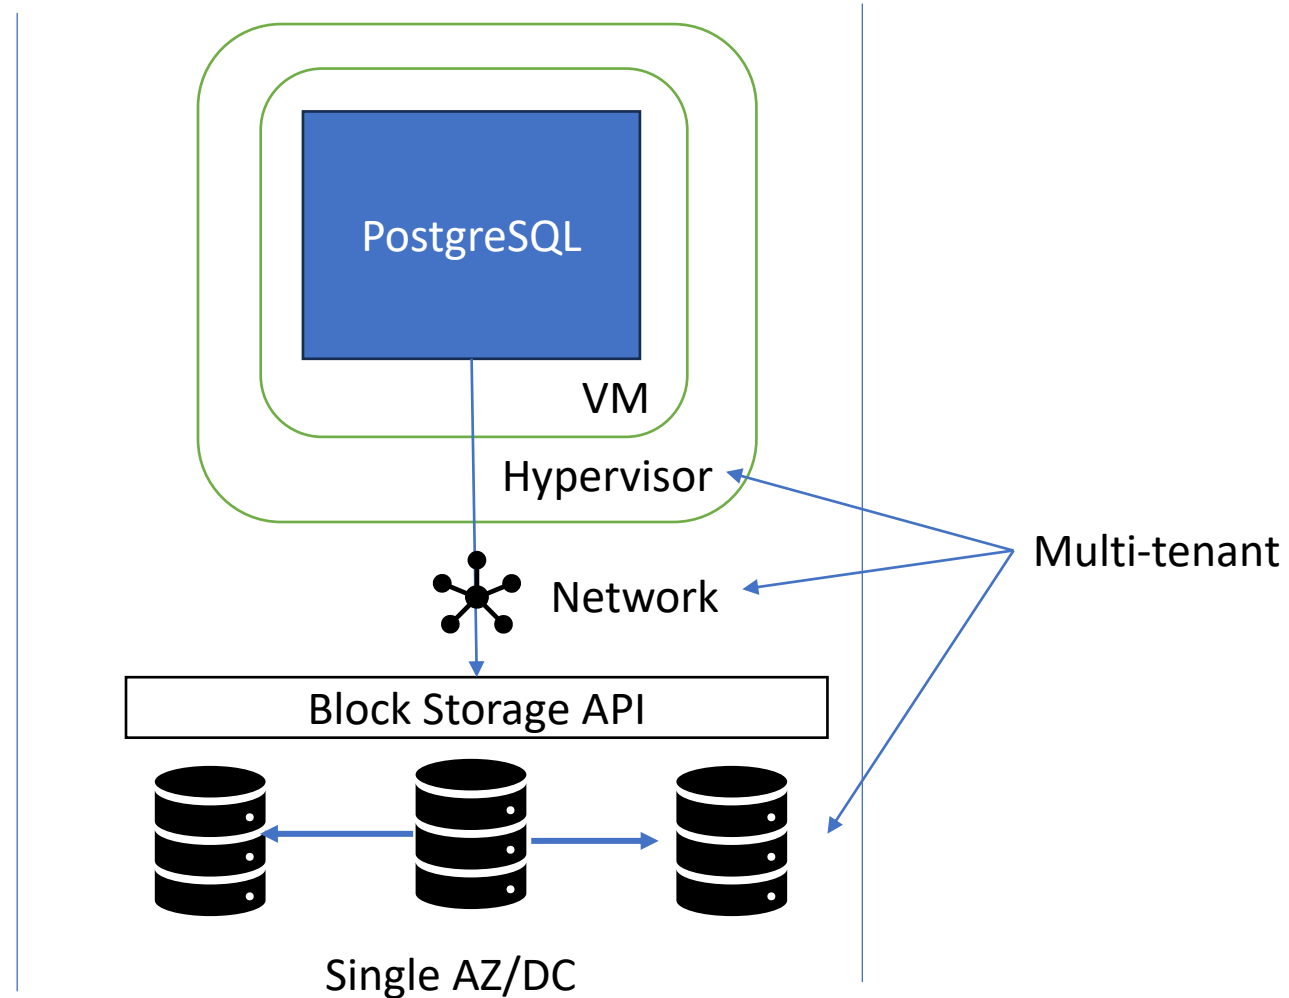
Max overall throughput:  $\# \text{sessions} / \text{avg.response time}$



# Network-attached block storage



# Network-attached block storage






# Network-attached storage

## Pros:


- Higher durability (replication)
- Higher uptime (replace VM, reattach)
- Fast backups and replica creation (snapshots)
- Disk is resizable

## Cons:

- Higher disk latency, esp. for writes ( $\sim 10\mu\text{s}$   $\rightarrow$   $\sim 1000\mu\text{s}$ )
- Lower IOPS ( $\sim 1\text{M}$   $\rightarrow$   $\sim 10\text{k}$  IOPS)
- Crash recovery on restart takes time
- Cost can be high

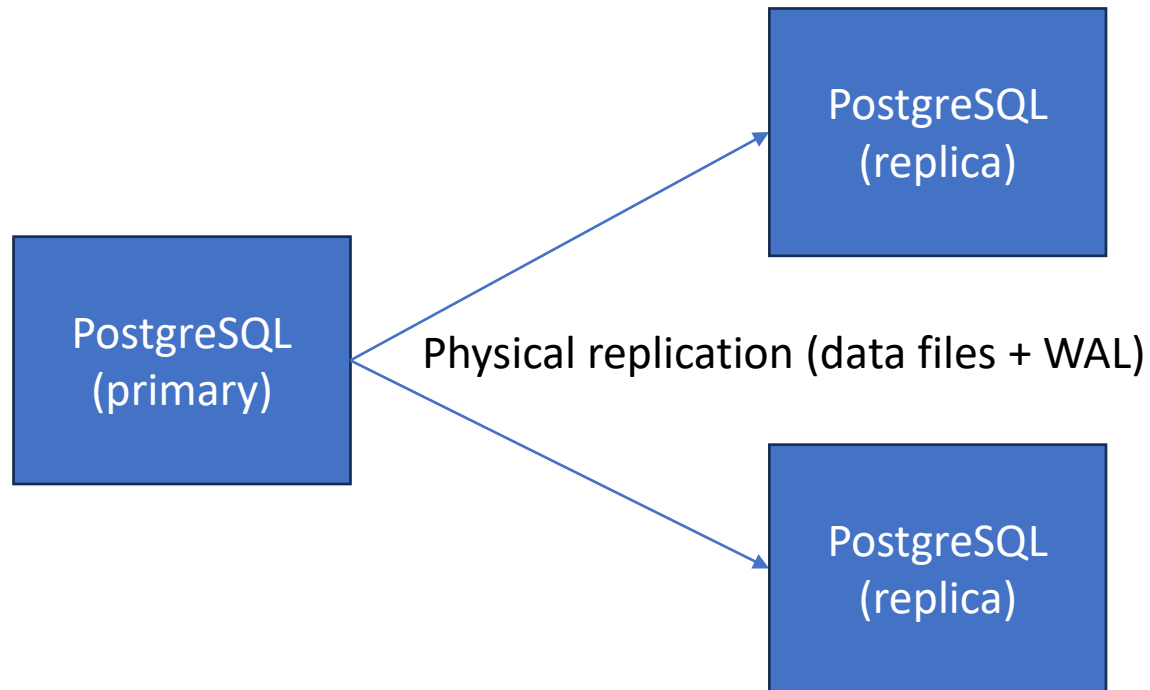


# Read replicas



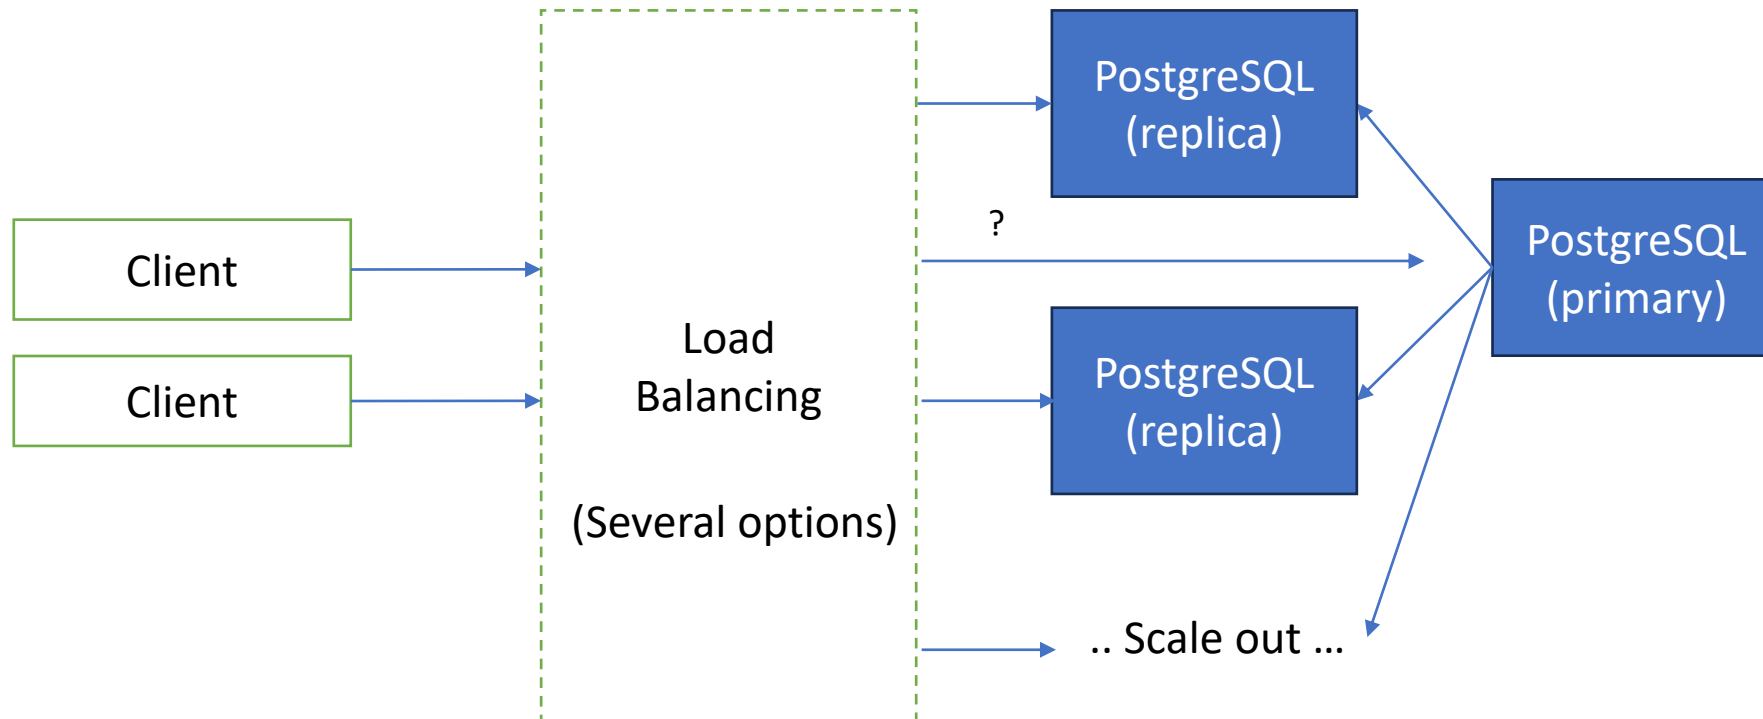
# Read replicas

Readable replicas can help you scale read throughput, reduce latency through cross-region replication, improve availability through auto-failover.



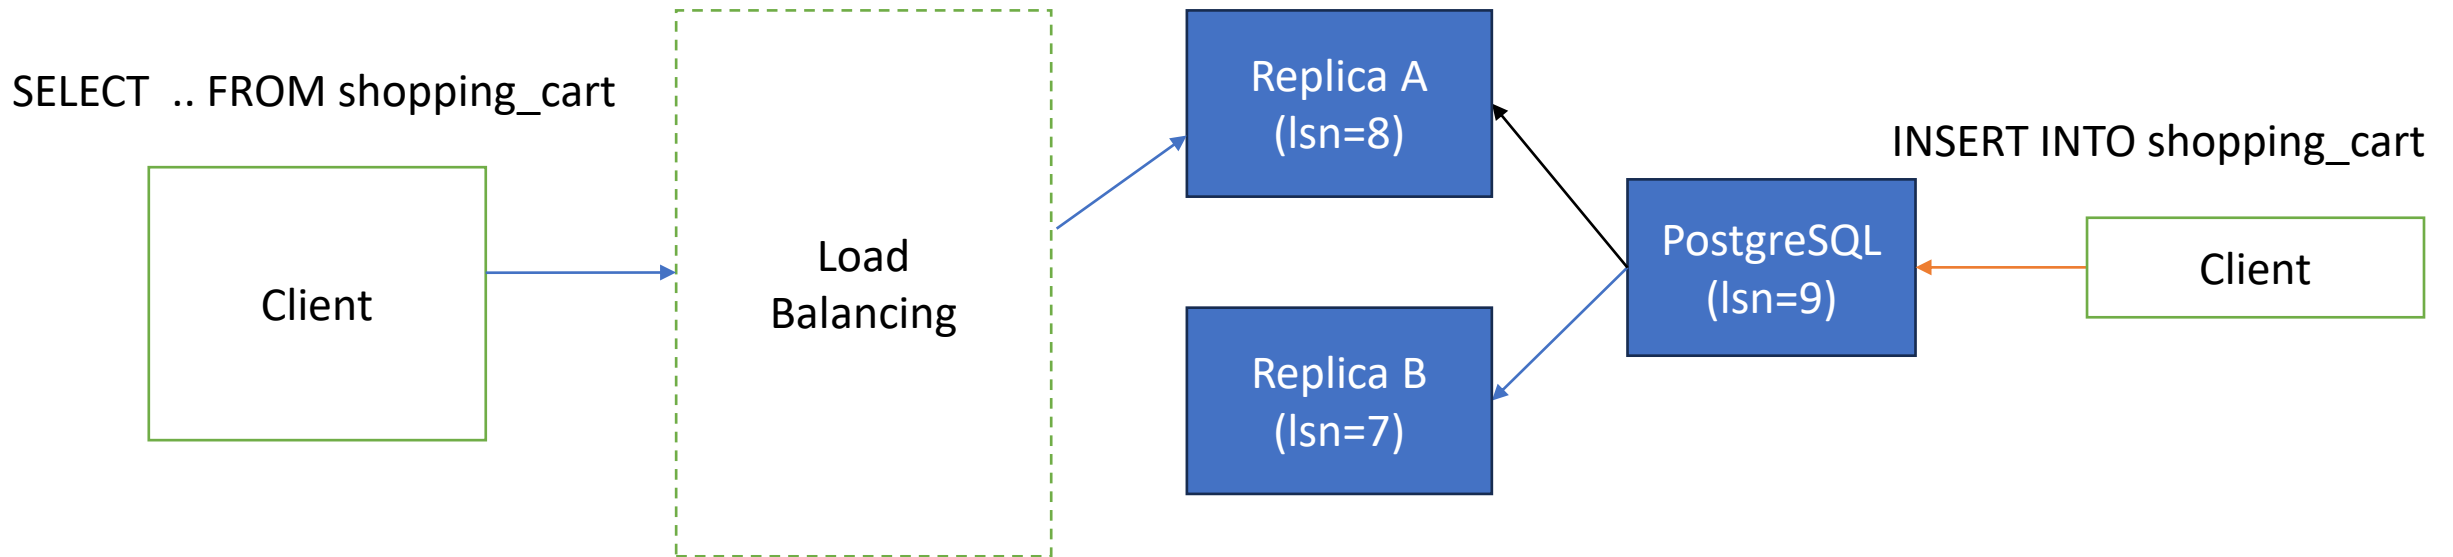
# Scaling read throughput

Readable replicas can help you scale read throughput (when reads are CPU or I/O bottlenecked) by load balancing queries across replicas.



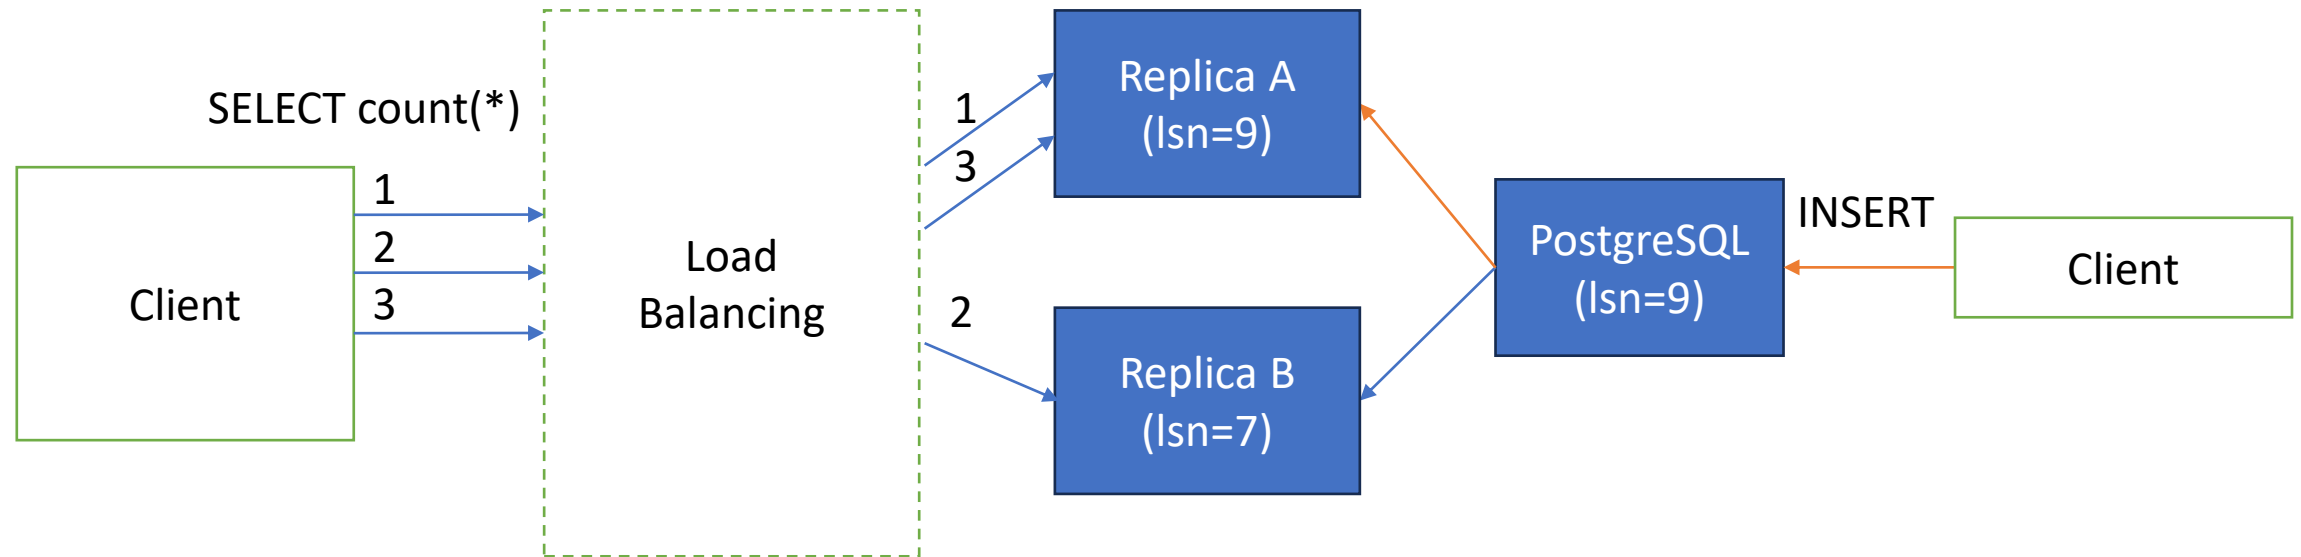
# Eventual read-your-writes consistency

Read replicas can be behind on the primary, cannot always read your writes.



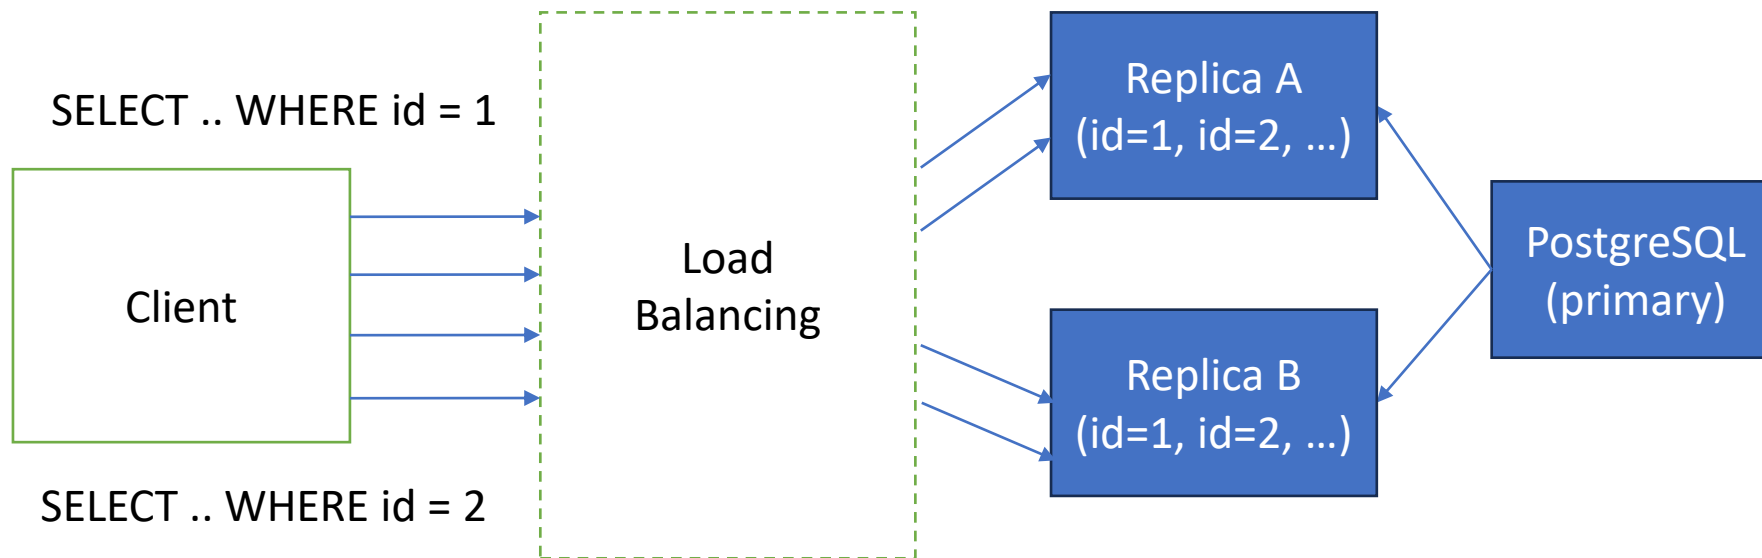
# No monotonic read consistency

Load-balancing across read replicas will cause you to go back-and-forth in time.



# Poor cache usage

If all replicas are equal, they all have the same stuff in cache



If working set  $\gg$  memory, all replicas get bottlenecked on disk I/O.



# Read scaling trade-offs

## **Pros:**

Read throughput scales linearly

Low latency stale reads if read replica is closer than primary

Lower load on primary


## **Cons:**

Eventual read-your-writes consistency

No monotonic read consistency

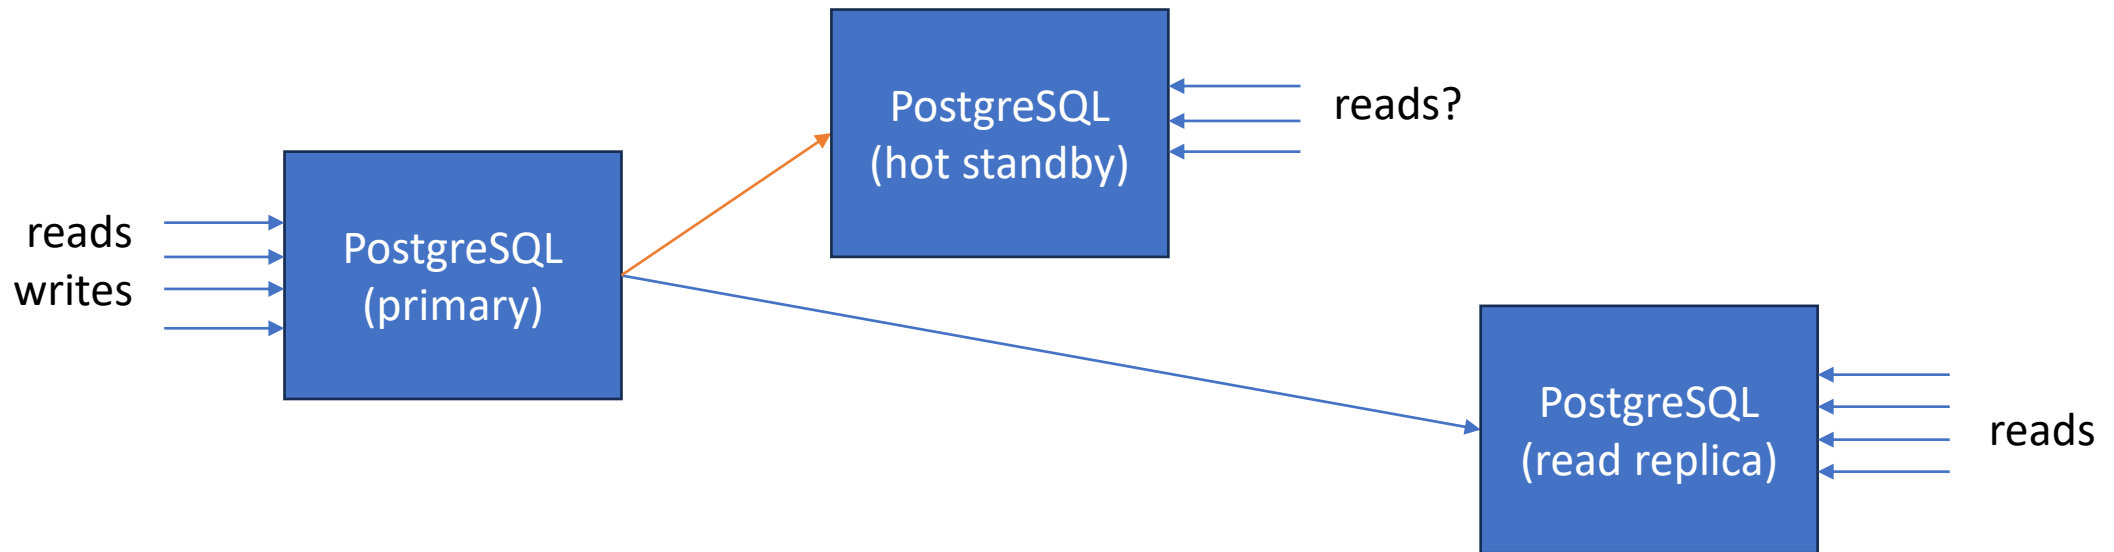
Poor cache usage

Hot standby




# Read replica vs. hot standby

“Hot standby” refers to a read replica that is a candidate for auto-failover, can also serve reads.



# Hot standby write latency

PostgreSQL first writes to primary WAL and then waits for replication to standby WAL.




**Batching writes is helpful!**

**Use COPY, multi-row inserts, etc.**

`synchronous_commit = on`

# Hot standby write latency

PostgreSQL first writes to primary WAL and then waits for replication to standby WAL.




`synchronous_commit = remote_write`

Some data lost only when simultaneous failure of primary and hot standby.

`synchronous_commit = remote_write`

# Hot standby write latency

PostgreSQL first writes to primary WAL and then waits for replication to standby WAL.




`synchronous_commit = local`

Some data lost when failover is triggered.

`synchronous_commit = local`

# Hot standby write latency

PostgreSQL first writes to primary WAL and then waits for replication to standby WAL.



`synchronous_commit = off`


Some data lost when primary fails.

`synchronous_commit = local`



# Primary failure

When a primary fails from the cluster manager point-of-view, it initiates failover.



**In a few seconds to minutes:**


- 1) Detect failure**
- 2) Demote primary (?)**
- 3) Promote standby**
- 4) Updating routing mechanism:**
  - Load balancer / VIP (requires infra to be up)
  - DNS (requires infra to be up, wait for TTLS)
  - Pooler (can fail, be stale)
  - Client-side (reliable, but clumsy)

**In a few minutes:**

- 5) Replace or rewind old primary**

# Primary network partition

Primary might not actually be down...



In a few seconds to minutes:

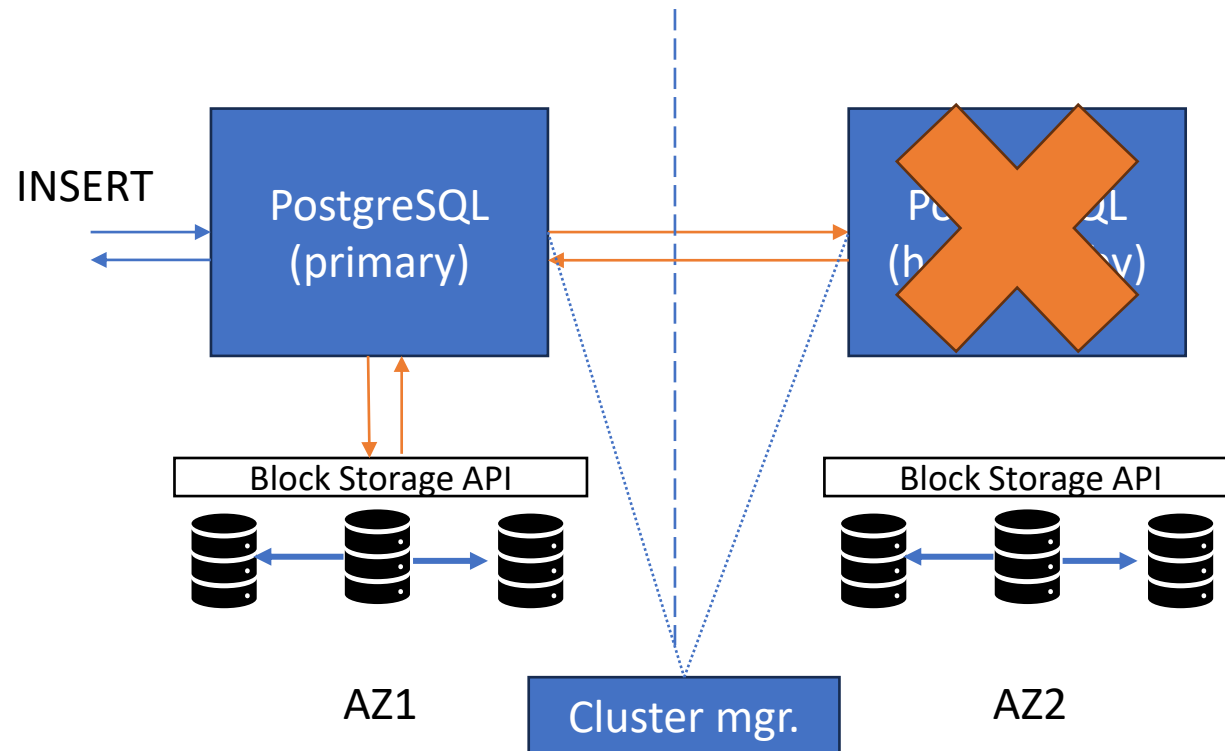
- 1) Detect failure
- 2) Demote primary (?)
- 3) Promote standby
- 4) Updating routing mechanism:
  - Load balancer / VIP (requires infra to be up)
  - DNS (requires infra to be up, wait for TTLS)
  - Pooler (can fail, be stale)
  - Client-side (reliable, but clumsy)

In a few minutes:

- 5) Replace or rewind old primary

# Hot standby failure

Writes get blocked when hot standby fails, but control plane can quickly react and make the standby ineligible.



In a few seconds:

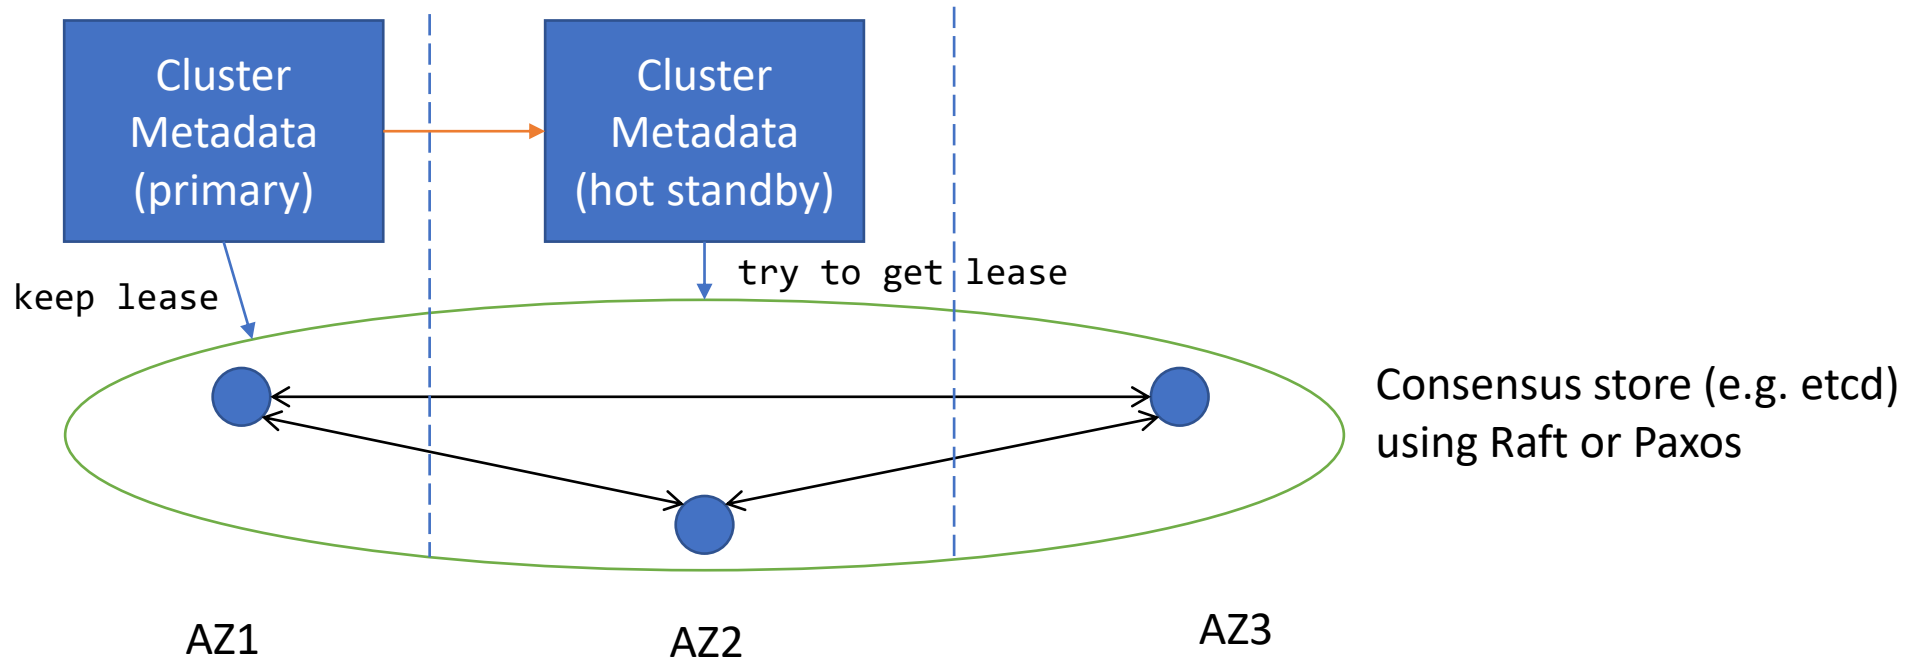
- 1) Pause failover (back online)
- 2) Reduce synchronous commit

In a few minutes:

- 2) Create new standby or wait for recovery
- 3) Resume failover

# Cluster manager

Cluster managers use consensus to decide a primary metadata store.



# Hot standby trade-offs

## **Pros:**

Can survive node and availability zone failure

No crash recovery when resuming from failure

## **Cons:**

Write latency can be very high

Cost is twice as high

Rerouting on failover can have issues / take a while

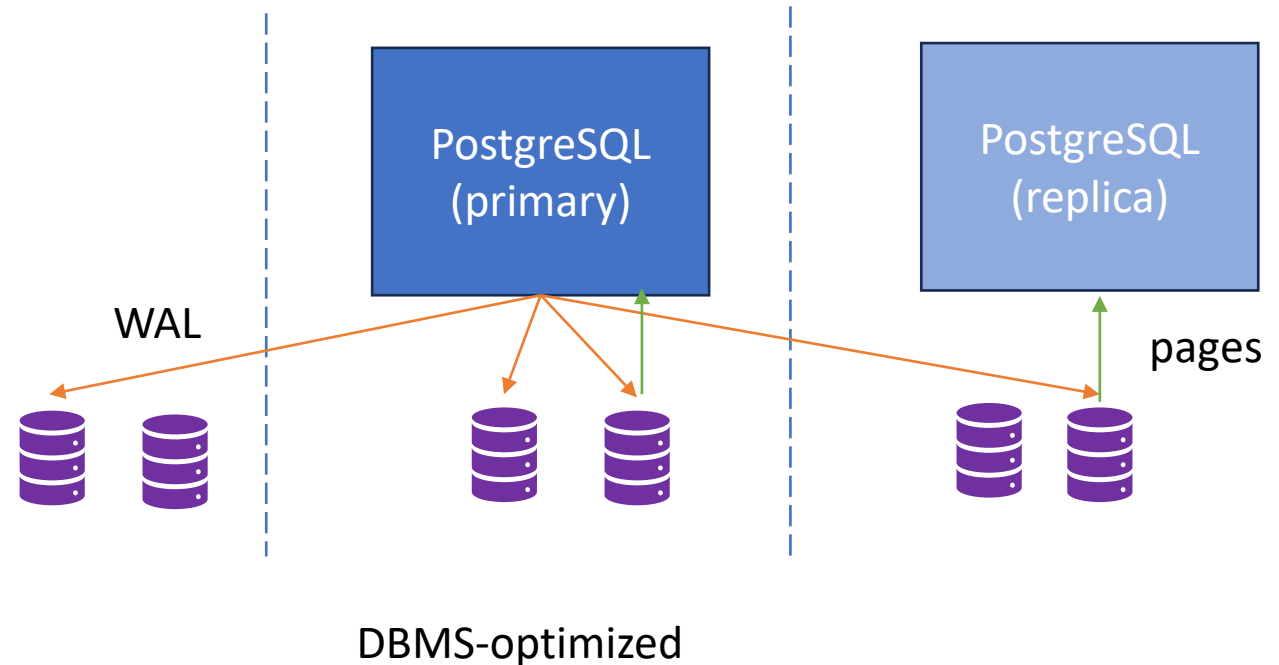

# DBMS-optimized storage

Like Aurora, Neon, AlloyDB



# DBMS-optimized storage

Cloud storage that can perform background page writes autonomously, which saves on write I/O from primary. Also optimized for other DBMS needs (e.g. read replicas).





# DBMS-optimized storage trade-offs

## **Pros:**

Potential performance benefits by avoiding page writes from primary

No crash recovery

Replicas can reuse storage, incl. hot standby

Less rigid than network-attached storage implementations (faster reattach, branching, ...)

## **Cons:**


Write latency is high by default

High cost / pricing

PostgreSQL is not designed for it, can be slower than regular storage

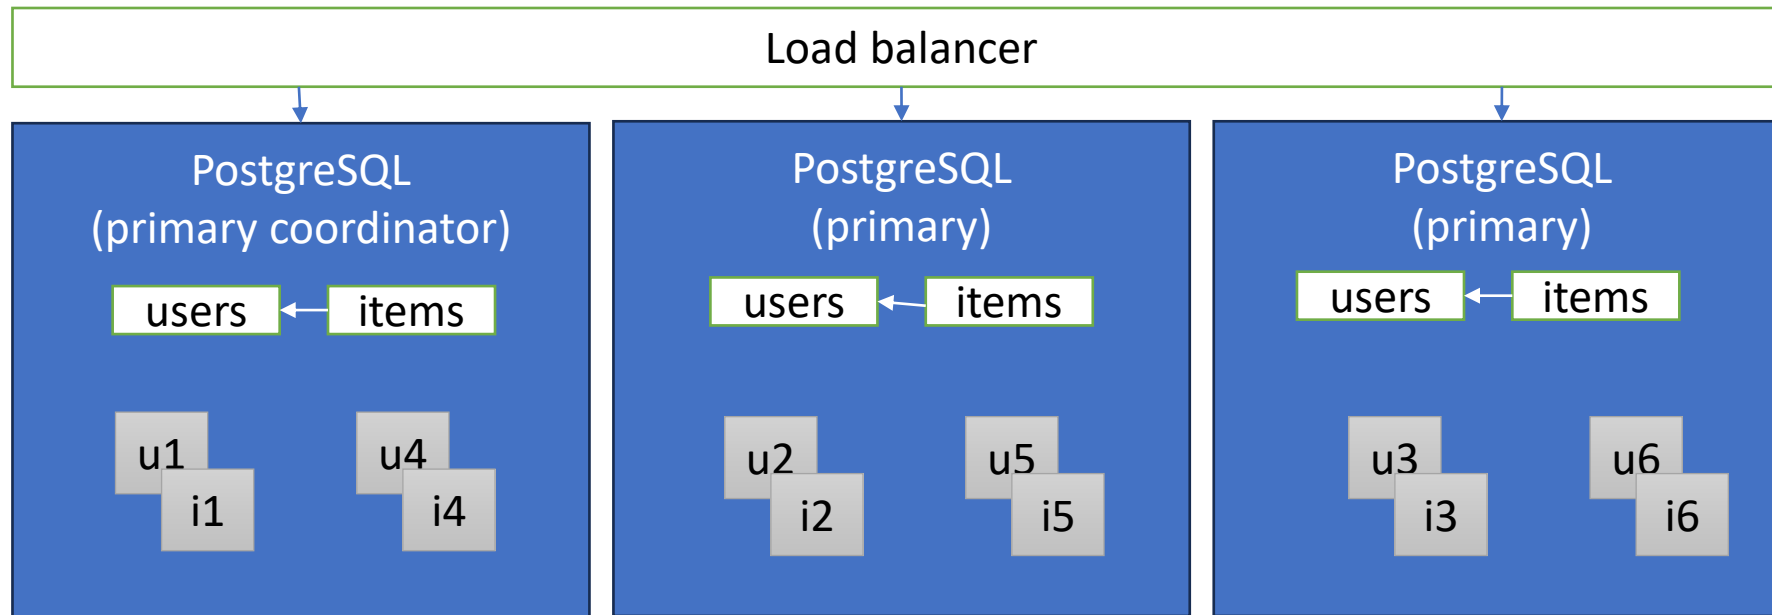
# Transparent sharding

Like Citus



# Transparent sharding

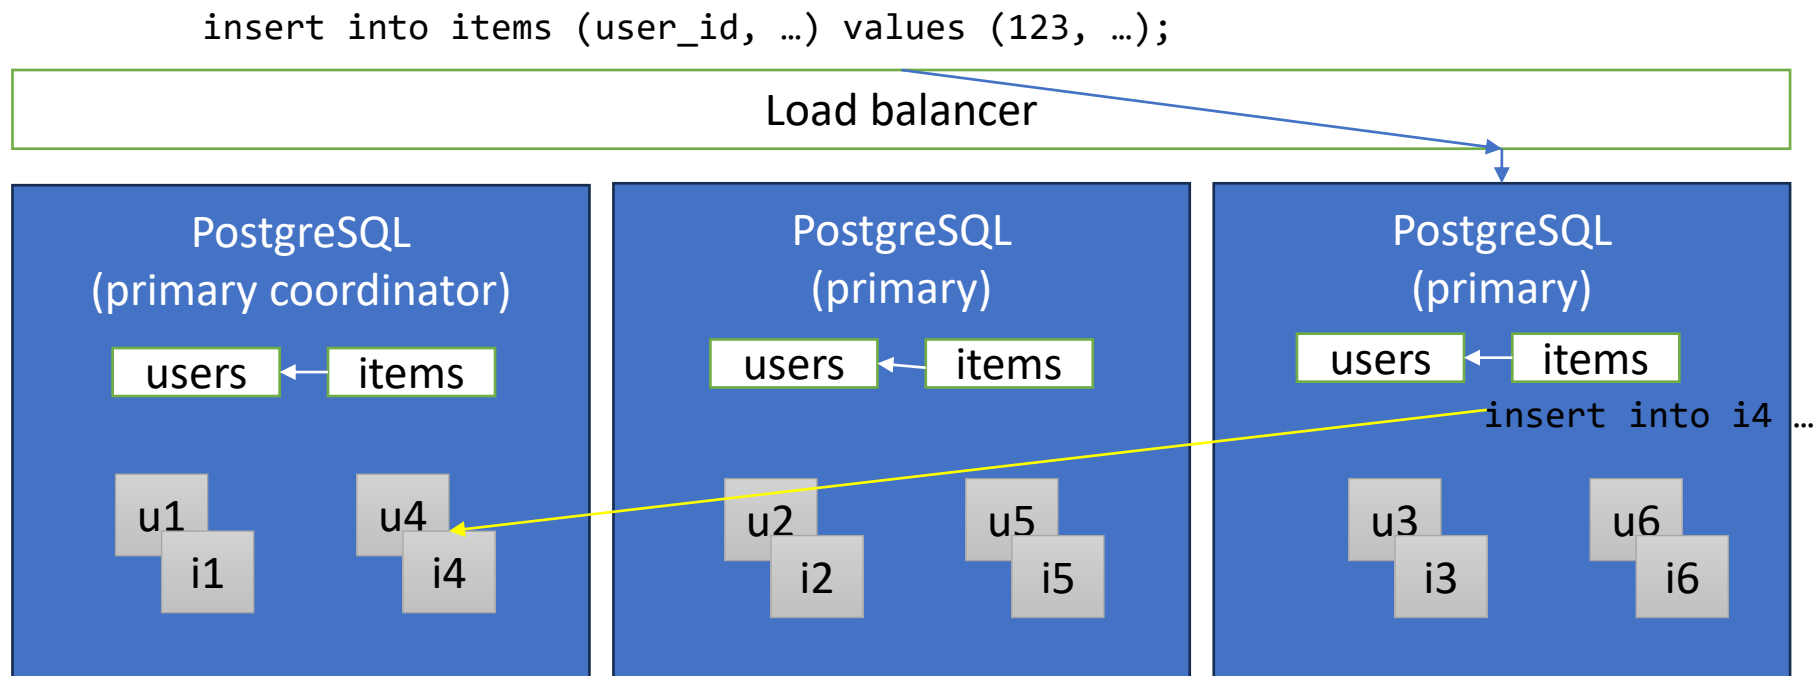
Distribute tables by a shard key and/or replicate tables across multiple (primary) nodes. Queries & transactions are transparently routed / parallelized.



Tables can be co-located to enable local joins, foreign keys, etc. by the shard key.

# Single shard queries for operational workloads

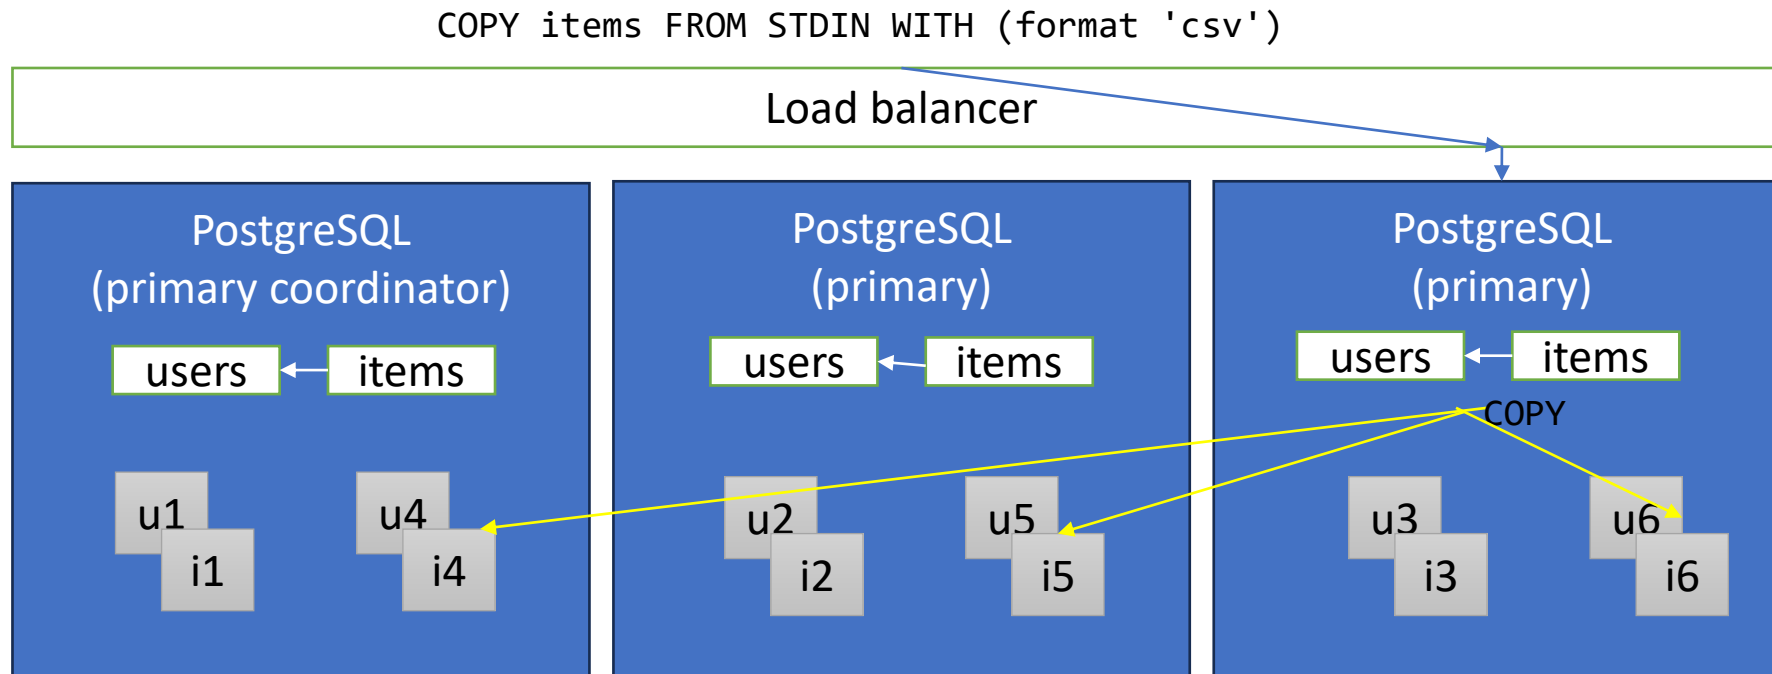
Scale capacity for handling a high rate of single shard key queries:



Per-statement latency can be a bottleneck!

# Data loading in sharded system

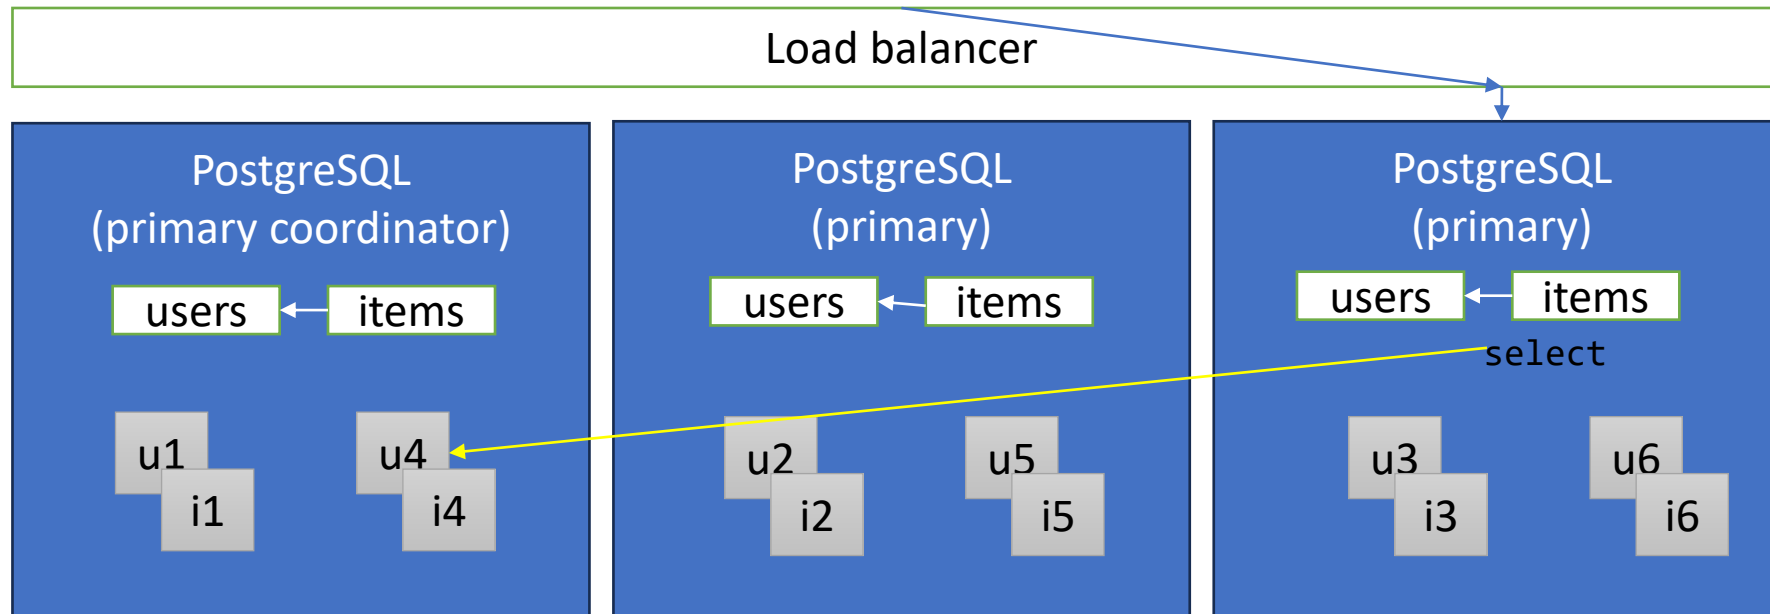
Pipelining through COPY can make data loading a lot more efficient and scalable



# Compute-heavy queries

Compute-heavy queries (shard key joins, json, vector, ...) get the most relative benefit

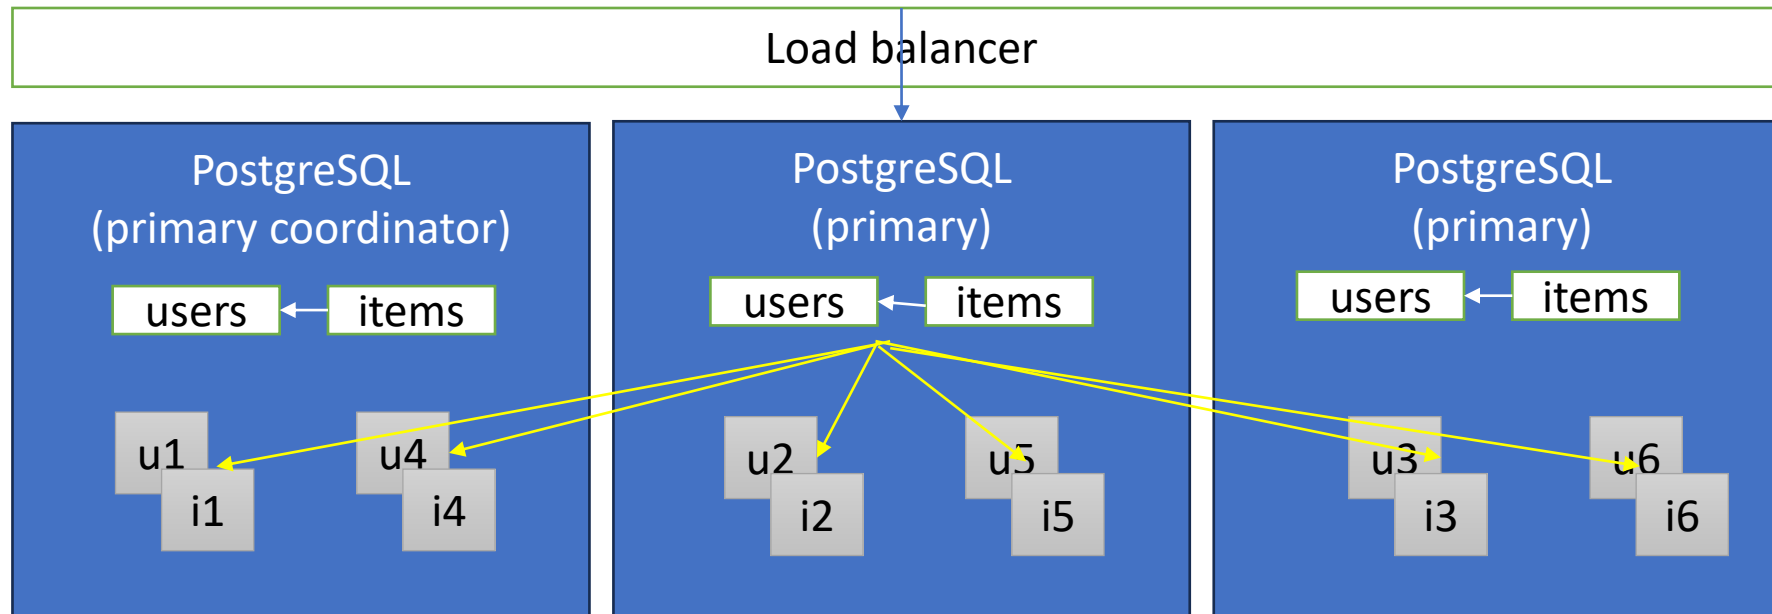
```
select compute_stuff(...) from users join items using (user_id) where user_id = 123 ...
```



# Multi-shard queries for analytical workloads

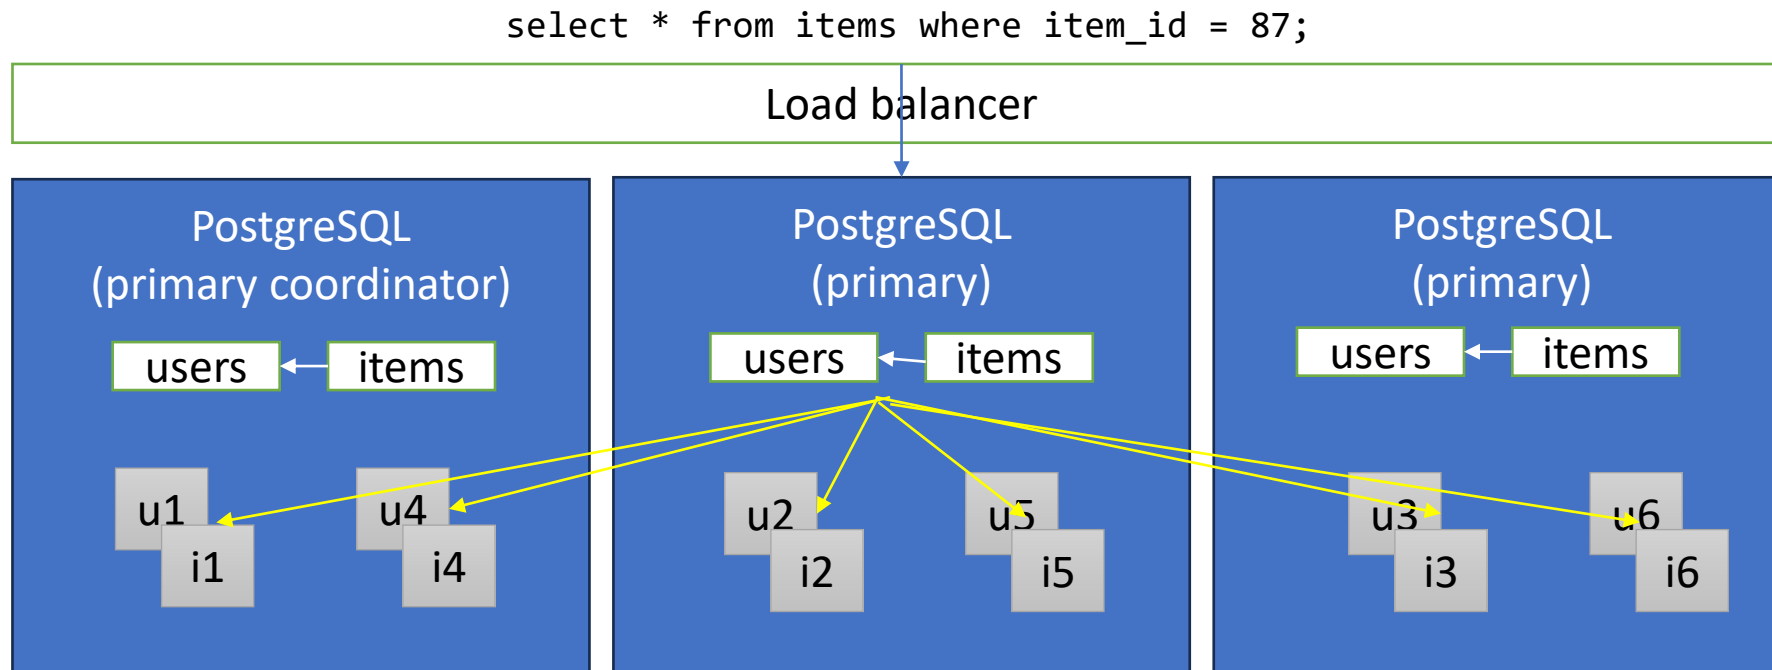
Parallel multi-shard queries can quickly answer analytical queries across shard keys:

```
select country, count(*) from items, users where ... group by 1 order by 2 desc limit 10;
```



# Multi-shard queries for operational workloads

Multi-shard queries add significant overhead for simple non-shard-key queries

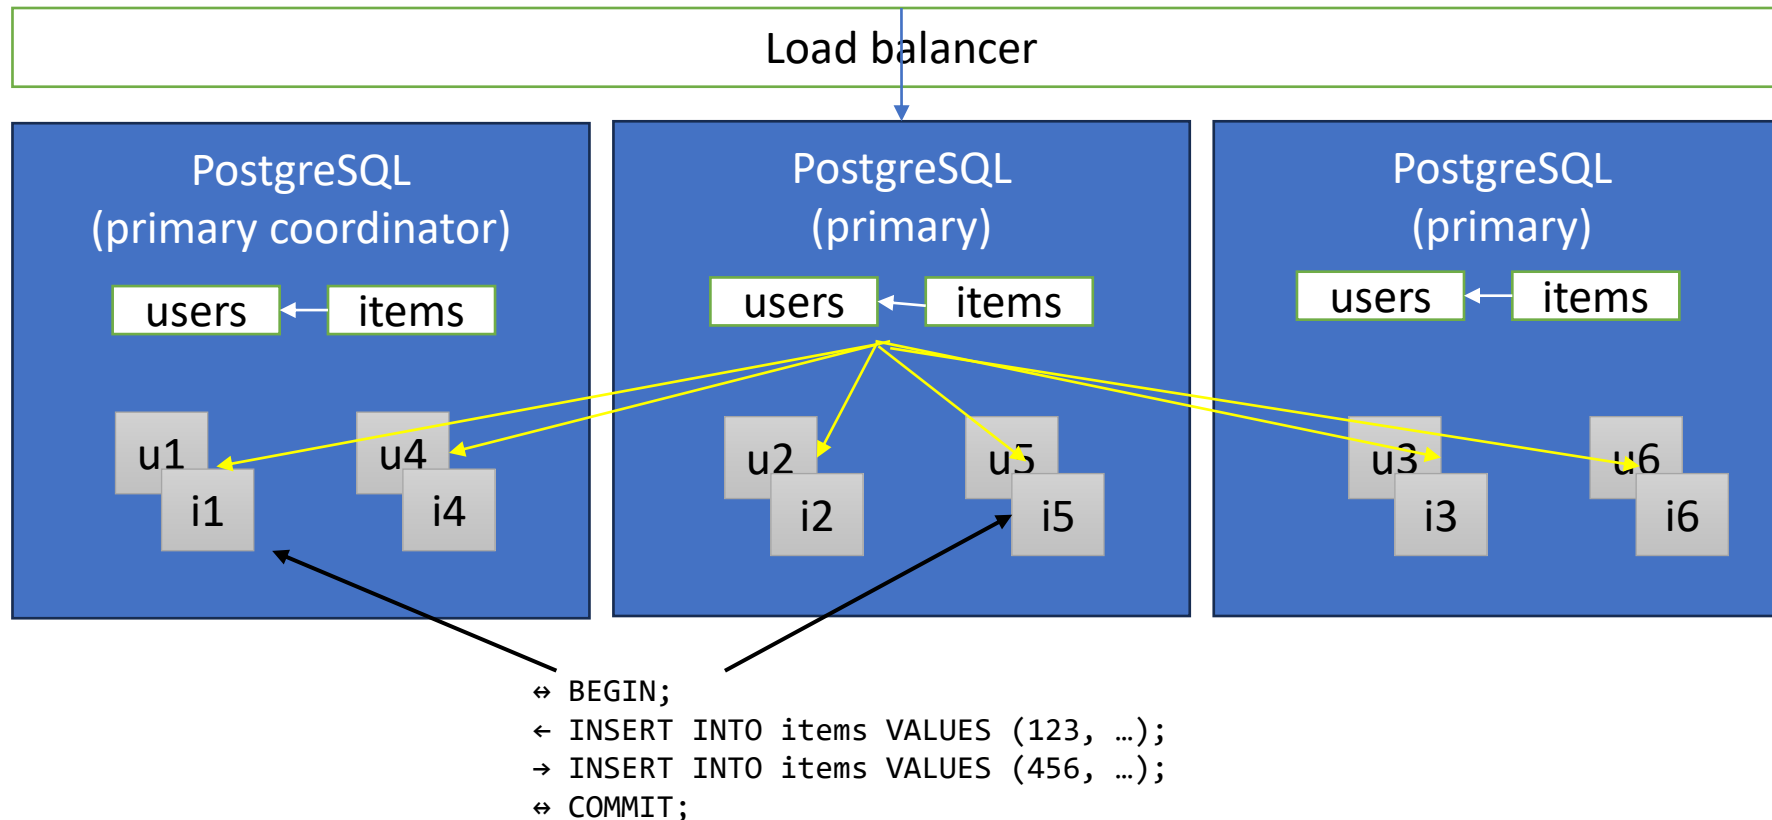




# Multi-shard queries for analytical workloads

Snapshot isolation is a challenge (involves trade-offs):

```
select country, count(*) from items, users where ... group by 1 order by 2 desc limit 10;
```



# Sharding trade-offs

## **Pros:**

Scale throughput for reads & writes (CPU & IOPS)

Scale memory for large working sets

Parallelize analytical queries, batch operations

## **Cons:**


High read and write latency

Data model decisions have high impact on performance

Snapshot isolation concessions

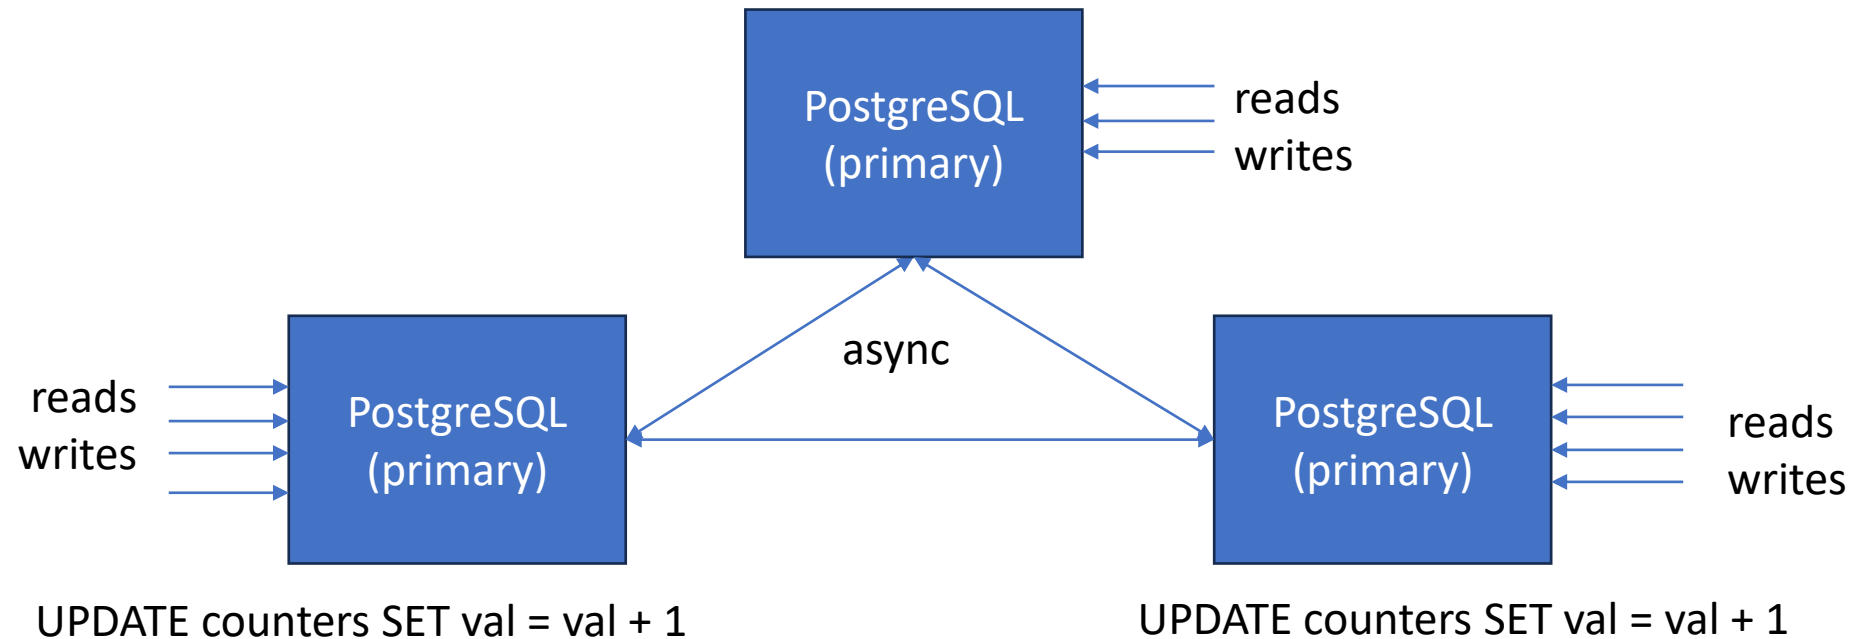
# Active-active

Like BDR, pgactive, pgEdge, ...




# Active-active / n-way replication

Accept writes from any node, use logical replication to asynchronously exchange and consolidate writes.



# Active-active / n-way replication

All nodes can survive network partitions by accepting writes locally, but no linear history (CAP).



# Active-active trade-offs

## **Pros:**

Very high read and write availability

Low read and write latency

*Read* throughput scales linearly

## **Cons:**


Eventual read-your-writes consistency

No monotonic read consistency

No linear history (updates might conflict after commit)

# Distributed SQL

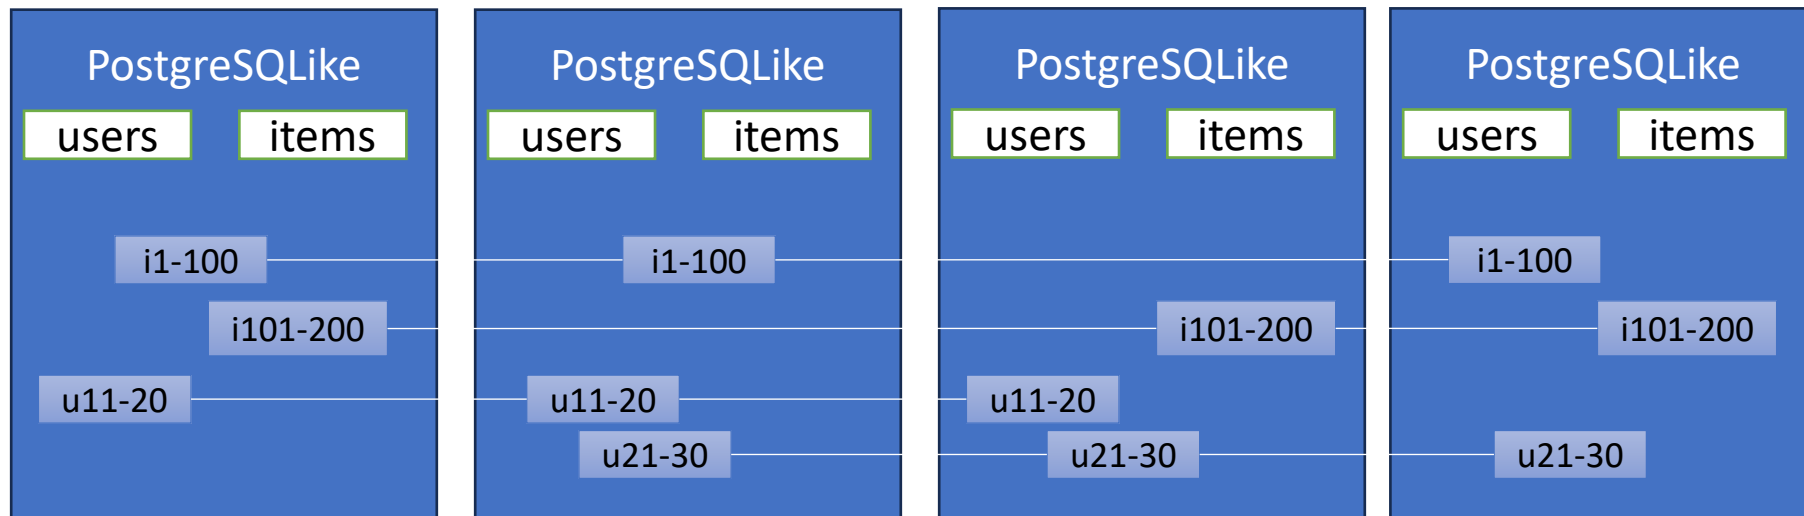
Like Yugabyte, CockroachDB, Spanner



# Distributed key-value storage with SQL (DSQL)

Tables are stored on distributed key-value stores, shards replicated using Paxos/Raft.

Distributed transactions with snapshot isolation via global timestamps (HLC or TrueTime).





# Distributed key-value storage trade-offs

## **Pros:**

Good read and write availability (shard-level failover)

Single table, single key operations scale well

No additional data modelling steps or snapshot isolation concessions

## **Cons:**

Many operations (index lookups, evaluating joins and foreign keys) incur high latency

No local joins (co-location, reference tables) in current implementations

Less mature and optimized than PostgreSQL

# Conclusion

PostgreSQL can be distributed at different layers.

Each architecture can introduce severe trade-offs.

Almost nothing comes for free..

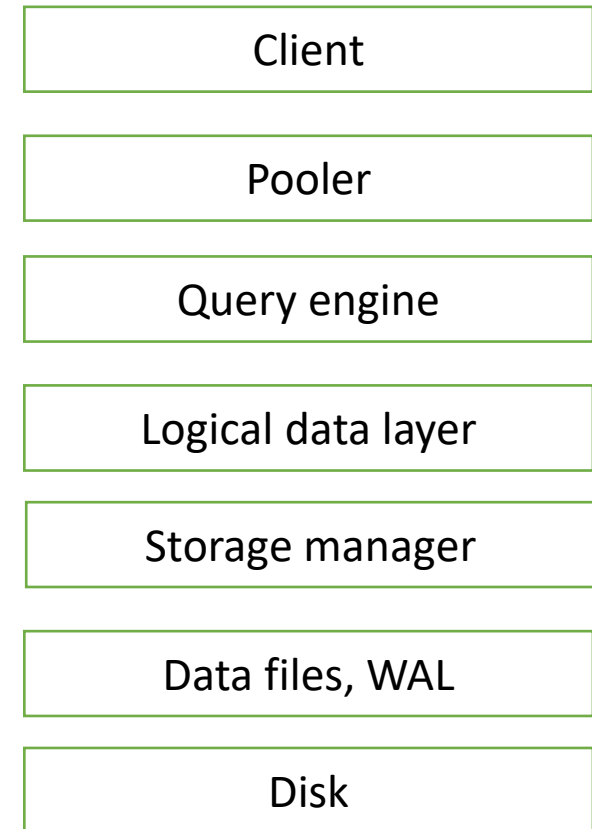
Keep asking:

What do I really want?

Which architecture achieves that?

What are the trade-offs?

What can my application tolerate? (can I change it?)



# General guidelines

- **Network-attached storage** operational benefits usually outweigh the perf. downsides
- **Read replicas** scale reads linearly, but anomalies can be cumbersome in practice
- **Hot standby** is good for availability, but costly and need to batch your writes
- **DBMS-optimized storage** makes subtle trade-offs, mileage may vary
- **Sharding** makes crude trade-offs to achieve scalability, requires data model shift
- **Active-active** anomalies usually outweigh its benefits
- **DSQL** works well as a versatile key-value store, but just use PostgreSQL

# Questions?

Marco.slot@gmail.com