

# Data Warehouse Systems: Design and Implementation

## Second Edition

**Alejandro VAISMAN**

Department of Information Engineering  
Instituto Tecnológico de Buenos Aires  
[avaisman@itba.edu.ar](mailto:avaisman@itba.edu.ar)

**Esteban ZIMÁNYI**

Department of Computer & Decision Engineering (CoDE)  
Université libre de Bruxelles  
[esteban.zimanyi@ulb.be](mailto:esteban.zimanyi@ulb.be)

## Chapter 4: Conceptual Data Warehouse Design

1. MultiDim: A Conceptual Model for Data Warehouses
2. Dimension Hierarchies
3. Advanced Modeling Aspects
4. Querying the Northwind Cube

## 1. MultiDim: A Conceptual Model for Data Warehouses

## 2. Dimension Hierarchies

## 3. Advanced Modeling Aspects

## 4. Querying the Northwind Cube

# Conceptual Multidimensional Models

## ◆ Conceptual models


- Allow better communication between designers and users to understand application requirements
  - More stable than implementation-oriented (logical) schema, which changes with the platform
  - Provide better support for visual user interfaces
- ◆ No well-established conceptual model for multidimensional data
- ◆ Several proposals based on UML, on the ER model, or using specific notations
- ◆ Problems:
- Cannot express complex kinds of hierarchies
  - Lack of a mapping to the implementation platform
- ◆ Currently, data warehouses are designed using mostly logical models (star and snowflake schemas)
- Difficult to express requirements (technical knowledge required)
  - Limit users to defining only elements that the underlying implementation systems can manage
  - Example: Users constrained to use only the simple hierarchies supported in current tools

# MultiDim: A Conceptual Multidimensional Model


- ◆ Based on the entity-relationship model
- ◆ Includes concepts like:
  - **dimensions**
  - **hierarchies**
  - **facts**
  - **measures**
- ◆ Supports various kinds of hierarchies existing in real-world applications
- ◆ Can be mapped to star or snowflake relational structures

# MultiDim Model: Notation

- ◆ **Dimension**: level or one or more hierarchies
- ◆ **Hierarchy**: several related levels
- ◆ **Level**: entity type
- ◆ **Member**: every instance of a level
- ◆ **Child** and **parent** levels: the lower and higher levels
- ◆ **Leaf** and **root** levels: first and last levels in a hierarchy
- ◆ **Cardinality**: Minimum/maximum numbers of members in a level related to members in another level
- ◆ **Criterion**: Expresses different hierarchical structures used for analysis
- ◆ **Key attribute**: Indicates how child members are grouped
- ◆ **Descriptive attributes**: Describe characteristics of members

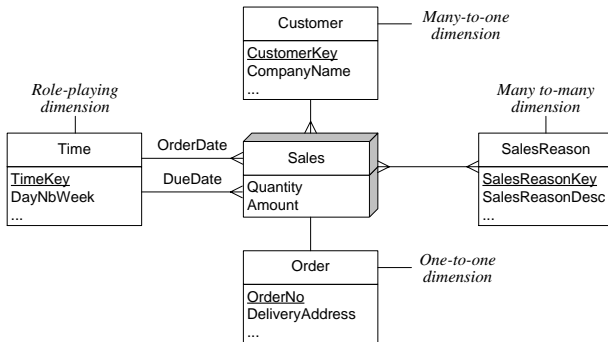


## Cardinalities




# MultiDim Model: Notation


- ◆ **Fact**: Relates measures to leaf levels in dimensions
- ◆ Dimensions can be related to fact with **one-to-one**, **one-to-many**, of **many-to-many**
- ◆ Dimension can be related several times to a fact with **different roles**





# MultiDim Model: Notation (Summary)



## Fact with measures and associated levels

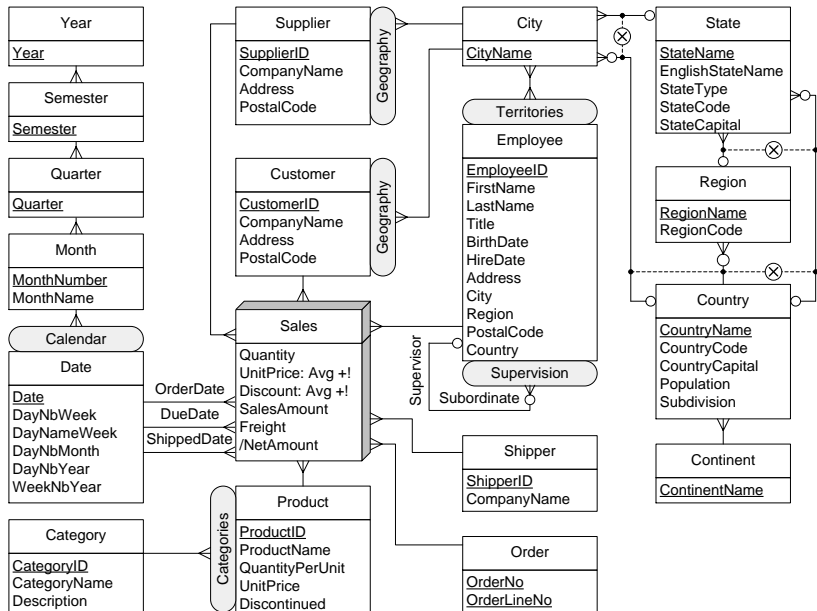


## Types of measures   Analysis criterion   Distributing factor   Exclusive relationships





# MultiDim Conceptual Schema of the Northwind Data Warehouse



# Table of Contents

1. MultiDim: A Conceptual Model for Data Warehouses

**2. Dimension Hierarchies**

3. Advanced Modeling Aspects


4. Querying the Northwind Cube

# Dimension Hierarchies


- ◆ Crucial in analytical applications
- ◆ Enable analysis at various abstraction levels
- ◆ In real-world situations, users must deal with complex hierarchies of various kinds
- ◆ Logical models of current DW and OLAP systems allow only a limited set of kinds of hierarchies
  - Users unable to capture the essential semantics of multidimensional applications
  - They must limit their analysis to the predefined set of hierarchies supported by the tools
- ◆ At the conceptual level, focus is to establish sequences of levels that should be traversed during roll-up and drill-down
- ◆ Distinction between the various kinds of hierarchies should also be made at the instance level
- ◆ Cardinalities in parent-child relationships must be considered
- ◆ MultiDim includes classification of hierarchies at the schema and instance level and proposes a graphical notation

# Balanced Hierarchies

- ◆ At **schema level**: only one path where all parent-child relationships are many-to-one and mandatory




- ◆ At **instance level**: members form a balanced tree (all the branches have the same length)
- ◆ All parent members have at least one child member, and a child belongs exactly to one parent




# Unbalanced Hierarchies

- At **schema level**: one path where all parent-child relationships are many-to-one, but some are optional




- At **instance level**: members form a unbalanced tree




# Recursive Hierarchies

- ◆ A special case of unbalanced hierarchies
- ◆ The **same level** is linked by the two roles of a parent-child relationship
- ◆ Used when all hierarchy levels express the same semantics
- ◆ The characteristics of the parent and child are similar (or the same)
- ◆ **Schema level**

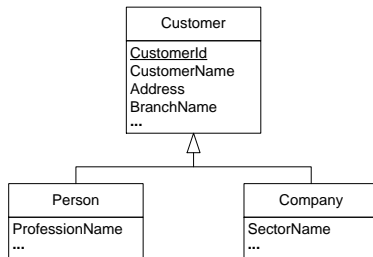


- ◆ **Instance level**



# Generalized Hierarchies


- ◆ Sometimes the members of a hierarchy are of different type
- ◆ Entity-relationship representation of customer types




- ◆ Hierarchical paths cannot be clearly distinguished
- ◆ Higher hierarchy levels (e.g., Branch) in a supertype not related to other hierarchy levels in its subtypes
- ◆ Not clear that measures related to a customer that is a person can be aggregated using a hierarchy Customer–Profession–Branch

# Generalized Hierarchies

- ◆ At **schema level**: multiple exclusive paths sharing at least the leaf level; may also share other levels
- ◆ Two aggregation paths, one for each type of customer



- ◆ At **instance level**: each member belongs to only one path



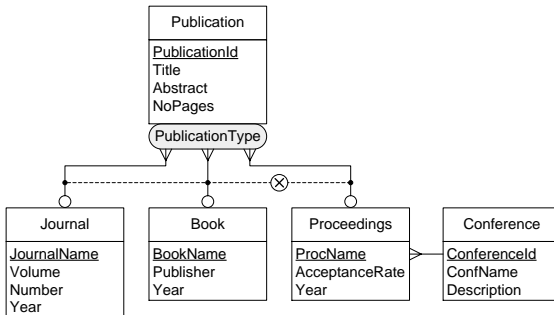


# Generalized Hierarchies

- ◆ **Supertype** of the generalization/specialization relationship is used in generalized hierarchies for representing a **leaf level**
- ◆ It only includes those attributes that represent concepts at the lowest granularity
  - E.g., CustomerId, CustomerName, and Address
- ◆ This kind of hierarchy **does not satisfy the summarizability conditions**
  - The mapping from the splitting level to the parent levels is incomplete
    - ◆ E.g., not all customers roll up to the Sector level
    - ◆ E.g., not all customers are mapped to the Profession level
- ◆ Conventional aggregation mechanism should be modified when a splitting and joining levels are reached in a drill-down and roll-up operations
- ◆ Traditional approach can be used for aggregating measures for common hierarchy levels

# Generalized Hierarchies

- ◆ In generalized hierarchies, it is not necessary that splitting levels must be joined




# Generalized Hierarchies

- ◆ Not all generalization/specialization hierarchies can be represented
- ◆ **Partial specializations**: Induce an additional path in the generalized hierarchy, relating the common levels
- ◆ **Overlapping specializations**: Various options are possible according to the users' requirements and the availability of measures
  - Example: An overlapping generalization where a person who owns a company buys products either for his/her individual use or for the company
  - If measures are known only for the superclass Customer, only the hierarchy with common levels will be represented, e.g., the Customer and Area levels
  - If measures are known only for each subclass, e.g., for Person and Company:
    - ◆ Separate dimensions and fact relationships with corresponding measures can be created for each specialization → difficult to manage dimensions with overlapping sets of members
    - ◆ Another solution: Disallow overlapping generalizations

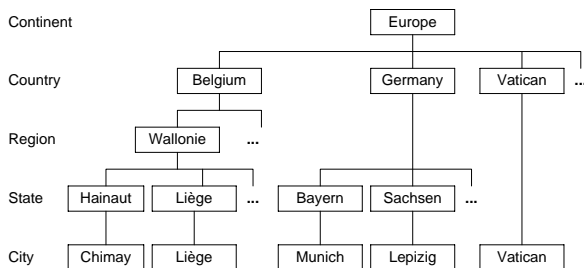
# Noncovering Hierarchies

- ◆ Also known as **ragged** or **level-skipping hierarchies**
- ◆ A **special case of generalized hierarchies**
- ◆ At the **schema level**: Alternative paths are obtained by skipping one or several intermediate levels




# Noncovering Hierarchies

- ◆ At **instance level**: Path length from the leaves to the same parent can be different for different members

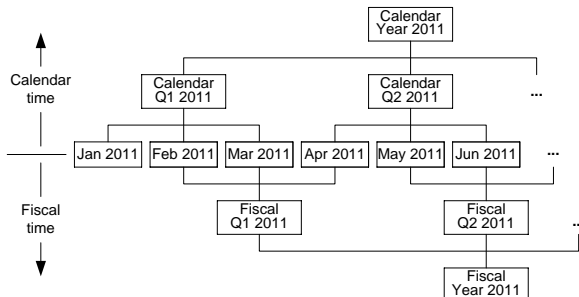


# Alternative Hierarchies

- ◆ At **schema level**: Multiple nonexclusive hierarchies that share at least the leaf level and account for the same analysis criterion



- ◆ At **instance level**: Members form graph

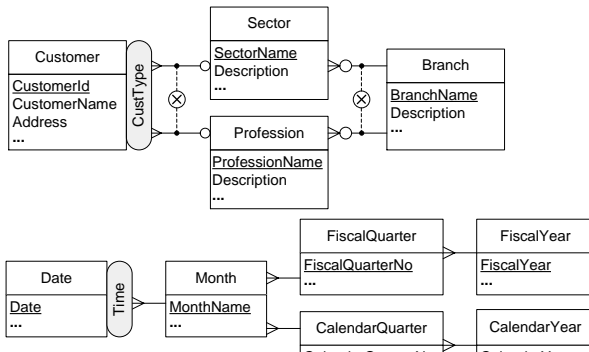


## Alternative Hierarchies

- ◆ Needed to analyze measures from an unique perspective (e.g., time) using alternative paths
- ◆ Measures will participate totally in each component hierarchy  $\Rightarrow$  conventional aggregation procedures
- ◆ It is not semantically correct to simultaneously combine different component hierarchies
- ◆ Combination can give meaningless intersections, i.e., a combination of members that do not have values for aggregated measures, e.g., B1-2001 and Q2-2001
- ◆ Users must choose only one of the alternative paths for their analysis and switch to other one if required

# Generalized vs. Alternative Hierarchies


- ◆ Both hierarchies
  - Share some levels
  - Use one analysis criterion
- ◆ A child member
  - Related to only one path in generalized hierarchies
  - Related to all paths in alternative hierarchies and users must choose one for analysis





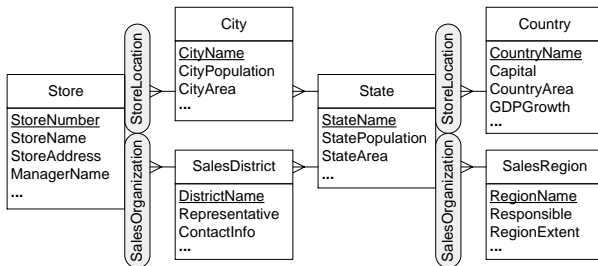
# Parallel Hierarchies

- ◆ Dimension has associated **several hierarchies** accounting for **different analysis criteria**
- ◆ Two different types
  - Parallel **independent** hierarchies
  - Parallel **dependent** hierarchies
- ◆ Parallel **independent** hierarchies
  - Composed of disjoint hierarchies, i.e., hierarchies that **do not share levels**
  - Component hierarchies may be of different kinds




# Parallel Hierarchies

- ◆ Parallel **dependent** hierarchies
- ◆ Composed of several hierarchies that account for different analysis criteria and **share some levels**
- ◆ Component hierarchies may be of different kinds



## Parallel Hierarchies


- ◆ Parallel dependent hierarchies leading to different parent members of the shared level



# Alternative vs. Parallel Hierarchies


- ◆ Both hierarchies
  - Share some levels
  - May include several simple hierarchies
- ◆ Criterion
  - Only one for alternative hierarchies
  - Several for parallel hierarchies
- ◆ Combining hierarchies
  - Meaningless for alternative hierarchies
  - Useful for parallel hierarchies
- ◆ Reusing aggregated measures for common levels
  - Can be done for alternative hierarchies
  - Cannot be done for parallel hierarchies

## Alternative vs. Parallel Hierarchies



- ◆ Aggregated measure for the Month level can be reused between both paths
- ◆ Traversing the Calendar hierarchy from a specific day in the Date level will end up in the same year independently of which path is used


## Alternative vs. Parallel Hierarchies




- ◆ Aggregated measure for State level cannot be reused between both paths
- ◆ Traversing the hierarchies Live and Work from the Employee to the State level will lead to different states for employees who live in one state and work in another

# Nonstrict Hierarchies

- ◆ At **schema level**: At least one many-to-many cardinality



- ◆ At **instance level**: Members form a graph



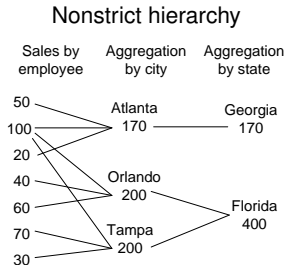

## ◆ **Nonstrict Hierarchies**

- ◆ Abuse of terminology: “nonstrict hierarchy” for “acyclic classification graph”
- ◆ The term “hierarchy” conveys the notion that users need to analyze measures at different levels of detail, which is less clear with the term “acyclic classification graph”
- ◆ The term “hierarchy” is already used by practitioners and some tools, in particular SQL Server Analysis Services, allow many-to-many parent-child relationships
- ◆ The term “hierarchy” is also used by several researchers



# Nonstrict Hierarchies: Double Counting


- ◆ Problem: **Double counting** of measures when a roll-up operation reaches a many-to-many relationship
- ◆ Examples of aggregation



# Nonstrict Hierarchies: Solutions for Double Counting


- ◆ Include a **distributing factor**
- ◆ Calculate **approximate** values of a distributing factor
- ◆ **Transform** a nonstrict hierarchy into a strict one:
  - **Create a new parent member** for each group of parent members linked to a single child member in a many-to-many relationship
  - **Choose one parent member as primary** and ignore the existence of other parent members
  - **Split the hierarchy in two** at the many-to-many relationship, where the levels from the parent level and beyond become a new dimension
- ◆ Each solution has its advantages and disadvantages and requires special aggregation procedures
- ◆ Appropriate solution must be chosen according to the situation at hand and user's requirements

## Nonstrict Hierarchies: Distributing Factor




- ◆ Employees may work in several sections
- ◆ A measure represents an employee's overall salary, i.e., the sum of the salaries paid in each section
- ◆ Distributing factor determines how measures are divided between several parent members
- ◆ Distributing factor is **not always known**
  - Percentage of time that an employee works in a section must be added to schema
- ◆ Sometimes this distribution is **impossible to specify**
  - E.g., participation of customer in joint account
- ◆ Distributing factor can be **approximated** by considering the total number of parent members with which the child member is associated
  - If an employee works in three sections, 1/3 of the value of the measure aggregated for each one

# Nonstrict Hierarchies: Splitting the Hierarchy



- ◆ Transform a nonstrict hierarchy into a strict one with an **additional dimension**
- ◆ **Focus of analysis** has changed from employee's salaries to employee's salaries by section
- ◆ Can only be applied when the **measure distribution is known**
- ◆ Nevertheless, double counting problem still remains
- ◆ Example: calculate the number of employees by section or by division



# Table of Contents

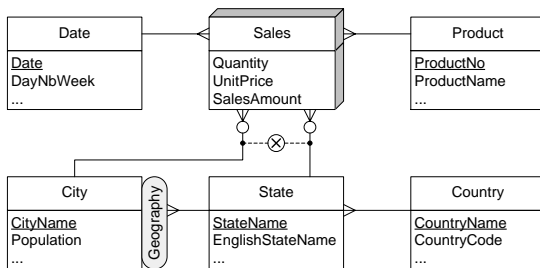
1. MultiDim: A Conceptual Model for Data Warehouses

2. Dimension Hierarchies

**3. Advanced Modeling Aspects**


4. Querying the Northwind Cube

## Advanced Modeling Aspects: Facts with Multiple Granularities



- ◆ Sales captured at the city level or at the state level

## Advanced Modeling Aspects: Many-to-Many Dimensions



- ◆ Multidimensional schema for the analysis of bank accounts

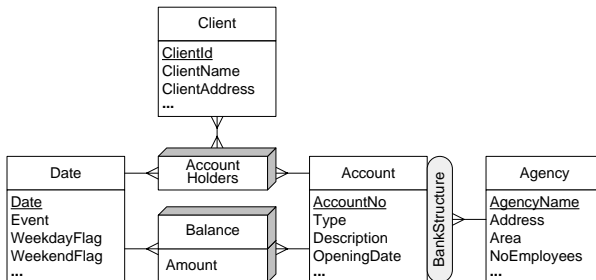
◆ Example of double-counting problem

Time	Account	Client	Balance
T1	A1	C1	100
T1	A1	C2	100
T1	A1	C3	100
T1	A2	C1	500
T1	A2	C2	500



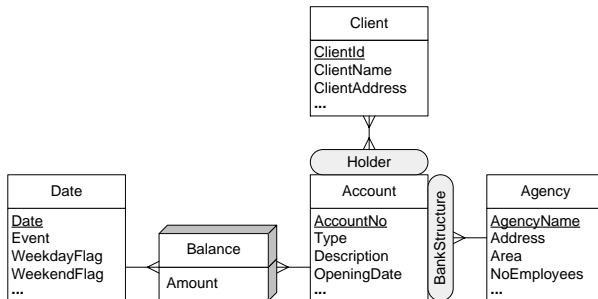
# Advanced Modeling Aspects: Many-to-Many Dimensions

- ◆ Two possible decompositions of the fact
  1. Creating two facts



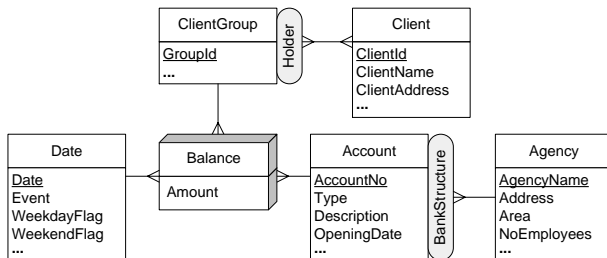
# Advanced Modeling Aspects: Many-to-Many Dimensions

- ◆ Two possible decompositions of the fact
  - 2. Including a nonstrict hierarchy



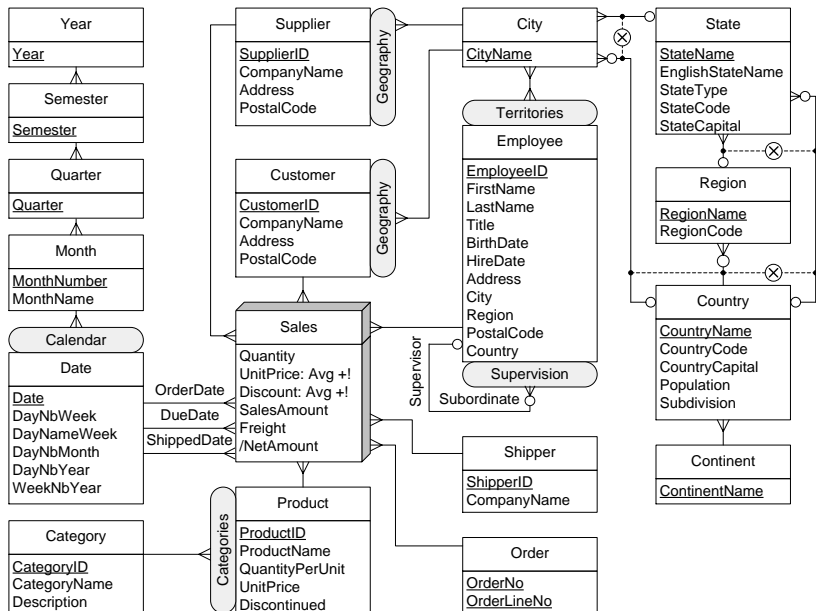
## Advanced Modeling Aspects: Many-to-Many Dimensions

### ◆ Alternative decomposition of the schema



1. MultiDim: A Conceptual Model for Data Warehouses
2. Dimension Hierarchies
3. Advanced Modeling Aspects
4. Querying the Northwind Cube

# Conceptual Schema of the Northwind Cube



## Querying the Northwind Cube

- ◆ **Query 4.1:** *Total sales amount per customer, year, and product category*

ROLLUP\*(Sales, Customer → Customer, OrderDate → Year, Product → Category,  
SUM(SalesAmount))

- ◆ **Query 4.2:** *Yearly sales amount per pair customer country and supplier country*

ROLLUP\*(Sales, OrderDate → Year, Customer → Country, Supplier → Country,  
SUM(SalesAmount))

- ◆ **Query 4.3:** *Monthly sales by customer state compared to those of the previous year*

Sales1 ← ROLLUP\*(Sales, OrderDate → Month, Customer → State, SUM(SalesAmount))  
Sales2 ← RENAME(Sales1, SalesAmount ← PrevYearSalesAmount)  
Result ← DRILLACROSS(Sales2, Sales1,  
Sales2.OrderDate.Month = Sales1.OrderDate.Month AND  
Sales2.OrderDate.Year+1 = Sales1.OrderDate.Year AND  
Sales2.Customer.State = Sales1.Customer.State)

- ◆ **Query 4.4:** *Total sales growth per month per product, that is, total sales per product compared to the previous month*

Sales1 ← ROLLUP\*(Sales, OrderDate → Month, Product → Product, SUM(SalesAmount))

Sales2 ← RENAME(Sales1, SalesAmount ← PrevMonthSalesAmount)

Sales3 ← DRILLACROSS(Sales2, Sales1,  
    ( Sales2.OrderDate.Month > 1 AND  
      Sales2.OrderDate.Month+1 = Sales1.OrderDate.Month AND  
      Sales2.OrderDate.Year = Sales1.OrderDate.Year ) OR  
    ( Sales2.OrderDate.Month = 1 AND  
      Sales1.OrderDate.Month = 12 AND  
      Sales2.OrderDate.Year+1 = Sales1.OrderDate.Year )

Result ← ADDMEASURE(Sales3, SalesGrowth = SalesAmount - PrevMonthSalesAmount)

# Querying the Northwind Cube

◆ **Query 4.5:** *Top three best-selling employees*

Sales1  $\leftarrow$  ROLLUP\*(Sales, Employee  $\rightarrow$  Employee, SUM(SalesAmount))  
Result  $\leftarrow$  MAX(Sales1, SalesAmount, 3)

◆ **Query 4.6:** *Best selling employees per product per year*

Sales1  $\leftarrow$  ROLLUP\*(Sales, Employee  $\rightarrow$  Employee, Product  $\rightarrow$  Product, OrderDate  
SUM(SalesAmount))  
Result  $\leftarrow$  MAX(Sales1, SalesAmount) BY Product, OrderDate

◆ **Query 4.7:** *Countries that account for top 50% of sales amount*

Sales1  $\leftarrow$  ROLLUP\*(Sales, Customer  $\rightarrow$  Country, SUM(SalesAmount))  
Sales2  $\leftarrow$  SORT(Sales1, Customer, SalesAmount DESC)  
Result  $\leftarrow$  TOPPERCENT(Sales2, Customer, SalesAmount, 50)



## Querying the Northwind Cube

- ◆ **Query 4.8:** *Total sales and average monthly sales by employee and year*

```
Sales1 ← ROLLUP*(Sales, Employee → Employee, OrderDate → Month, SUM(SalesAmount))  
Result ← ROLLUP*(Sales1, Employee → Employee, OrderDate → Year,  
SUM(SalesAmount), AVG(SalesAmount))
```

- ◆ **Query 4.9:** *Total sales amount and total discount amount per product and month*

```
Sales1 ← ROLLUP*(Sales, Product → Product, OrderDate → Month, SUM(SalesAmount))  
Result ← ADDMEASURE(Sales1, TotalDisc = Discount * Quantity * UnitPrice)
```

## Querying the Northwind Cube

- ◆ **Query 4.10:** *Monthly year-to-date sales for each product category*

```
Sales1 ← ROLLUP*(Sales, Product → Category, OrderDate → Month, SUM(SalesAmount))  
Result ← ADDMEASURE(Sales1, YTD = SUM(SalesAmount) OVER Date.Year  
ALL CELLS PRECEDING)
```

- ◆ **Query 4.11:** *Three-month moving average of the sales amount by product category*

```
Sales1 ← ROLLUP*(Sales, Product → Category, OrderDate → Month, SUM(SalesAmount))  
Result ← ADDMEASURE(Sales1, MovAvg3M = AVG(SalesAmount) OVER Date  
2 CELLS PRECEDING)
```

- ◆ **Query 4.12:** *Total sales amount made by an employee and his/her subordinates during 2017*

```
Sales1 ← ROLLUP*(Sales, Employee → Employee, OrderDate → Year, SUM(SalesAmount))  
Sales2 ← SLICE(Sales1, OrderDate.Year = 2017)  
Result ← RECROLLUP(Sales2, Employee → Employee, SUM(SalesAmount))
```

## Querying the Northwind Cube

- ◆ **Query 4.13:** *Total sales amount, number of products, and number of units sold (i.e., the sum of the quantities) by order*

ROLLUP\*(Sales, Order → Order, SUM(SalesAmount), COUNT(Product) AS CountP

- ◆ **Query 4.14:** *Total number of orders, total sales amount, and average sales amount by order, all by month*

Sales1 ← ROLLUP\*(Sales, OrderDate → Month, Order → Order, SUM(SalesAmount

Result ← ROLLUP\*(Sales1, OrderDate → Month, SUM(SalesAmount),  
AVG(SalesAmount) AS AvgSales, COUNT(Order) AS CountOrders)

- ◆ **Query 4.15:** *Number of cities and number of states assigned to each employee*

ROLLUP\*(Employee, Employee → State, COUNT(City), COUNT(State))