

Data Warehousing View Materialization

Esteban Zimányi

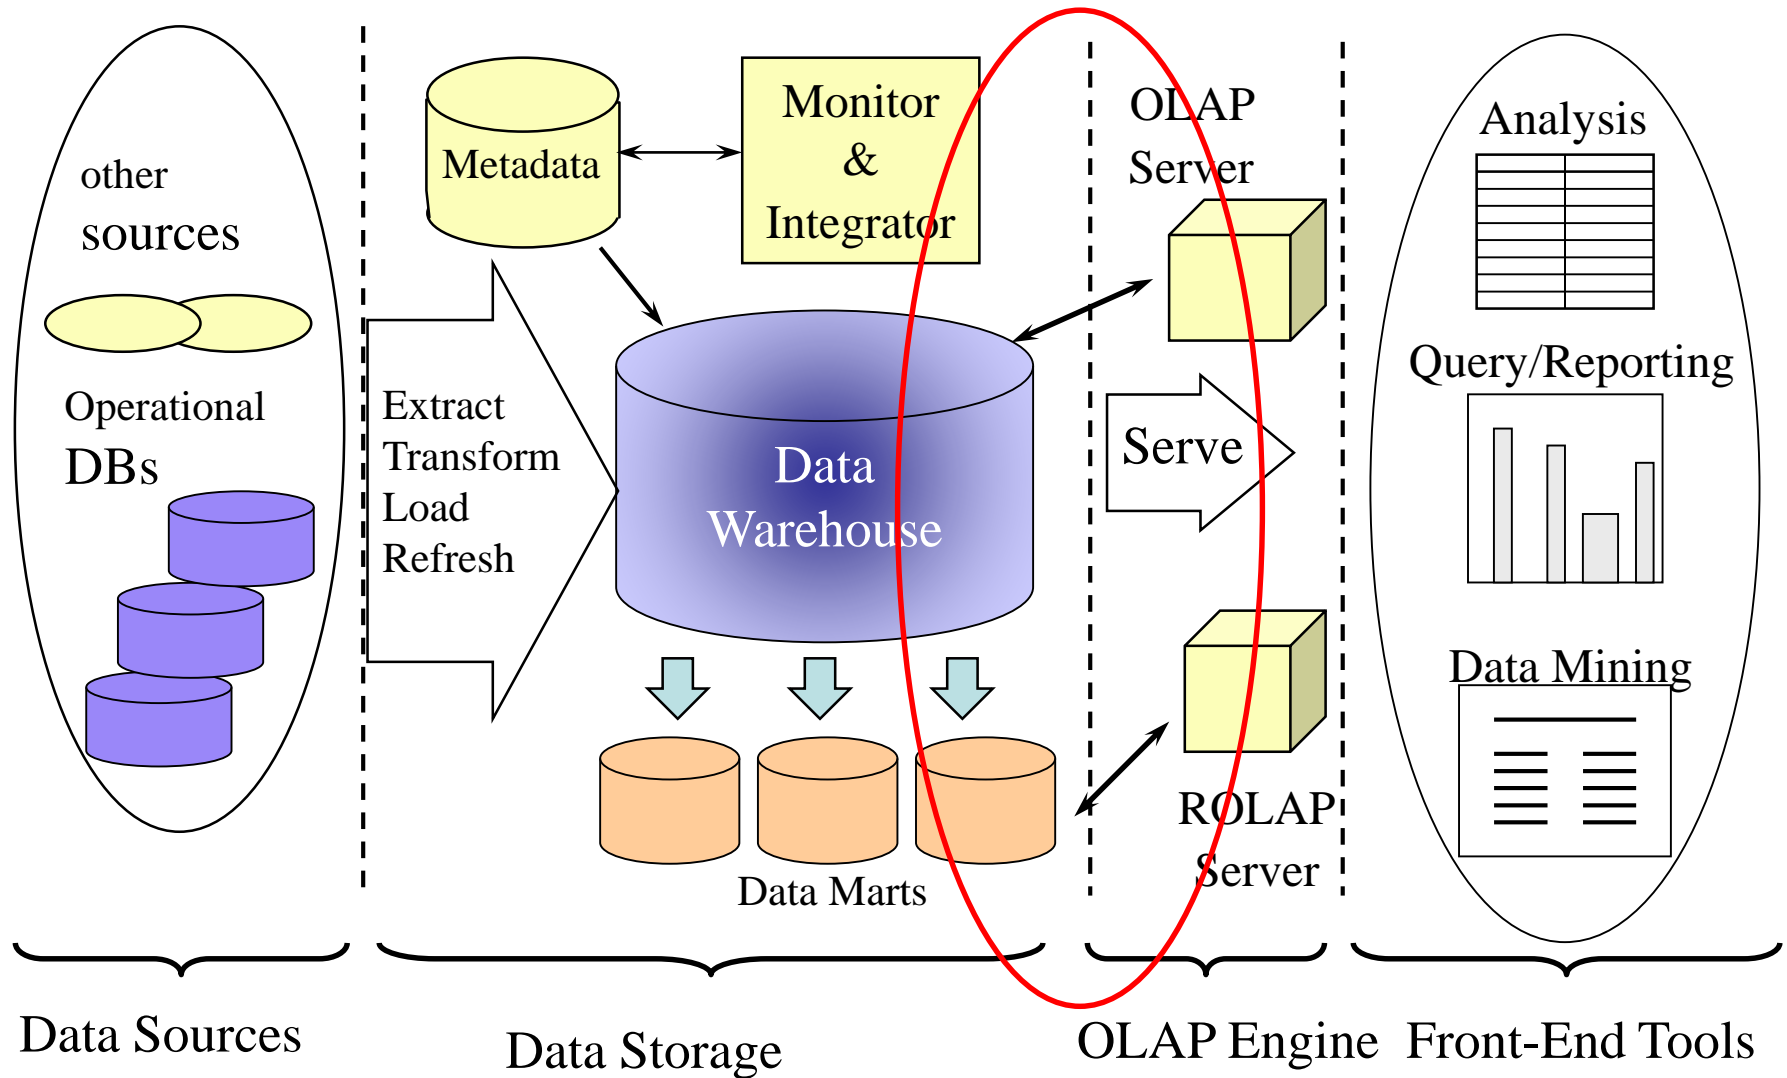
ezimanyi@ulb.ac.be

Slides by Toon Calders

Outline


- Problem statement
- Formal model of computation
 - Partial order on Queries
 - Cost model
- Greedy solution
 - Performance guarantee
- Extensions
- Conclusion

Previous Lecture




Data Explosion Problem

Size of cube w.r.t. number of dimensions (500 data points)



Data Explosion Problem

Country	Brand	Sales
FR	A	123
NL	B	456
BE	C	678
US	D	254
US	E	134



View materialization

- One of the major optimization techniques for OLAP
- Main question: Which parts to materialize to optimize performance?
 - Explosion problem: Pre-computing everything is impossible

Techniques in this slide set are based upon the following seminal work:

V. Harinarayan, A. Rajaraman, and J. D. Ullman: Implementing data cubes efficiently. In: ACM SIGMOD 1996.

The problem: informally

- Relation Sales
 - Dimensions: part, supplier, customer
 - Measure: sales
 - Aggregation: sum
- The data cube contains all aggregations on part, supplier, customer; e.g.:
 - (p1,s1,<all>,165)
 - (p1,s2,<all>,145)

Queries

- We are interested in answering the following *type* of queries efficiently:

(part, supplier)

```
SELECT part, supplier, sum(sales)
FROM Sales
GROUP BY part, supplier
```

(part)

```
SELECT part, sum(sales)
FROM Sales
GROUP BY part
```

(supplier)

```
SELECT supplier, sum(sales)
FROM Sales
GROUP BY supplier
```


(part, customer)

```
SELECT part, customer, sum(sales)
FROM Sales
GROUP BY part, customer
```

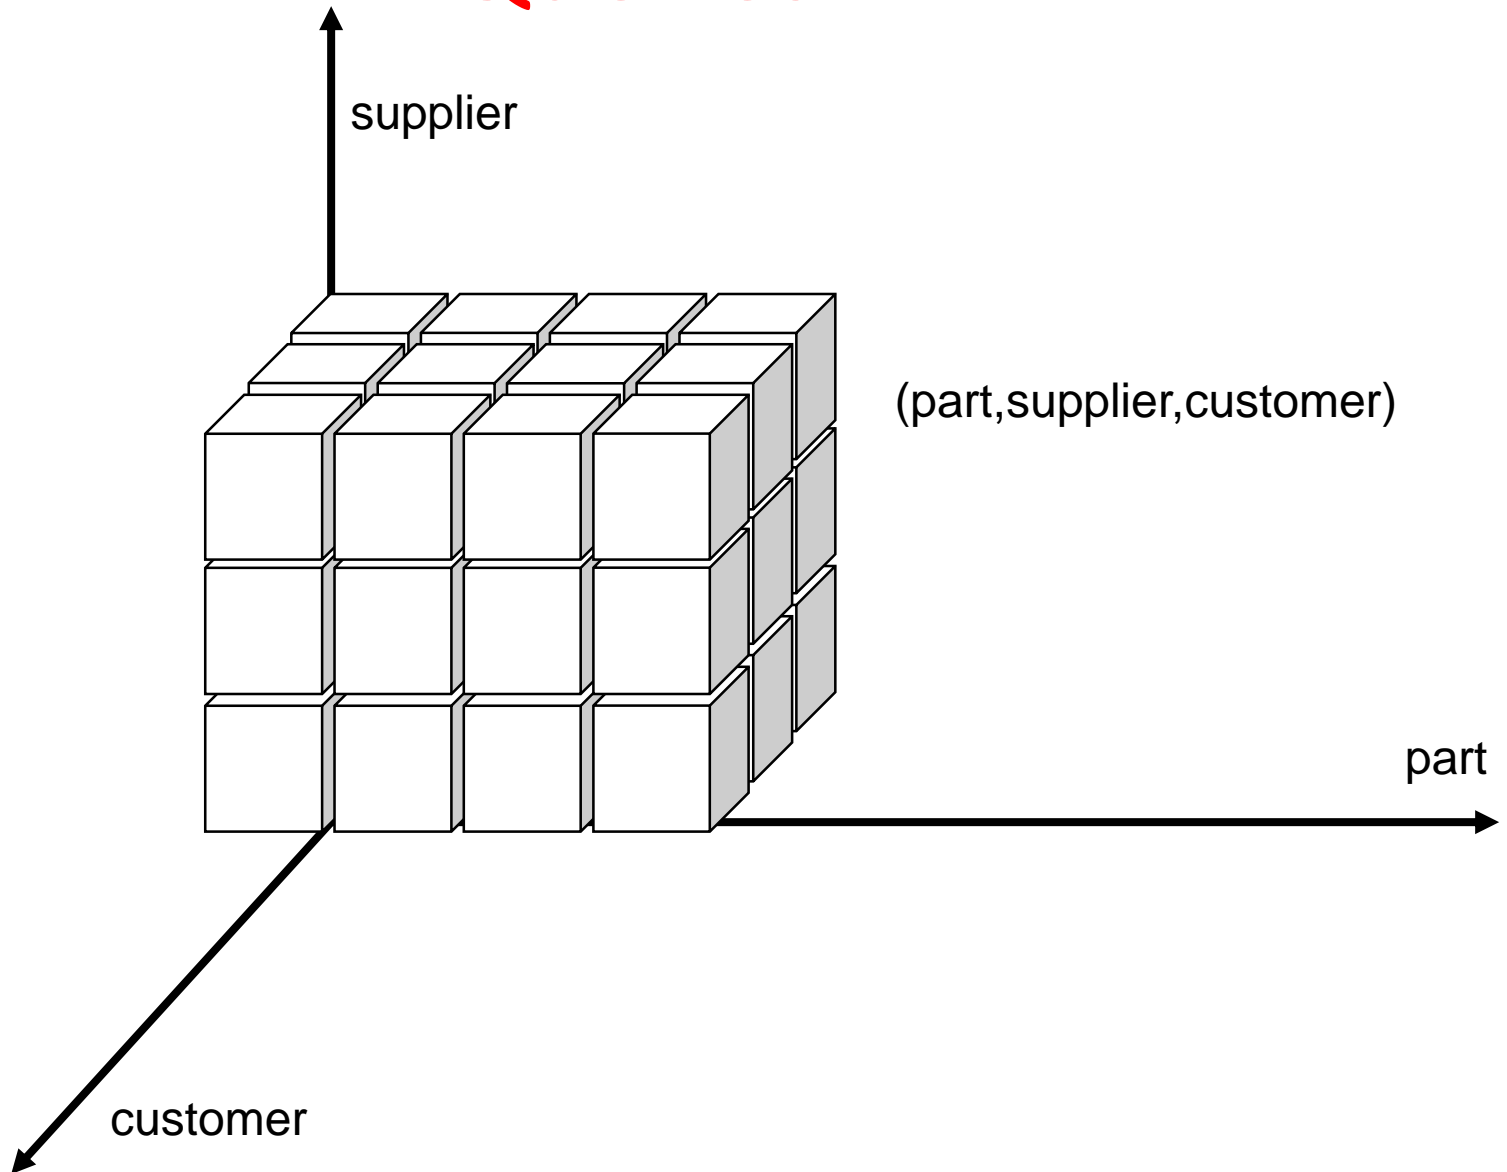

Queries

- The queries are quantified by their *grouping attributes*
- These queries select *disjoint* parts of the cube.
- The *union* of all these queries is the complete cube.

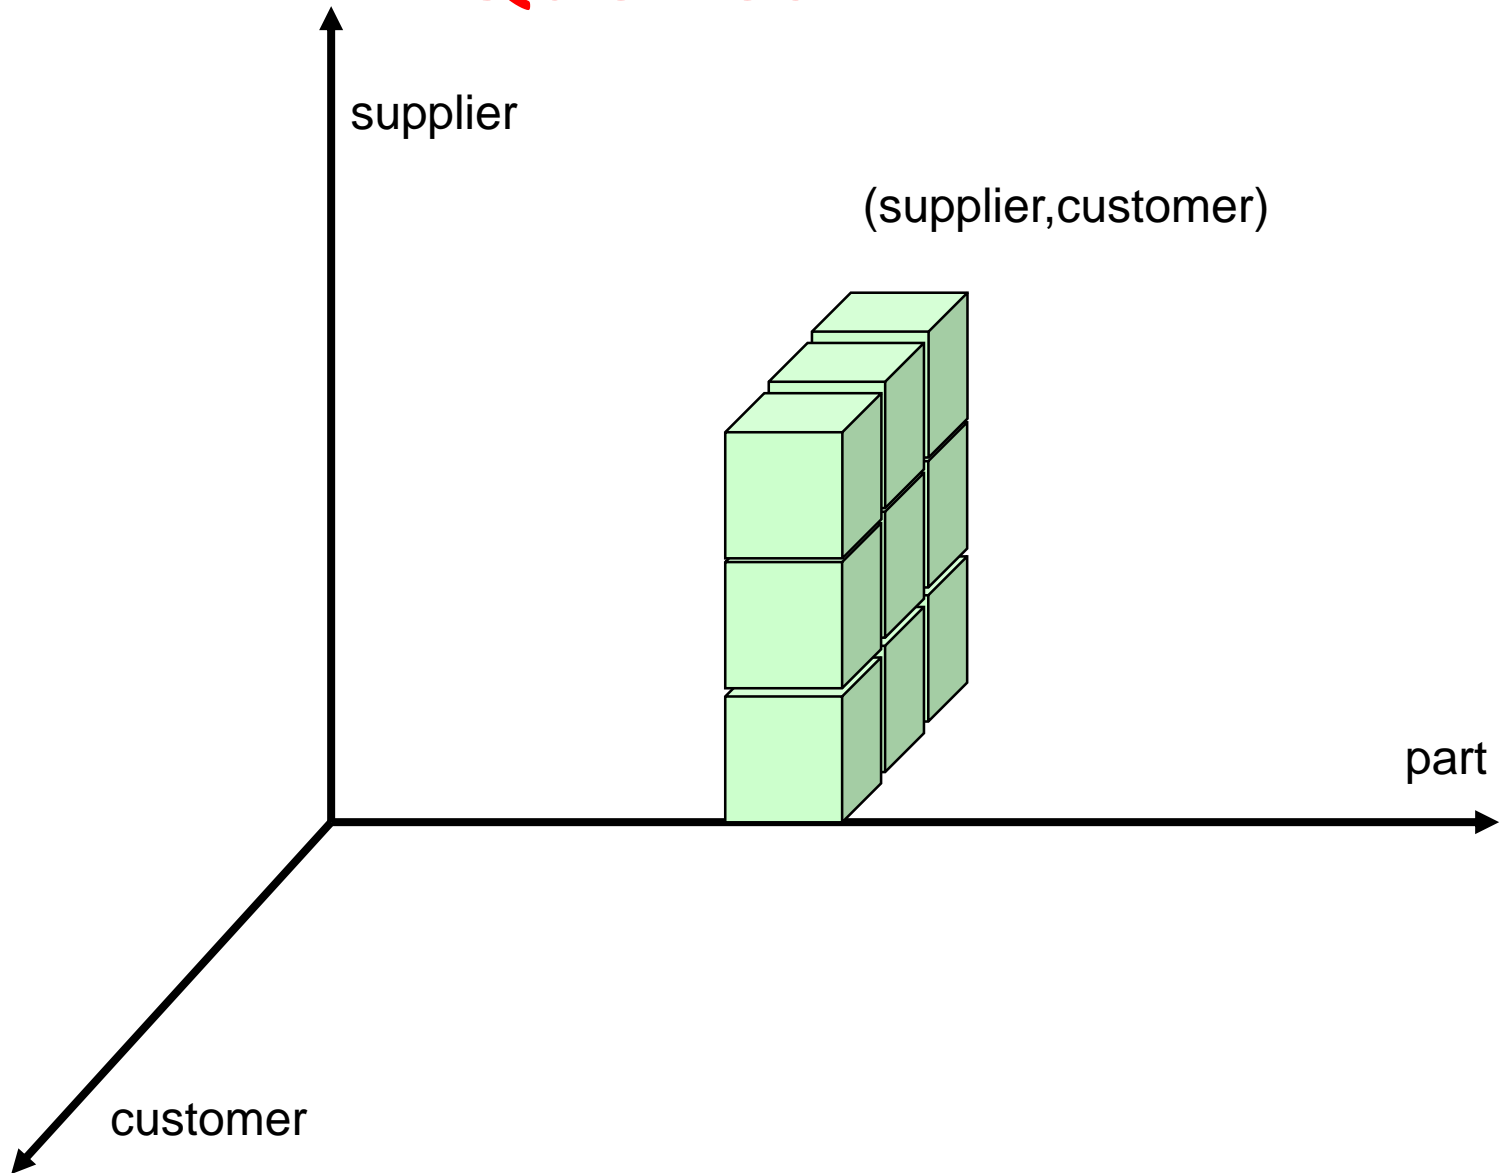
Queries




Queries



Queries



Queries



Outline

- Problem statement
- Formal model of computation
 - Partial order on Queries
 - Cost model
- Greedy solution
 - Performance guarantee
- Extensions
- Conclusion

Materialization

- Materializing everything is impossible
- Cost of evaluating following query directly:

```
SELECT D1, ..., Dk, sum(M)
FROM R
GROUP BY D1, ..., Dk
```

- in practice roughly linear in the size of R
(worst case: $|R| \log(|R|)$)
- Materializing *some* queries can help the other queries too

Materialization

Example:

```
( part, customer )  
SELECT customer, part, sum(sales)  
FROM Sales | Sales |  
GROUP BY customer, part
```

```
( part )  
SELECT part, sum(sales)  
FROM Sales | Sales |  
GROUP BY part
```


Materialization

Example:

(part, customer)

```
SELECT customer, part, sum(sales)
FROM Sales
GROUP BY customer, part
```

materialized as PC

|PC|

(part)

```
SELECT part, sum(sales)
FROM PC
GROUP BY part
```

|PC|

Caveat!

- For *additive measures-dimension* pairs and *distributive aggregation operators*
 - Measures that can be aggregated meaningful in any way
E.g., amount of sales, number of orders
- Careful with semi-additive
 - E.g., inventory level
- For algebraic aggregation operators: store support measures

Example

Query	Answer
• (part,supplier,customer)	6M
• (part,customer)	6M
• (part,supplier)	0.8M
• (supplier,customer)	6M
• (part)	0.2M
• (supplier)	0.01M
• (customer)	0.1M
• ()	1

Example: nothing materialized

Query	Answer	Cost
• (part,supplier,customer)	6M	6M
• (part,customer)	6M	6M
• (part,supplier)	0.8M	6M
• (supplier,customer)	6M	6M
• (part)	0.2M	6M
• (supplier)	0.01M	6M
• (customer)	0.1M	6M
• ()	1	6M

Total cost: 48M

Example: some materialized

Query	Answer	Cost
• <u>(part,supplier,customer)</u>	6M	
• (part,customer)	6M	
• <u>(part,supplier)</u>	0.8M	
• (supplier,customer)	6M	
• (part)	0.2M	
• (supplier)	0.01M	
• <u>(customer)</u>	0.1M	
• ()	1	

Example: some materialized

Query	Answer	Cost
• <u>(part,supplier,customer)</u>	6M	<u>6M</u>
• (part,customer)	6M	6M
• <u>(part,supplier)</u>	0.8M	<u>0.8M</u>
• (supplier,customer)	6M	6M
• (part)	0.2M	<u>0.8M</u>
• (supplier)	0.01M	<u>0.8M</u>
• <u>(customer)</u>	0.1M	<u>0.1M</u>
• ()	1	<u>0.1M</u>

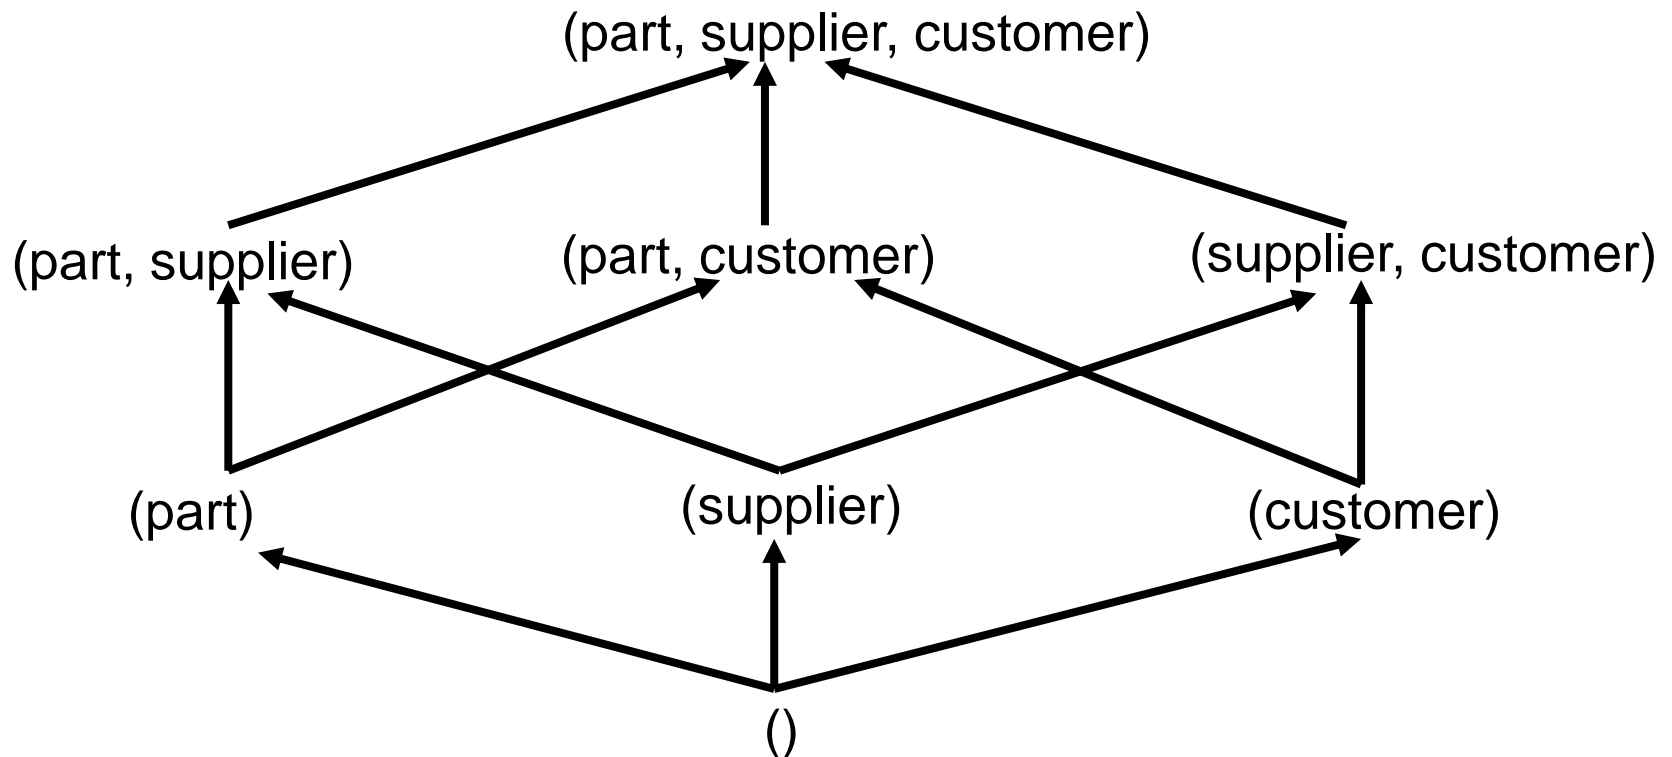
Total cost: 20.6M

Research Questions

- What is the *optimal set of views* to materialize?
 - *How many views* must we materialize to get reasonable performance?
 - Given bounded space S , what is the *optimal choice* of views to materialize?
- First important step towards a solution: cost model
 - Estimate the benefit of materializing a view – associated with an evaluation strategy
 - Often *impossible to know* the sizes of the views without actually computing them

Partial Order on Queries

- partial order on the views: $Q1 \leq Q2$:
 - Q1 can be answered using the query results of Q2
 - *materialized view* of Q2 can be used to answer Q1



Partial order on queries

- Can be generalized to hierarchies:
 $(\text{product}, \text{country}, -) \leq (\text{product}, \text{city}, -)$
 $(-, \text{country}, -) \leq (\text{product}, \text{city}, -)$
 $(-, -, -) \leq (\text{product}, \text{city}, -)$
 $(-, \text{country}, \text{year}) \leq (-, \text{city}, \text{month})$
- This is actually the normal specialization relation of the roll-up lattice

Cost Model

- Hence, given a set of materialized views S

$$S = \{\text{BaseRelation}, V_1, \dots, V_n\},$$

Unique top element in the lattice

- The *cost of query Q given S* is

$$\text{cost}(Q|S) := \min(\{ |V_i| : i = 1 \dots n, Q \leq V_i \})$$

- The *total cost given S* is:

$$\text{cost}(S) := \sum_Q \text{cost}(Q|S)$$

- The *benefit of S* is:

$$B_S := \text{cost}(\{\text{BaseRelation}\}) - \text{cost}(S)$$

Outline

- Problem statement
- Formal model of computation
 - Partial order on Queries
 - Cost model
- Greedy solution
 - Performance guarantee
- Extensions
- Conclusion

Greedy Solution for One Variant

- Class of *view selection problems*:

$$\begin{array}{ll} \text{minimize cost}(S) \text{ s.t.} & \\ |S| \leq k & \text{(bound \# views)} \\ \sum_{V \in S} |V| \leq M & \text{(space bound)} \\ \dots & \end{array}$$

- We will treat one particular case:
 - Given: for every view its size, k
Find: k views giving the highest
- This problem is NP-complete*
- greedy algorithm with good guarantees

Greedy algorithm

- Input:
 - Size for every view
 - Relation \leq between views
 - Constant k

Greedy:

$S = \{ \text{top level view} \}$

Repeat k times:

 Select W that minimizes $\text{cost}(S \cup \{W\})$

$S := S \cup \{W\}$


Return S

“Benefit” of W given S :

$$\begin{aligned} B(W|S) &= \text{Cost}(S) - \text{Cost}(S \cup \{W\}) \\ &= \sum_Q \text{cost}(Q|S) - \sum_Q \text{cost}(Q|S \cup \{W\}) \\ &= \sum_Q [\text{cost}(Q|S) - \text{cost}(Q|S \cup \{W\})] \\ &= \sum_Q B_Q(W|S) \end{aligned}$$

Example


- Only base table a is materialized
 - Total: $8 \times 100 = 800$



Query	Size	Uses?	Cost
a	100	a	100
b	50	a	100
c	75	a	100
d	20	a	100
e	30	a	100
f	40	a	100
g	1	a	100
h	10	a	100

Example

- Base table a and table b are materialized
 - Total: $3 \times 100 + 5 \times 50 = 550$




Query	Size	Uses?	Cost	Benefit
a	100	a	100	-
b	50	b	50	50
c	75	a	100	-
d	20	b	50	50
e	30	b	50	50
f	40	a	100	-
g	1	b	50	50
h	10	b	50	50

- Benefit of storing b = $250 = 5 \times (100-50)$

Example

- Benefit for materializing the other tables:




query	Materialized view						
	b	c	d	e	f	g	
a	-	-	-	-	-	-	-
b	50	-	-	-	-	-	-
c	-	25	-	-	-	-	-
d	50	-	80	-	-	-	-
e	50	25	-	70	-	-	-
f	-	25	-	-	60	-	-
g	50	25	80	70	-	99	-
h	50	25	-	70	60	-	90
Total	<u>250</u>	125	160	210	120	99	90

a	100
b	50
c	75
d	20
e	30
f	40
g	1
h	10

Example

- Base table a, table b, and f are materialized
 - Total: $2 \times 100 + 4 \times 50 + 2 \times 40 = 480$




Query	Size	Uses?	Cost	Benefit
a	100	a	100	-
b	50	b	50	-
c	75	a	100	-
d	20	b	50	-
e	30	b	50	-
f	40	f	40	60
g	1	b	50	-
h	10	f	40	10

- Additional benefit of materializing f
 - $= 70 = 1 \times (100 - 40) + 1 \times (50 - 40)$

Example

- Benefit for materializing the other tables:




query	Materialized view					
	c	d	e	f	g	h
a	-	-	-	-	-	-
b	-	-	-	-	-	-
c	25	-	-	-	-	-
d	-	30	-	-	-	-
e	-	-	20	-	-	-
f	25	-	-	60	-	-
g	-	30	20	-	49	-
h	-	-	20	10	-	40
Total	50	60	60	<u>70</u>	49	40

a	100
b	50
c	75
d	20
e	30
f	40
g	1
h	10

Example

- Benefit for materializing a third table:




query	Materialized view				
	c	d	e	g	h
a	-	-	-	-	-
b	-	-	-	-	-
c	25	-	-	-	-
d	-	30	-	-	-
e	-	-	20	-	-
f	-	-	-	-	-
g	-	30	20	49	-
h	-	-	10	-	30
Total	25	<u>60</u>	50	49	30

a	100
b	50
c	75
d	20
e	30
f	40
g	1
h	10

Example

- Choice of 3 views according to the Greedy algorithm:
 - Benefit = $250 + 70 + 60 = 380$
 - Cost = 420



Outline

- Problem statement
- Formal model of computation
 - Partial order on Queries
 - Cost model
- Greedy solution
 - Performance guarantee
- Extensions
- Conclusion


Some theoretical results

- It can be proved that we can get at least $(e - 1) / e$ (which is about 63%) of the benefit of the optimal algorithm
- When selecting k views, bound is:

$$B_{\text{greedy}} / B_{\text{opt}} \geq 1 - \left(\frac{k-1}{k} \right)^k$$

There are lattices for which this ratio is arbitrarily close to this bound.


Example



a	200
b	100
c	99
d	100
e	50
f	50
g	50
h	50

Node	Benefit
b, d	$3 \times (200-100) = 300$
<u>c</u>	$3 \times (200-99) = 303$
e, f, g, h	$1 \times (200-50) = 150$


Example



a	200
b	100
c	99
d	100
e	50
f	50
g	50
h	50

Node	Benefit
b, d	$2 \times (200 - 100) = 200$
e, h	$1 \times (200 - 50) = 150$
f, g	$1 \times (99 - 50) = 49$


Example



a	200
b	100
c	99
d	100
e	50
f	50
g	50
h	50

- Hence $B_{\text{greedy}} = 503$


Example



a	200
b	100
c	99
d	100
e	50
f	50
g	50
h	50

- Hence $B_{\text{greedy}} = 503$
- $B_{\text{opt}} = 6 \times (200 - 100) = 600$
- $B_{\text{greedy}} / B_{\text{opt}} = 503/600 \geq 75\%$

Example



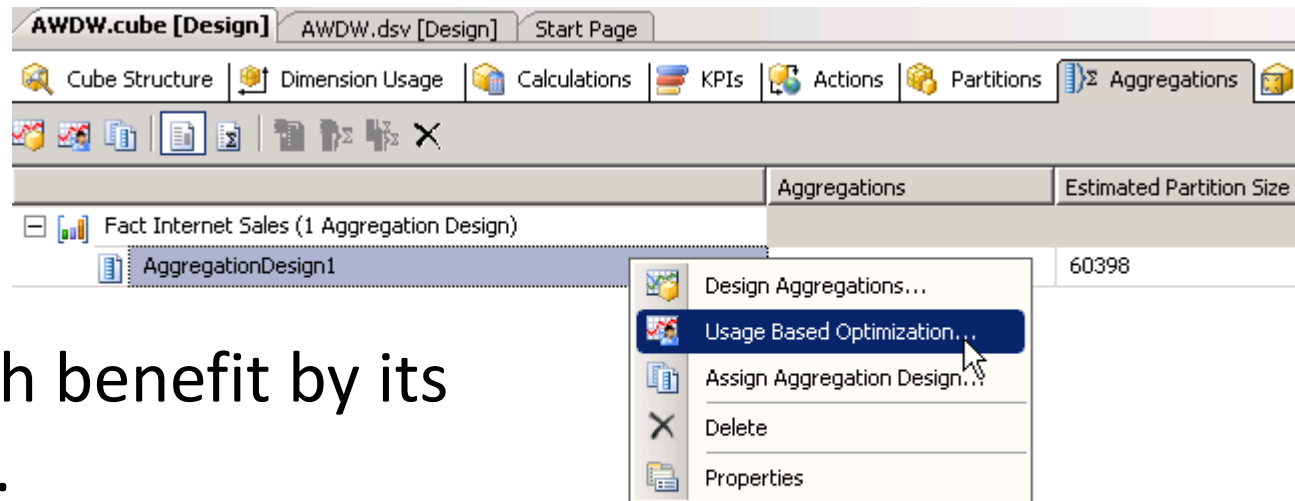
- Hence $B_{\text{greedy}} = (2N + 1) \times (100 + \varepsilon) + (N + 1) \times 100 \rightarrow 300N + 200$
- $B_{\text{opt}} = (4N + 2) \times 100 \rightarrow 400N + 200$
- $B_{\text{greedy}} / B_{\text{opt}} \rightarrow (3N + 2) / (4N + 2) \rightarrow 75\%$

Outline

- Problem statement
- Formal model of computation
 - Partial order on Queries
 - Cost model
- Greedy solution
 - Performance guarantee
- Extensions
- Conclusion

Extensions (1)


- Problem
 - The views in a lattice are unlikely to have the same probability of being requested in a query.
- Solution:
 - We can weight each benefit by its probability.



Example

- Suppose following distribution:

a: 0%, b: 10%, c:0%, d:10%, e:5%, f:20%, g: 40% , h:15%



a	100
b	50
c	75
d	20
e	30
f	40
g	1
h	10


Query	Uses?	Cost	Exp. Cost
a	a	100	0
b	b	50	5
c	a	100	0
d	b	50	5
e	b	50	2.5
f	f	40	8
g	b	50	20
h	f	40	6
Total			46.5

- Expected cost given $S = 46.5$

Example

- Solution found by greedy:

a: 0%, b: 10%, c:0%, d:10%, e:5%, f:20%, g: 40% , h:15%



Query	Uses?	Cost	Exp. Cost
a	a	100	0
b	a	100	10
c	a	100	0
d	d	20	2
e	e	30	1.5
f	a	100	20
g	d	20	8
h	e	30	4.5
Total			46


a	100
b	50
c	75
d	20
e	30
f	40
g	1
h	10

- Expected cost given $S = 46$

Example

- Optimal Solution:

a: 0%, b: 10%, c:0%, d:10%, e:5%, f:20%, g: 40% , h:15%



Query	Uses?	Cost	Exp. Cost
a	a	100	0
b	a	100	10
c	a	100	0
d	d	20	2
e	a	100	5
f	f	40	8
g	d	20	8
h	f	40	6
Total			39

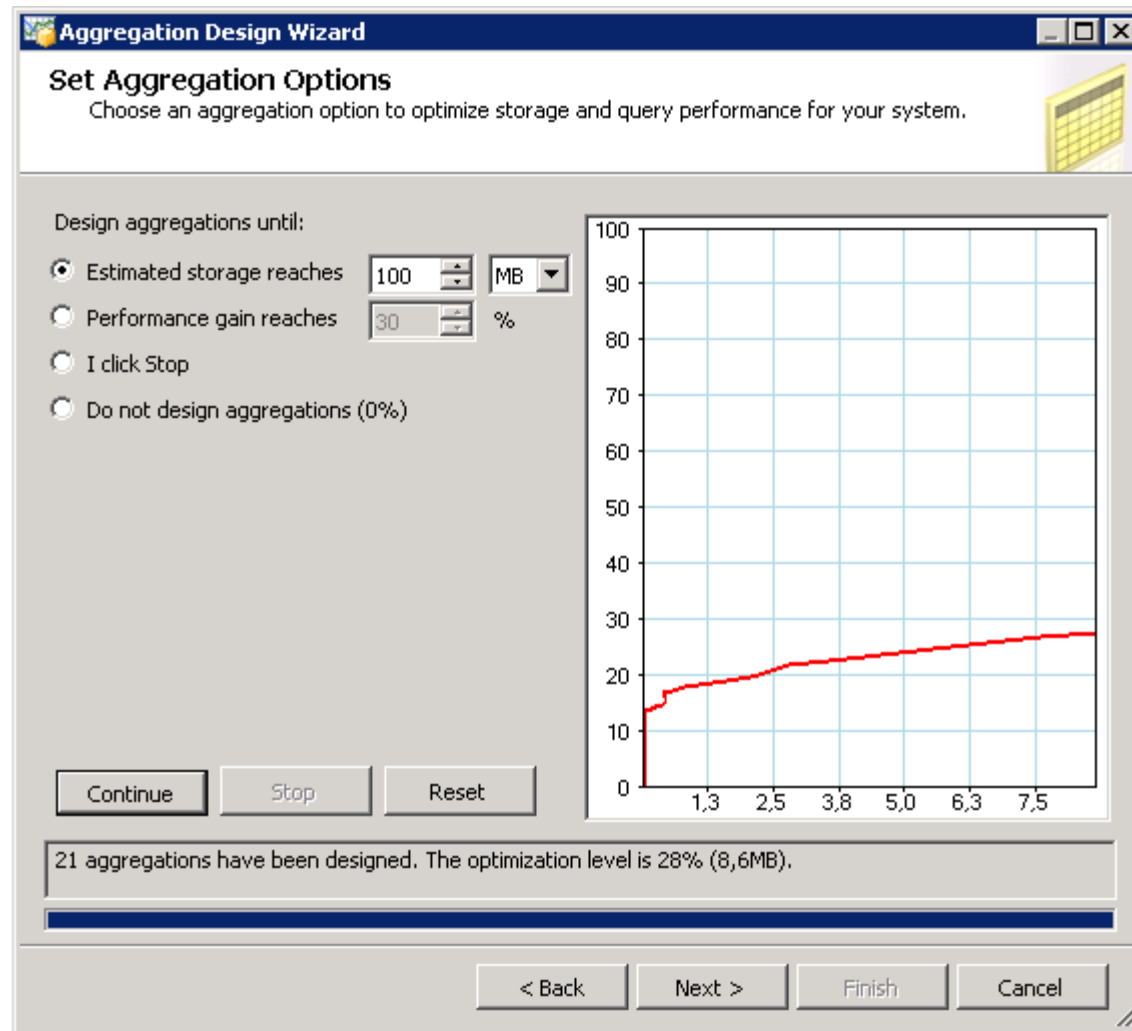
a	100
b	50
c	75
d	20
e	30
f	40
g	1
h	10

- Expected cost given $S = 39$

Extensions (2)

- Problem
 - Instead of asking for some fixed number (k) of views to materialize, we might instead allocate a fixed amount of space to views.
- Solution
 - We can consider the “benefit of each view per unit space”.

View Materialization



Conclusion: View Materialization

- Materialization of views is an essential query optimization strategy for decision-support applications.
 - Finding optimal solution is NP-hard.
- Introduction of greedy algorithm
 - There exists cases which greedy algorithm fails to produce optimal solution
 - Yet, greedy algorithm has a good guarantee
- Expansion of greedy algorithm.