

DATA MINING (INFO-H-423)
Mahmoud SAKR
Intended Learning Outcomes

Introduction & Decision Tree
<ul style="list-style-type: none"> • Explain the steps of supervised learning. • Describe classification as one of the DM tasks. • Explain/Illustrate the concepts of Entropy, Gain, Gain-ratio. • Explain and Apply the ID3 Algorithm to a given dataset.
Data Preparation and Distance Measures
<ul style="list-style-type: none"> • Illustrate and compare the different model validation methods: holdout, cross validation, bootstrap. • Describe the different data preparation tasks, their goals, and few examples. The details of the methods are not required. It is enough to know the general idea. • Describe without detail the concepts and methods of data cleaning, scaling, normalization, and distance measures. • Explain and apply the Lp-norm distance, and the Edit distance.
Clustering
<ul style="list-style-type: none"> • Explain clustering as an unsupervised learning method. • Apply the K-Means, the K-Medians, and the DBSCAN Algorithms to given datasets. • Compare K-Means, K-Medians and K-Medoids Algorithms. • Compare between the different families of clustering Algorithms (representative-based, density based, probabilistic model-based). • Analyze the complexity of the studied clustering Algorithms. • Explain how to assess clustering quality <p>Note: the details of the probabilistic model-based clustering are not required for this exam. Only the general concept of fuzzy clustering is required.</p>
Timeseries Forecasting
<ul style="list-style-type: none"> • Explain the different types of forecasting models • Illustrate with examples the trend, seasonality, and cycle components of a time series • Apply classical timeseries decomposition for a given timeseries • Illustrate timeseries forecasting with decomposition • Describe mathematical concepts for analyzing time series, including concepts of white noise, stationarity, and autocorrelation • Explain AR model, and parameters • Explain MA model, and parameters • Explain ARMA, ARIMA models, and parameters • Explain the Prophet model
Outlier Mining
<ul style="list-style-type: none"> • Illustrate with examples the types of outliers: global, contextual, and collective • Illustrate the use of grids in speeding up distance based outlier detection • Discuss in general terms the different methods of outlier mining: distance based-clustering based, classification based.
Frequent itemsets & Association rules Mining
<ul style="list-style-type: none"> • Illustrate the uses of frequent pattern mining. • Explain and compare the measures: support, confidence, lift, correlation analysis.

- Explain the Apriori property, and its applications in optimizing the search for frequent itemsets.
- Apply the Apriori and the FP-Growth Algorithms to given datasets.
- Compare the frequent pattern mining Algorithms: brute force enumeration of the itemset lattice, Apriori, and FP-Growth Algorithms.
- Apply your understanding of the different frequent itemset mining methods to reason about further optimizations.
- Extract association rules from frequent itemsets, and assess their quality.

Classification

- Explain the naive Bayes classifier, and apply it to a given dataset.
- Explain the concept of the confusion matrix, and use it to define the different classification quality measures.
- Motivate ensemble learning, as a method to improve classification accuracy.
- Illustrate the difference between bagging and boosting.
- Describe random forest induction, and the tuning of parameters L, D.

Stream Mining

- Describe and illustrate Bloom filter, Count-Min, Flajolet-Martin, and hyperloglog
- Apply these methods to query real data streams

Note: No need to memorize equations and Algorithms. They will be given if needed.

With my best wishes.