# Distributed Data Management

Big Data Management

Knowledge objectives

1. Give a definition of Distributed System
2. Enumerate the 6 challenges of a Distributed System
3. Give a definition of Distributed Database
4. Explain the different transparency layers in DDBMS
5. Identify the requirements that distribution imposes on the ANSI/SPARC architecture
6. Draw a classical reference functional architecture for DDBMS
7. Enumerate the 8 main features of Cloud Databases
8. Explain the difficulties of Cloud Database providers to have multiple tenants
9. Enumerate the 4 main problems tenants/users need to tackle in Cloud Databases
10. Distinguish the cost of sequential and random access
11. Explain the difference between the cost of sequential and random access
12. Distinguish vertical and horizontal fragmentation
13. Recognize the complexity and benefits of data allocation
14. Explain the benefits of replication
15. Discuss the alternatives of a distributed catalog

Understanding Objectives

- Decide when a fragmentation strategy is correct

Distributed Systems

Distributed System

Distributed DBMS

Cloud DBMS

www.essi.upc.edu/dtim

Distributed system

"One in which components located at networked computers communicate and coordinate their actions only by passing messages."

G. Coulouris et al.

- Characteristics:
  - Concurrency of components
  - Independent failures of components
  - Lack of a global clock

Challenges of distributed systems

- ~~Openness~~
- Scalability
- Quality of service
  - Performance/Efficiency
  - Reliability/Availability
  - ~~Confidentiality~~
- Concurrency
- Transparency
- ~~Heterogeneity of components~~

Scalability

Cope with large workloads
- Scale up ❌
- Scale out

- Use:
  - Automatic load-balancing

- Avoid:
  - Bottlenecks
  - Unnecessary communication
    - Peer-to-peer

Performance/Efficiency

Efficient processing
- Minimize latencies
- Maximize throughput

- Use
  - Parallelism
  - Network optimization
 - Specific techniques

Reliability/Availability

a) Keep consistency

b) Keep the system running
- Even in the case of failures

- Use
  - Replication
  - Flexible routing
  - Heartbeats
 - Automatic recovery

Concurrency

Share resources as much as possible

- Use
 - Consensus Protocols

- Avoid
  - Interferences
  - Deadlocks

Transparency

a) Hide implementation (i.e., physical) details to the users

b) Make transparent to the user all the mechanisms to solve the other challenges

Further objectives

- Use
  - Platform-independent software

- Avoid
  - Complex configurations
  - Specific hardware/software

Distributed Database Systems

Distributed database

"A Distributed DataBase (DDB) is an <u>integrated collection of databases</u> that is <u>physically distributed</u> across sites in a computer network. A Distributed DataBase Management System (DDBMS) is the software system that manages a distributed database such that the <u>distribution aspects are transparent</u> to the users."

Encyclopedia of Database Systems

Transparency layers (I)

- Fragmentation transparency
  - The user must not be aware of the existence of different fragments
- Replication transparency
  - The user must not be aware of the existing replicas
- Network transparency
  - Data access must be independent regardless where data is located
  - Each data object must have a unique name
- Data independency at the logical and physical level must be guaranteed
  - Inherited from centralized DBMSs (ANSI SPARC)

Transparency layers (II)
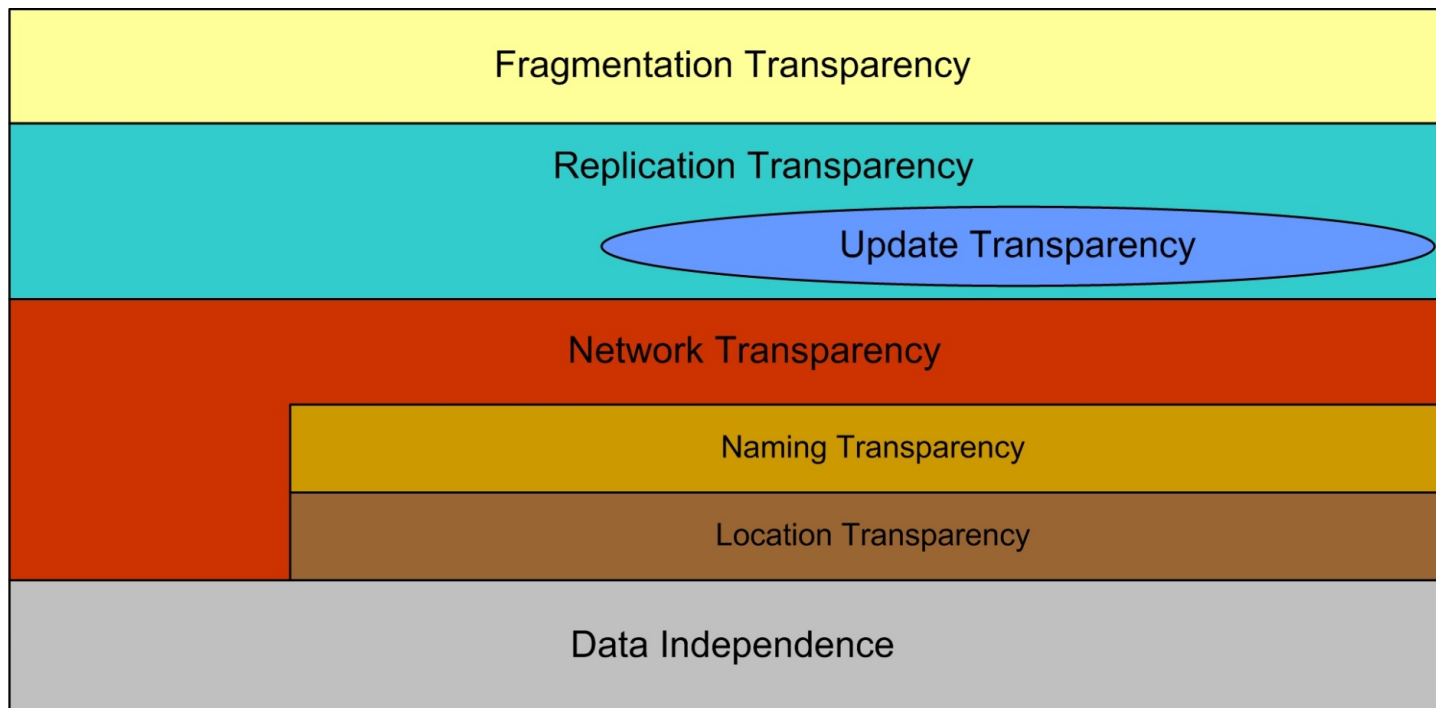
Classification According to Degree of Autonomy

	Autonomy	Central schema	Query transparency	Update transparency
DDBMS	No	Yes	Yes	Yes
T.C. Federated	Low	Yes	Yes	Limited
L.C. Federated	Medium	No	Yes	Limited
Multi-database	High	No	No	No

Extended ANSI-SPARC Architecture of Schemas
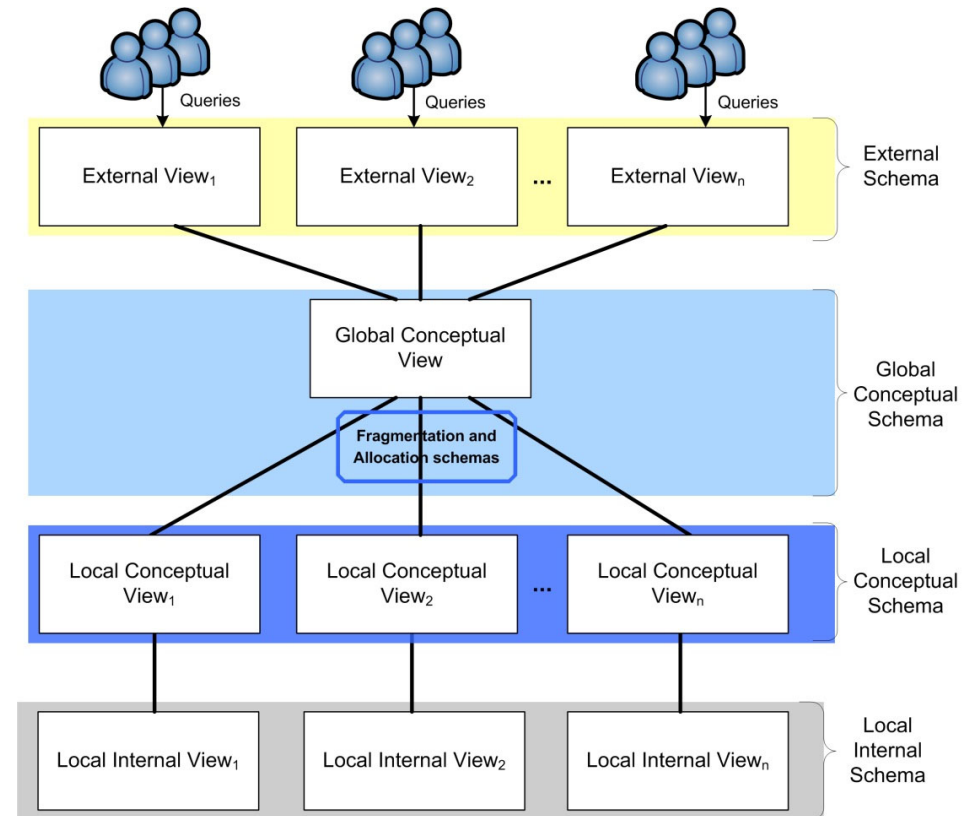
- Global catalog (Mappings between ESs – GCS and GCS – LCSs)
- Each node has a local catalog (Mappings between LCS_i – IS_i)

Centralized DBMS Functional Architecture
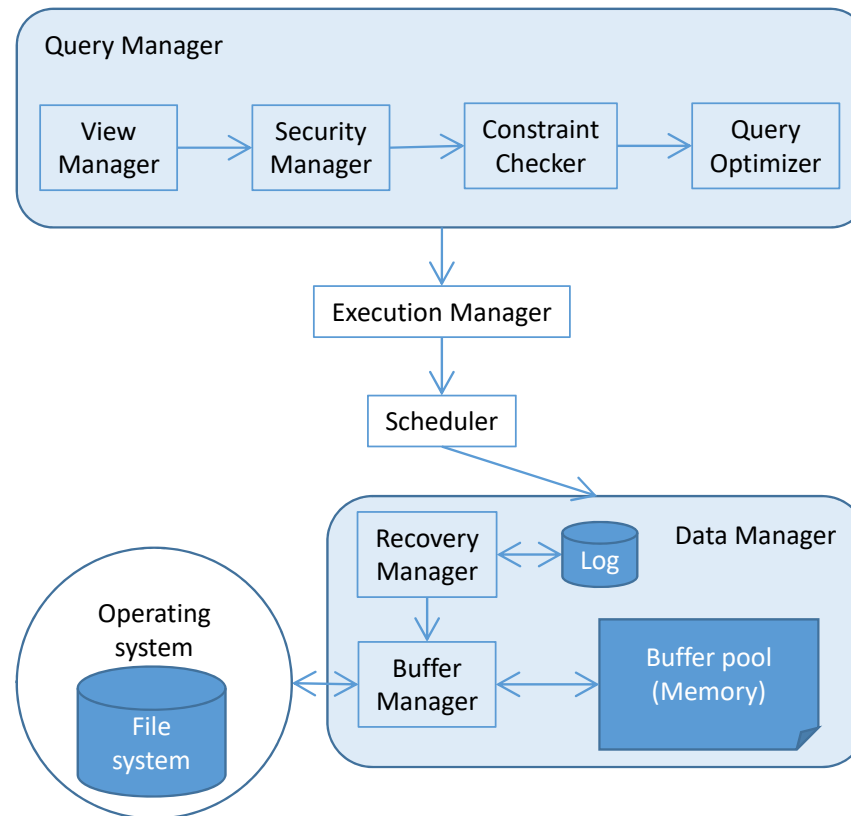
Distributed DBMS Functional Architecture

Cloud Databases

Distributed System

Distributed DBMS

Cloud DBMS

UNIVERSITAT POLITÈCNICA
DE CATALUNYA
BARCELONATECH
UPC

DTIM
www.essi.upc.edu/dtim
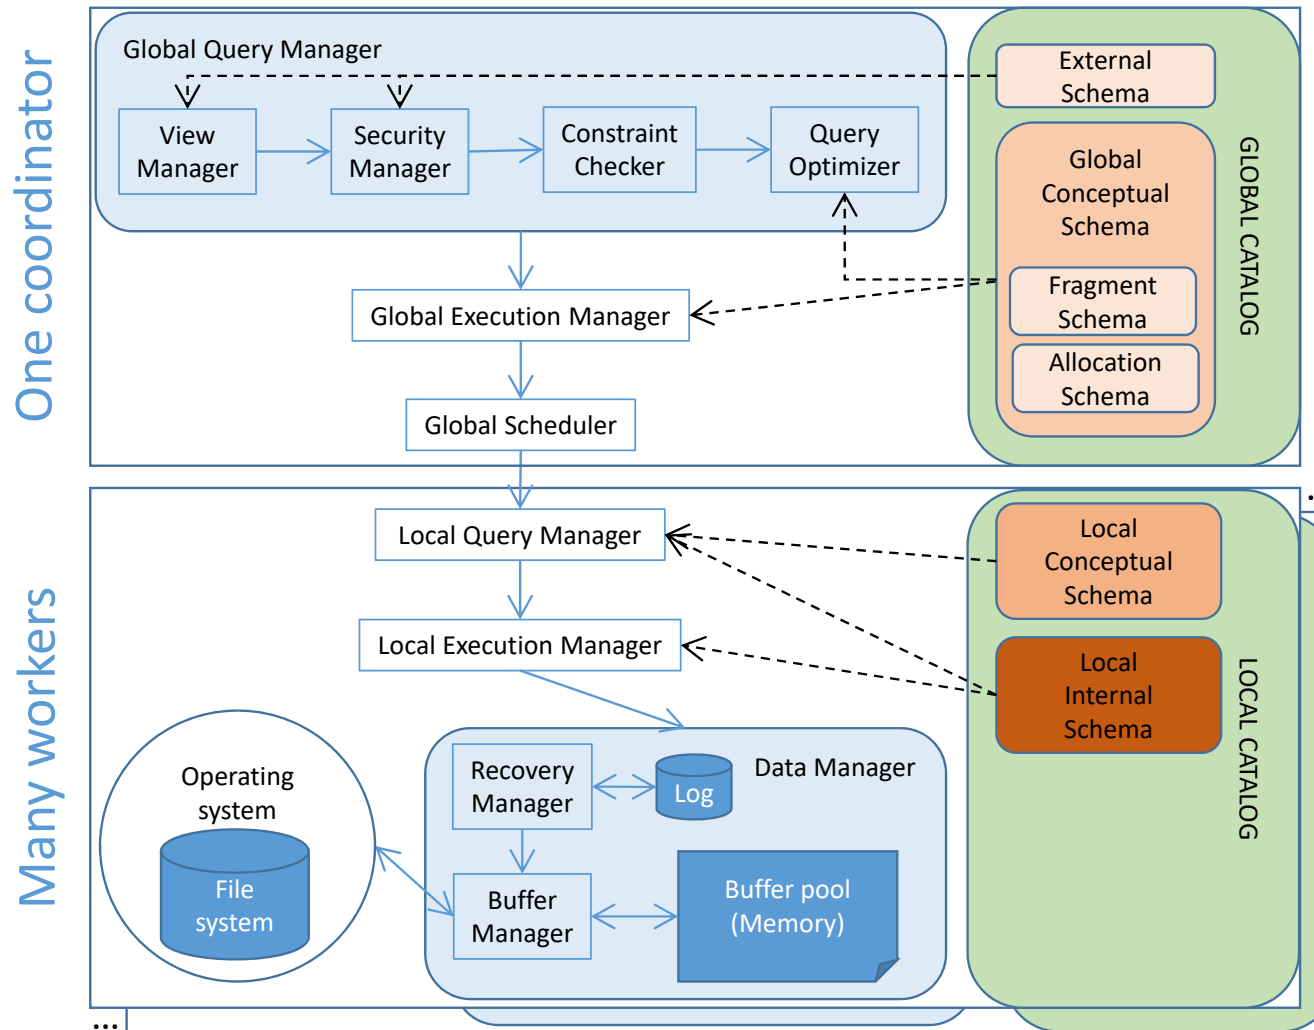
Parallel database architectures

D. DeWitt & J. Gray. Figure by D. Abadi

Key Features of Cloud Databases

- Scalability
    a) Ability to horizontally scale (scale out)
- Quality of service
    - Performance/Efficiency
        b) Fragmentation: Replication & Distribution
        c) Indexing: Distributed indexes and RAM
    - Reliability/Availability
- Concurrency
    d) Weaker concurrency model than ACID
- Transparency
    e) Simple call level interface or protocol
        - No declarative query language
- Further objectives
    f) Flexible schema
        - Ability to dynamically add new attributes
    g) Quick/Cheap set up
    h) Multi-tenancy

Multi-tenancy platform problems (provider side)

- Difficulty: Unpredictable load characteristics
  - Variable popularity
  - Flash crowds
  - Variable resource requirements
- Requirement: Support thousands of tenants
  a) Maintain metadata about tenants (e.g., activated features)
  b) Self-managing
  c) Tolerating failures
  d) Scale-out is necessary (sooner or later)
     - Rolling upgrades one server at a time
  e) Elastic load balancing
     - Dynamic partitioning of databases

DTIM
www.essi.upc.edu/dtim

Data management problems (tenant side)

I.   (Distributed) data design
 - Data fragmentation
 - Data allocation
 - Data replication

II.   (Distributed) catalog management
 - Metadata fragmentation
 - Metadata allocation
 - Metadata replication

III.   (Distributed) transaction management
 - Enforcement of ACID properties
   - Distributed recovery system
   - Distributed concurrency control system
   - Replica consistency
     - Latency&Availability vs. Update performance

IV.   (Distributed) query processing
 - Optimization considering
   1)   Distribution/Parallelism
     - Communication overhead
   2)   Replication

DTIM
www.essi.upc.edu/dtim

(Distributed) Data Design

Challenge I

DDB Design

- Given a DB and its workload, how should the DB be split and allocated to sites as to optimize certain objective functions
  - Minimize resource consumption for query processing

- Two main issues:
  - Data fragmentation
  - Data allocation
    - Data replication

Vertical fragmentation (join)

Primary horizontal fragmentation (union)

Data Fragmentation

- Usefulness
  - An application typically accesses only a subset of data
  - Different subsets are (naturally) needed at different sites
  - The degree of concurrency is enhanced
    - Facilitates parallelism
      - Fragments can be even defined dynamicaly (i.e., at query time, not at design time)
- Difficulties
  - Complicates the catalog management
  - May lead to poorer performance when multiple fragments need to be joined
    - Fragments likely to be used jointly can be colocated to minimize communication overhead
  - Costly to enforce the dependency between attributes in different fragments

Fragmentation Correctness

- Completeness
  - Every datum in the relation must be assigned to a fragment

- Disjointness
  - There is no redundancy and every datum is assigned to only one fragment
    - The decision to replicate data is in the allocation phase

- Reconstruction
  - The original relation can be reconstructed from the fragments
    - Union for horizontal fragmentation
    - Join for vertical fragmentation

Finding the best fragmentation strategy

- Consider it per table
  - Computational cost is NP-hard
- Needed information
  - Workload
    - Frequency of each query
    - Access plan and cost of each query
      - Take intermediate results and repetitive access into account
  - Value distribution and selectivity of predicates
- Work in three phases
  1. Determine primary partitions (i.e., attribute subsets often accessed together)
  2. Generate a disjoint and covering combination of primary partitions
  3. Evaluate the cost of all combinations generated in the previous phase

UNIVERSITAT POLITÈCNICA DE CATALUNYA BARCELONATECH
UPC

DTIM
www.essi.upc.edu/dtim

Data Allocation

- Given a set of <u>fragments</u>, a set of <u>sites</u> on which a number of <u>applications</u> are running, **allocate** each fragment such that some <u>optimization criterion</u> is met (subject to certain <u>constraints</u>)

- It is known to be an NP-hard problem
  - The optimal solution depends on many factors
    - Location in which the query originates
    - The query processing strategies (e.g., join methods)
  - Furthermore, in a dynamic environment the workload and access patterns may change
- The problem is typically simplified with certain assumptions
  - E.g., only communication cost considered
- Typical approaches build *cost models* and any optimization algorithm can be adapted to solve it
  - Sub-optimal solutions
  - Heuristics are also available
    - E.g., best-fit for non-replicated fragments

UNIVERSITAT POLITÈCNICA DE CATALUNYA BARCELONATECH
UPC

DTIM
www.essi.upc.edu/dtim

Data Replication

- Generalization of Allocation (for more than one location)

- Provides execution alternatives

- Improves availability

- Generates consistency problems
  - Specially useful for read-only workloads
    - No synchronization required

(Distributed) Catalog Management

Challenge II

DDBMS Catalog Characteristics

- Fragmentation
  - Global metadata
    - External schemas
    - Global conceptual schema
      - Fragment schema
      - Allocation schema
  - Local metadata
    - Local conceptual schema
    - Physical schema

- Allocation
  - Global metadata in the coordinator node
  - Local metadata in the workers

- Replication
  a) Single-copy (Coordinator node)
    - Single point of failure
    - Poor performance (potential bottleneck)
  b) Multi-copy (Mirroring, Secondary node)
    - Requires synchronization
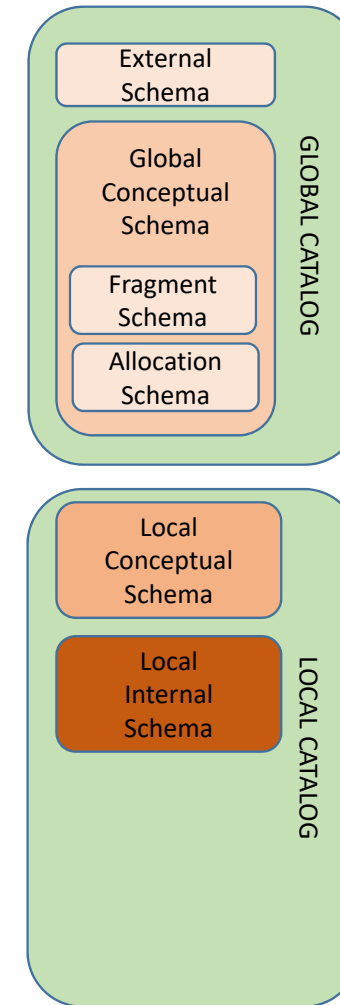
Closing

Summary

- Distributed Systems
- Distributed Database Systems
    - Distributed Database Systems Architectures
- Cloud Databases
- Distributed Database Design
    - Fragmentation
        - Kinds
        - Characteristics
    - Allocation
        - Replication
- Distributed Catalog

References

- D. DeWitt & J. Gray. *Parallel Database Systems: The future of High Performance Database Processing*. Communications of the ACM, June 1992

- N. J. Gunther. *A Simple Capacity Model of Massively Parallel Transaction Systems*. CMG National Conference, 1993

- L. Liu, M.T. Özsu (Eds.). *Encyclopedia of Database Systems*. Springer, 2009

- M. T. Özsu & P. Valduriez. *Principles of Distributed Database Systems*, 3rd Ed. Springer, 2011

- G. Coulouris et al. *Distributed Systems: Concepts and Design*, 5th Ed. Addisson-Wesley, 2012

UNIVERSITAT POLITÈCNICA DE CATALUNYA BARCELONATECH

DTIM
www.essi.upc.edu/dtim