

Spark II

Big Data Management

Knowledge objectives

1. Define RDD
2. Distinguish between Base RDD and Pair RDD
3. Distinguish between transformations and actions
4. Explain available transformations
5. Explain available actions
6. Name the main Spark runtime components
7. Explain how to manage parallelism in Spark
8. Explain how recoverability works in Spark
9. Distinguish between narrow and wide dependencies
10. Name the two mechanisms to share variables
11. Enumerate some abstraction on top of Spark

Application Objectives

- Provide the Spark pseudo-code for a simple problem using RDDs

Resilient Distributed Datasets

Resilient Distributed Datasets

- RDD
 - Resilient: Fault-tolerant
 - Distributed: Partitioned and parallel
 - Dataset: a set of data



"Unified abstraction for cluster computing, consisting in a read-only, partitioned collection of records. Can only be created through deterministic operations on either (1) data in stable storage or (2) other RDDs."

```
rdd := spark.textFile("hdfs://...")
```

M. Zaharia

Types of RDDs in Spark

- Base RDD
 - $\text{RDD}\langle T \rangle$
- Pair RDDs
 - $\text{RDD}\langle K, V \rangle$
 - Particularly important for MapReduce-style operations
- Other specific types
 - Structured Stream
 - VertexRDD
 - EdgeRDD
 - ...

RDD

Object
Object
...
Object

Dataframe

T_1	T_2	...	T_n
$1/A_1$	$2/A_2$...	n/A_n
x	"x"	...	T
y	"y"	...	F
...
z	None	...	T

Characteristics

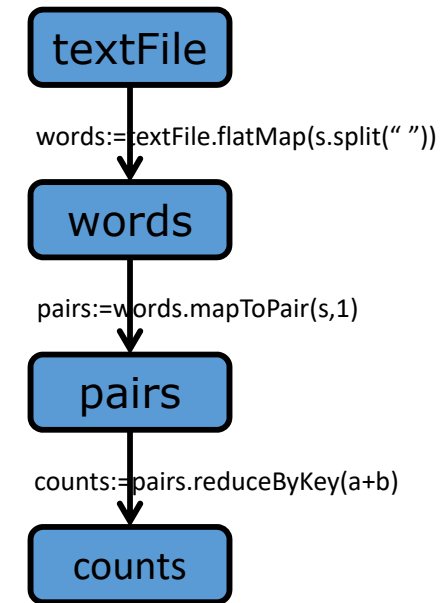
- Statically typed
- Parallel data structures
 - Disk
 - Memory
- User controls ...
 - Data sharing
 - Partitioning (fixed number per RDD)
 - Repartition (shuffles data through disk)
 - Coalesce (reduces partitions in the same worker)
- Rich set of coarse-grained operators
 - Simple and efficient programming interface
- Fault tolerant
- Baseline for more abstract applications

MapReduce vs Spark

	MapReduce	Spark RDD
Records	Key-Value pairs	Arbitrary
Storage	Results always in disk	Results can simply stay in memory
Functions	Only two	Rich palette
Partitioning	Statically decided	Dynamically decided

Example: Word count (Java)

```
JavaRDD<String> textFile = sc.textFile("hdfs://...");
JavaRDD<String> words = textFile.flatMap(s -> {
    return Arrays.asList(s.split(" "))
});
JavaPairRDD<String, Integer> pairs = words.mapToPair(s -> {
    return new Tuple2<String, Integer>(s, 1);
});
JavaPairRDD<String, Integer> counts = pairs.reduceByKey(a,b -> {
    return a + b;
});
counts.saveAsTextFile("hdfs://...");
```



Transformations and Actions

Apache Spark

Transformations vs. Actions

- Transformations
 - Applied to RDDs and generate new RDDs
 - They are run lazily
 - Only run when required to complete an action
- Actions
 - Trigger the execution of a pipeline of transformations
 - The result is ...
 - a) ... a primitive data type (not an RDD)
 - b) ... data written to an external storage system

Transformations on base RDDs


`map(f:T→U): RDD[T]→RDD[U]`

`filter(f:T→bool): RDD[T]→RDD[T]`

`sample(fraction: Float): RDD[T]→RDD[T]` (deterministic)

`flatMap(f:T→seq[U]): RDD[T]→RDD[U]`

`union/intersection/substract(): (RDD[T],RDD[T])→RDD[T]`

 `cartesian(): (RDD[T1],RDD[T2])→RDD[(T1,T2)]`

`partitionBy(p:partitioner[T]): RDD[T]→RDD[T]`

`sort(c:comparator[T]): RDD[T]→RDD[T]`

`distinct(T): RDD[T]→RDD[T]`

`persist(): RDD[T]→RDD[T]`

`mapToPair(f:T→(K,V)): RDD[T]→RDD[(K,V)]` (can be implicit)

<https://spark.apache.org/docs/latest/api/java/index.html?org/apache/spark/api/java/RDD.html>

Added transformations on pair RDDs

- ! `mapValues(f:V→W): RDD[(K,V)]→RDD[(K,W)]`
- `reduceByKey(f:(V,V)→V): RDD[(K,V)]→RDD[(K,V)]`
- `groupByKey(): RDD[(K,V)]→RDD[(K,seq(V))]`
- `join(): (RDD[(K,V)],RDD[(K,W)])→RDD[(K,(V,W))]`
- `cogroup(): (RDD[(K,V)],RDD[(K,W)])→RDD[(K,(seq[V],seq[W])]`
- `partitionBy(p:partitioner[K]): RDD[(K,V)]→RDD[(K,V)]`
- `sortByKey(): RDD[T]→RDD[T]`
- `keys(): RDD[(K,V)] → RDD[K]` (can be implicit)
- `values(): RDD[(K,V)] → RDD[V]` (can be implicit)

<https://spark.apache.org/docs/latest/api/java/index.html?org/apache/spark/api/java/JavaPairRDD.html>

Actions on base RDDs

`save(path: String)`: Writes the RDD to external storage (e.g., HDFS)

☠ `collect()`: $\text{RDD}[T] \rightarrow \text{seq}[T]$

`take(k)`: $\text{RDD}[T] \rightarrow \text{seq}[T]$

`first()`: $\text{RDD}[T] \rightarrow T$

`count()`: $\text{RDD}[T] \rightarrow \text{Long}$

`countByKey()`: $\text{RDD}[T] \rightarrow \text{seq}[(T, \text{Long})]$

`reduce(f: (T, T) \rightarrow T)`: $\text{RDD}[T] \rightarrow T$

`foreach(f: \rightarrow -)`: $\text{RDD}[T] \rightarrow -$ (executes in the workers)

Added actions on pair RDDs

`countByKey(): RDD[(K,V)]→seq[(K,Long)]`

`lookup(k: K): RDD[(K,V)]→seq[V]`

<https://spark.apache.org/docs/latest/api/java/index.html?org/apache/spark/api/java/JavaPairRDD.html>

Example

Analyzing HR data with RDDs

Average satisfaction level

- Does the number of projects an employee works on affect their satisfaction level?
- CSV Dataset (HR_comma_sep.csv)
 - Satisfaction Level
 - Last evaluation
 - Number of projects
 - Time spent at the company (in months)
 - Salary

Sample data

```
satisfaction_level, ...  
0.38,0.53,2,3,low  
0.8,0.86,5,6,medium  
0.11,0.88,7,4,medium  
0.72,0.87,5,5,low  
0.37,0.52,2,3,low  
0.41,0.5,2,3,low  
0.1,0.77,6,4,low  
0.92,0.85,5,5,high  
...
```

<https://www.kaggle.com/liujiaqi/hr-comma-sepcsv>

Implementation (Python)







"Average satisfaction level per number of projects, ordered from lowest to highest"

```
sc = pyspark.SparkContext.getOrCreate()

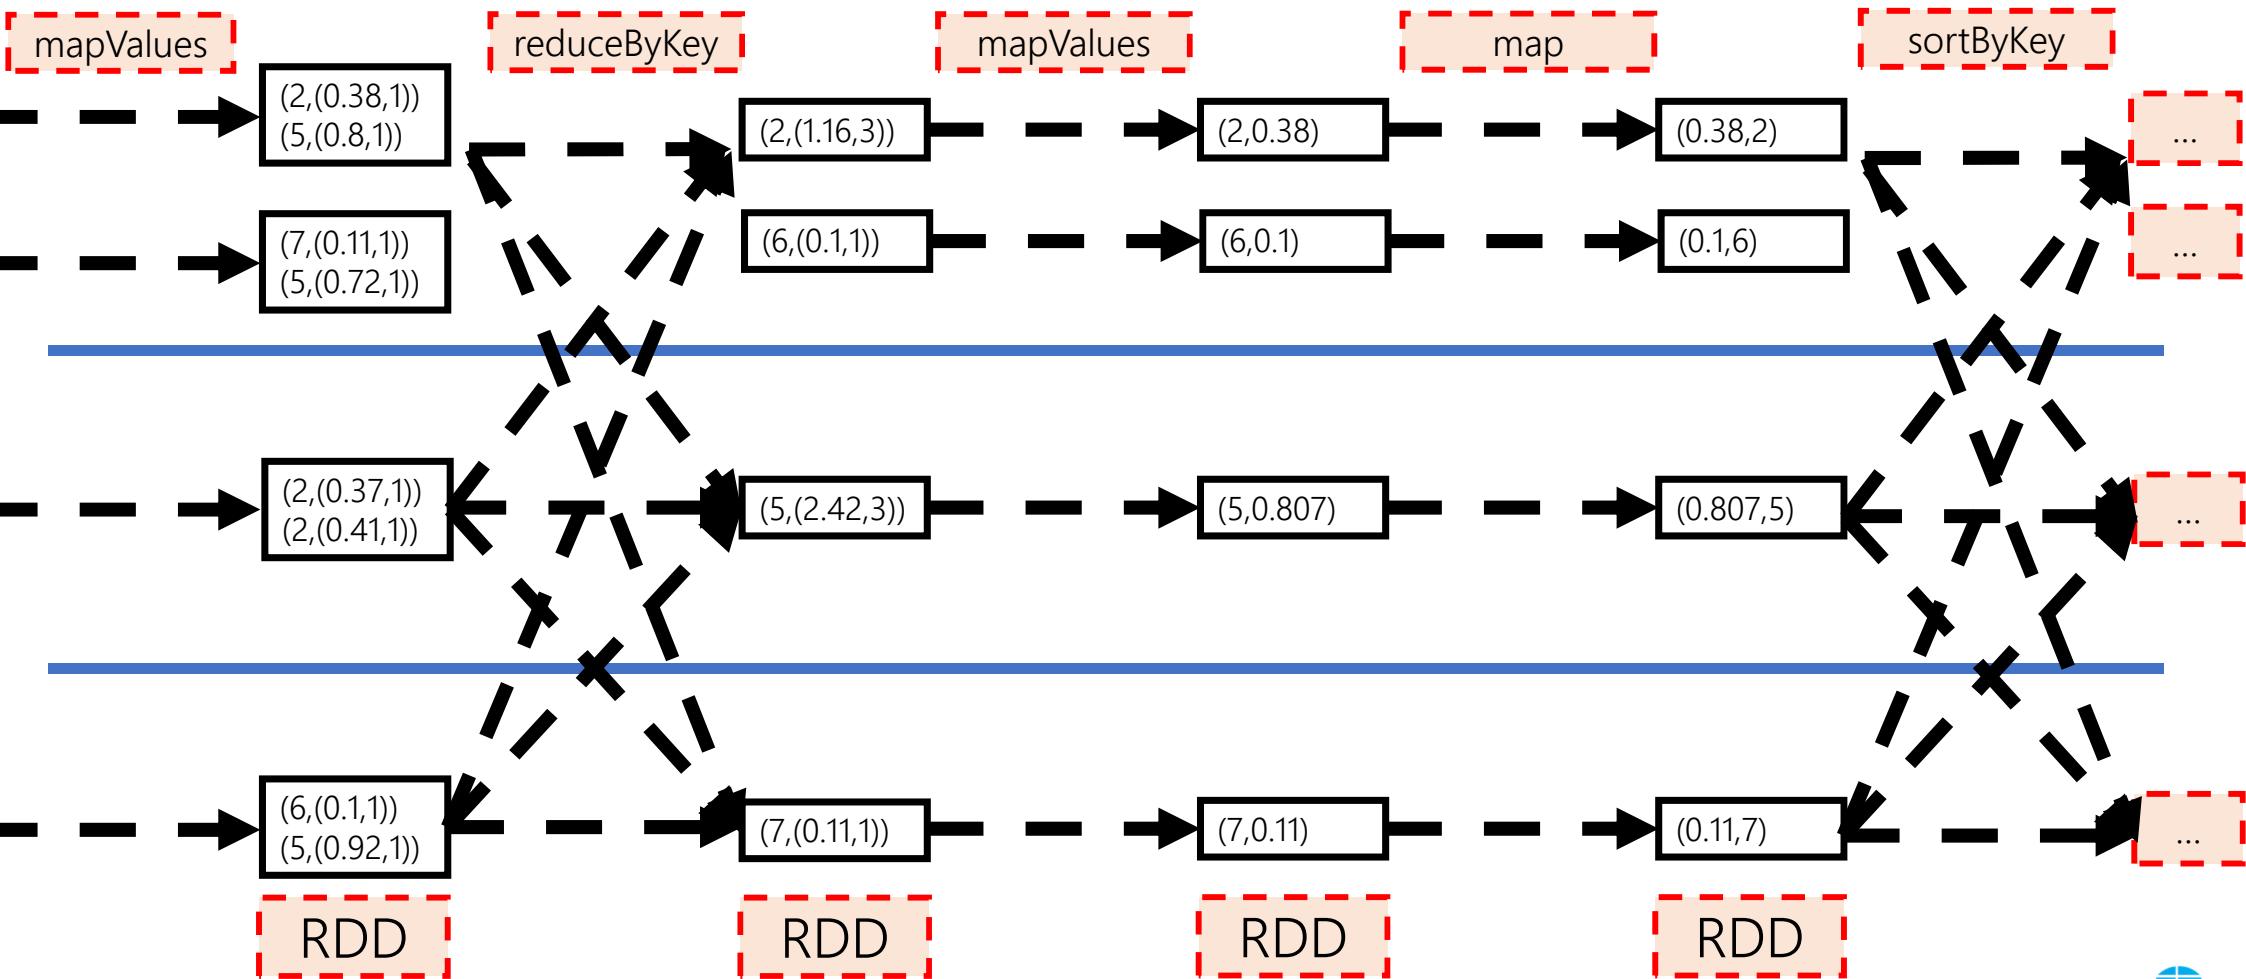
out = sc.textFile("HR_comma_sep.csv") \
    .filter(lambda t: "satisfaction level" not in t) \
    .map(lambda t: (int(t.split(",")[2]), float(t.split(",")[0]))) \
    .mapValues(lambda t: (t,1)) \
    .reduceByKey(lambda a,b: (a[0]+b[0],a[1]+b[1])) \
    .mapValues(lambda t: t[0]/t[1]) \
    .map(lambda t: (t[1],t[0])) \
    .sortByKey()

for x in out.collect():
    print(x)
```

Runtime execution (I)



Runtime execution (II)

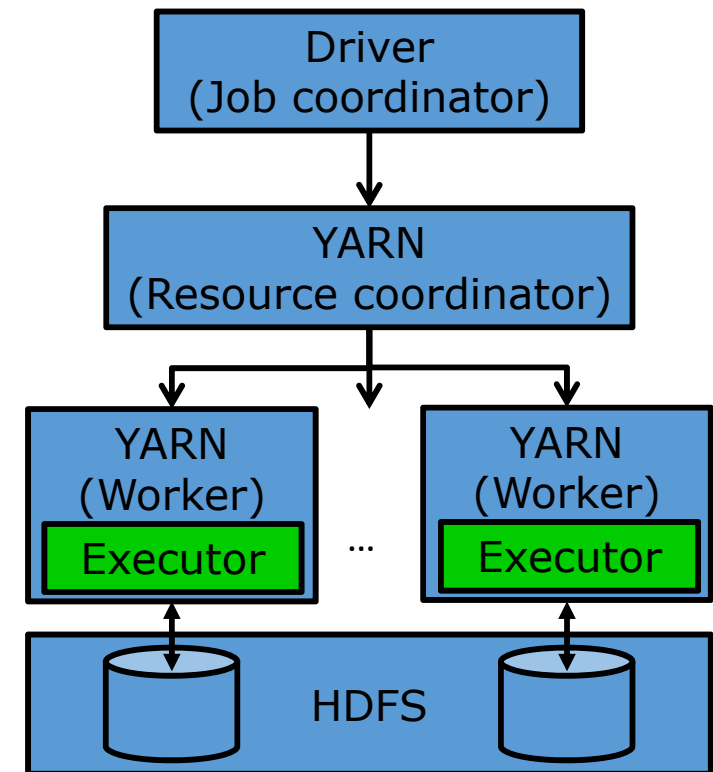


Under the hood

Apache Spark

Runtime architecture

- Driver (Job coordinator)
 - Creates the context
 - Decides on RDDs
 - Converts a program into tasks
 - Schedules tasks
 - Tracks location of cached data
- YARN (Resource coordinator)
 - Resource manager
- Executors (Job worker)
 - Run tasks
 - Store data



RDD Abstraction Representation

- A set of dependencies on parent RDDs
- A function for computing the dataset
- Partitioning schema/metadata
 - Hash
 - Range
- A set of partitions
- Data placement
 - Partitions per node

Parallelism

- Too few parallelism
 - Wastes resources
 - Hinders work balance
- Too much parallelism
 - May generate significant overheads
- Degree is automatically inferred from partitions


Partitioning

- Initially based on data locality
 - Useful based on keys
 - Hash
 - *partitionBy*
 - *groupByKey*
 - Range
 - *sortByKey*
- Partitions kept in workers' memory
- Different RDDs can use the same key
 - Similar to vertical partitioning
- Transformations lose partitioning information
 - *mapValues* retains partitioning information


Optimization

- Lineage graph is translated into a physical execution plan:
 - Truncate the lineage graph to use cached results
 - Pipeline or collapse several RDD into one stage
 - If no data movement needed
- Decompose one job into several stages
 - Stages are decomposed into tasks per partition
 - Each task has three phases:
 1. Fetch data (from either local or remote disk)
 2. Execute operations
 3. Write result (for shuffling or returning results to driver)


Narrow Dependencies



MapValue/Filter




Union




Join with inputs co-partitioned

Wide Dependencies




groupByKey



Join with inputs not co-partitioned

Scheduling



Recovery


- An RDD has enough information to be reconstructed after a failure
 - Lineage graph (logging not needed)
- Data can be cached/persisted (in up to two nodes)
 - Orthogonal to persistency options
 - Rule of thumb: cache an RDD if it is parent of more than one RDD

Storage Level	Meaning
MEMORY_ONLY	Store RDD as deserialized Java objects in the JVM. If the RDD does not fit in memory, some partitions will not be cached and will be recomputed on the fly each time they're needed. This is the default level.
MEMORY_AND_DISK	Store RDD as deserialized Java objects in the JVM. If the RDD does not fit in memory, store the partitions that don't fit on disk, and read them from there when they're needed.
MEMORY_ONLY_SER (Java and Scala)	Store RDD as <i>serialized</i> Java objects (one byte array per partition). This is generally more space-efficient than deserialized objects, especially when using a fast serializer , but more CPU-intensive to read.
MEMORY_AND_DISK_SER (Java and Scala)	Similar to MEMORY_ONLY_SER, but spill partitions that don't fit in memory to disk instead of recomputing them on the fly each time they're needed.
DISK_ONLY	Store the RDD partitions only on disk.
MEMORY_ONLY_2, MEMORY_AND_DISK_2, etc.	Same as the levels above, but replicate each partition on two cluster nodes.
OFF_HEAP (experimental)	Similar to MEMORY_ONLY_SER, but store the data in off-heap memory . This requires off-heap memory to be enabled.

<https://spark.apache.org/docs/latest/programming-guide.html#rdd-persistence>

Shared variables

- Broadcast variables (`sparkContext.broadcast()`)
 - Usage
 - Passed as a serializable object to the context
 - Accessed by workers (read-only)
 - Guarantees
 - The value is sent only once to each worker
- Accumulators (`sparkContext.accumulator()`)
 - Usage
 - Initialized by the driver
 - Incremented by workers (write-only)
 - Value accessed by driver
 - Guarantees
 - Consistent inside actions
 - Unpredictable result inside transformations
 - In case of reexecution



Closing

Summary

- Abstractions
- Resilient Distributed Datasets
 - Operations
 - Transformations
 - Actions
 - Persisting
 - Architecture
 - Dependencies
 - Scheduling
 - Partitioning

References

- H. Karau et al. *Learning Spark*. O'Really, 2015
- M. Zaharia. *An Architecture for Fast and General Data Processing on Large Clusters*. ACM Books, 2016
- A. Hogan. *Procesado de Datos Masivos* (Universidad de Chile). <http://aidanhogan.com/teaching/cc5212-1-2020>