(only FALSE ones are indicated or when there has been some ambiguity.)

1. **Complexity control and all that jazz**.

   - ❏ The VC dimension of a hypothesis class tells us how capable its members are to fit data.
   - ❏ In machine learning it is necessary to have classifiers with zero training error. **FALSE, zero training error may lead to overfitting for example**
   - ❏ Regularization is a framework in which poor goodness-of-fit can be compensated with complexity. **FALSE, here when I wrote "complexity" I meant high complexity, but I have also accepted TRUE answers if explained that low complexity was interpreted.**
   - ❏ The VC dimension can only be infinite for models with infinite parameters. **FALSE, in fact in class we saw an example of a function class with 1 parameter and infinite VC dimension**
   - ❏ Test data is typically used to estimate generalization error.
   - ❏ Validation data is typically used to perform model selection.
   - ❏ The VC dimension of a function strongly affects its variance during learning. **FALSE, "VC dimension of a function" makes no sense as VC dimension applies to a family of functions or hypothesis class.**
   - ❏ The VC dimension is related to the number of parameters of functions in a hypothesis class. **FALSE, not necessarily (example with 1 parameter and infinite dimension for example.)**
   - ❏ Empirical risk, the opposite of training error, serves as an approximation to the true risk. **FALSE, empirical risk is not the opposite of training error.**
   - ❏ In order to prove that VC dimension is strictly smaller than n, we have to show that every set of cardinality n cannot be shattered.

2. **Bayesian classifiers**.

   - ❏ Linear models are called linear because they are linear in their parameters.
   - ❏ By definition, Bayes rule chooses the class with highest posterior probability.
   - ❏ Bayesian classifiers lead to linear decision boundaries. **FALSE, not always, for example QDA leads to quadratic decision boundaries.**
   - ❏ The bayesian classifier is optimal in the sense that it has the smallest generalization error among all classifiers.
   - ❏ The Mahalanobis distance is a generalized form of Euclidean distance that takes into account correlations among variables.
   - ❏ If the class-conditional distributions in a classification problem are normal, then linear functions can achieve the best possible generalization error. **FALSE, in QDA we assume normal class-conditional distributions but optimal classifier is not linear in general.**
   - ❏ When using QDA or LDA in a practical problem with finite data that is normally distributed, there is no need to estimate generalization error because we know that these methods are optimal. **FALSE, with finite data we still need to estimate the parameters of the gaussians and therefore we still need to check quality of the trained models.**
   - ❏ Laplace correction is particularly necessary in Naive Bayes classification if the dimensionality of the input data is very large. **My particular answer here would be TRUE, however I have accepted some FALSE ones that have argued that the problem is with 0 counts. The reasoning behind my TRUE is that the "probabiltiy" that 0 counts happen grows exponentially with the dimensionality of the input, and therefore we need to deal with this.**
   - ❏ The VC dimension of the k-NN classifier is infinite.

- ❏ Tuning the "k" in k-NN trades off bias and variance.

3. **Maximum Likelihood and GLMs**.

   - ❏ The maximum likelihood method is a general method based on the optimization of the likelihood function which is a function determined by a data sample that maps parameter values to non-negative real values.
   - ❏ The likelihood function is a probability distribution over all possible parameter values for finite sets of data. **FALSE, if you add (or integrate) over all parameters values the results does not have to be 1 (for example).**
   - ❏ GLMs are typically trained using the maximum likelihood method which involves solving some form of linear optimization problem. **FALSE, optimizing likelihood typically involves non-linear optimization techniques.**
   - ❏ Logistic regression is considered a (generalized) linear model because the classes' posteriors are linear functions of the predictors. **FALSE, the log of the odds is a linear function of the predictors, not the posteriors.**
   - ❏ The Iterated Reweighted Least Squares algorithm is used to find the maximum likelihood optimum in logistic regression.
   - ❏ The decision boundary of a logistic regression model has a sigmoidal shape. **FALSE, the decision boundary is linear.**
   - ❏ The likelihood function to be optimized when modelling with GLM is determined by the target distribution $P(T|\mathbf{X};\theta)$ assumed in the model.
   - ❏ Zero bias and variance of a linear regression model implies zero training error. **FALSE, we still have noise in the training data.**
   - ❏ Fitting non-linear models requires non-linear optimization techniques. **FALSE, this is not always so, e.g. in SVMs we still use linear techniques to fit nonlinear functions by making use of the kernel trick.**
   - ❏ The "generalized" part of GLM is due to their use of link functions.

4. **Neural networks**.

   - ❏ MLPs generalize GLMs
   - ❏ Once trained, MLP networks are deterministic while RBF networks are not. **FALSE, both MLPs and RBFs are deterministic.**
   - ❏ The only differences between MLP and RBF networks are the number of hidden layers they allow and their training mechanisms. **FALSE, they use different "types" of nodes as well.**
   - ❏ Since GLMs are particular cases of MLPs, the backpropagation algorithm is enough to train a logistic model algorithm. **FALSE, the backpropagation algorithm only computes gradients, we need to use these gradients in some way to update parameters.**
   - ❏ Any cost function, as long as it is differentiable, can be used in the backpropagation algorithm to find suitable weights.
   - ❏ Computing the gradients during MLP training requires an iterative procedure which may not always converge. **FALSE, we can use backpropagation which always finishes, all it does is one forward and one backward step through the layers of the network.**
   - ❏ Backpropagation is a non-deterministic procedure which, depending on how initial values are chosen may give different results. **FALSE, backpropagation is totally deterministic.**
   - ❏ Gradient descent is a non-deterministic procedure which, depending on how initial values are chosen may give different results.
   - ❏ The number of neurons in hidden layers may have a large impact on the quality of the predictions of a neural network.

- ❑ RBF neural networks can be trained using backpropagation. **FALSE, the first layer is "trained" using some sort of clustering, and the second layer can be computed using some sort of gradient descent method, for example.**

5. **Kernels and SVMs**.

   - ❑ Adding two kernels k_1 and k_2 with corresponding transformations ϕ_1 and ϕ_2 yields a new kernel k_3 which corresponds to the data transformation $\phi_3 : \mathbf{x} \mapsto \phi_1(\mathbf{x}) + \phi_2(\mathbf{x})$. **FALSE, adding two kernels corresponds to the concatenation of transformations.**
   - ❑ Multipying a kernel k by a strictly positive scalar yields a new kernel k' whose corresponding transformation is of dimensionality at most that of k's transformation.
   - ❑ Since adding kernels and multiplying them by scalars yield new kernels, any polynomial on a given kernel is also a kernel. **FALSE, for example -k is a polynomial on k (k being a kernel) and it is not a kernel.**
   - ❑ There is no danger of overfitting in SVMs because seeking large margins prevents it. **FALSE, using some kernels effectively adds complexity to the models and therefore there is danger of overfitting.**
   - ❑ Increasing the value of C (all else being equal) in an SVM may decrease training error.
   - ❑ Kernel matrices are always positive semi-definite because they correspond to inner products in some Hilbert space. **My answer here is TRUE but I also accepted FALSE answers that pointed out the existence of CPSD kernels.**
   - ❑ SVMs produce linear classifiers with high margin if we chose an appropriate kernel. **FALSE, for some data this is not possible**
   - ❑ An algorithm can be kernelized if its solution can be expressed as a linear combination of input vectors.
   - ❑ Kernels are defined as functions of two input vectors in Euclidean space whose corresponding kernel matrices are positive semi-definite. **FALSE, kernels do not necessarily need to operate on vectors in Euclidean space.**
   - ❑ Linear SVMs tend to generalize better because simpler models have lower VC dimension. **FALSE, depends on the data.**

6. **Miscellaneous**.

   - ❑ Unsupervised learning is about making predictions on unseen future examples. **FALSE, in unsupervised learning there is no notion of making predictions.**
   - ❑ High training and test error is an indicator of high bias in a model.
   - ❑ High training error and low test error is an indicator of overfitting in a model. **FALSE, it would be low training and high test error.**
   - ❑ Random forests are typically better predictors than individual decision trees due to the high variance of individual trees, which random forests reduce.
   - ❑ The E-M algorithm is an optimization algorithm that maximizes some likelihood function.
   - ❑ The E-M algorithm is guaranteed to find an optimum solution if the input data is distributed according to a mixture of gaussians. **FALSE, the E-M algorithm only finds local optima.**
   - ❑ If all elements in a mixture of gaussians have equal covariance matrices, then E-M behaves like k-means. **FALSE, the covariance matrices also need to be proportional to the diagonal matrix**
   - ❑ Clustering makes little sense in practice since we have no way of knowing whether the result is fully truthful. **FALSE, clustering can be very useful in practice and it is very much used.**
   - ❑ Classification is always easier than regression or clustering. **FALSE, it depends on the particular problem (data).**
   - ❑ Clustering can be hard because in most cases we have no gold standard.
